## Computer Organization and Architecture

Instructor: Dr. Rushdi Abu

Zneit

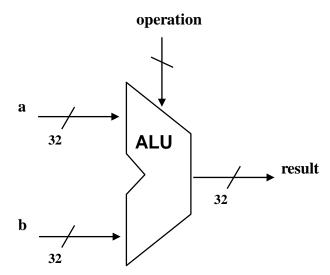
Slide Sources: Patterson &

Hennessy

# Ch. 4 Arithmetic for Computers

#### Arithmetic

- Where we've been:
  - performance
  - abstractions
    - instruction set architecture
    - assembly language and machine language
- What's up ahead:
  - implementing the architecture



#### Numbers

- Bits are just bits (no inherent meaning)
  - conventions define relationship between bits and numbers
- Binary integers (base 2)
  - 0000 0001 0010 0011 0100 0101 0110 0111 1000 1001...
  - decimal: 0, ..., 2<sup>n</sup>-1

n bits

- Of course it gets more complicated:
  - bit strings are finite, but
    - for some fractions and real numbers, finitely many bits is not enough, so
    - overflow & approximation errors: e.g., represent 1/3 as binary!
  - negative integers
- How do we represent negative integers?
  - which bit patterns will represent which integers?

#### Possible Representations



$$000 = 0$$
  
 $001 = +1$   
 $010 = +2$   
 $011 = +3$   
 $100 = 0$   
 $101 = -1$   
 $110 = -2$   
 $111 = -3$ 

One's Complement

$$000 = 0$$
  
 $001 = +1$   
 $010 = +2$   
 $011 = +3$   
 $100 = -3$   
 $101 = -2$   
 $110 = -1$   
 $111 = 0$ 

Two's Complement

```
and positives, unique zero and positives, unique zero and positives; unique zero 000 = +1 010 = +3 101 = -3 111 = -1 111 = -1
```

- Issues:
  - balance equal number of negatives and positives
  - ambiguous zero whether more than one zero representation
  - ease of arithmetic operations
- Which representation is best? Can we get both balance and non-ambiguous zero?

### Representation Formulae

Two's complement:

$$x_n \, x_{n-1} ... x_0 \; = \; x_n \, * \, -2^n \, + \, x_{n-1} \, * \, 2^{n-1} \, + \, ... \, + \, x_0 \, * \, 2^0$$
 
$$or$$
 
$$x_n X' \; = \; x_n \, * \, -2^n \, + \, X' \quad \text{(writing rightmost n bits $x_{n-1} ... x_0$ as $X'$)}$$
 
$$= \; \left\{ \begin{array}{l} X', \qquad \text{if $x_n = 0$} \\ -2^n \, + \, X', & \text{if $x_n = 1$} \end{array} \right.$$

One's complement:

$$x_n X' = \begin{cases} X', & \text{if } x_n = 0 \\ -2^n + 1 + X', & \text{if } x_n = 1 \end{cases}$$

#### MIPS – 2's complement

32 bit signed numbers:

```
0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 =
0000 0000 0000 0000 0000 0000 0000 1_{two} = + 1_{ten}
0000 0000 0000 0000 0000 0000
                               0010_{two} = + 2_{ten}
                                                     maxint
   0000 0000 0000 0000 0000 0000 0000_{two} = -2,147,483,648_{ten}
1000
                               0001_{\text{two}} = -2,147,483,647_{\text{ten}}
    0000 0000 0000 0000 0000
1000
                               0010_{two} = -2,147,483,646_{ten}
    0000 0000 0000 0000 0000
        1111 1111 1111 1111 1111 1101_{two} = -3_{ten}
                                                     minint
```

Negative integers are exactly those that have leftmost bit 1

#### Two's Complement Operations

- Negation Shortcut: To negate any two's complement integer
   (except for minint) invert all bits and add 1
  - note that negate and invert are different operations!
  - why does this work? Remember we don't know how to add in 2's complement yet! Later...!
- <u>Sign Extension Shortcut</u>: To convert an n-bit integer into an integer with more than n bits i.e., to make a narrow integer fill a wider word *replicate the most significant bit* (*msb*) of the original number to fill the new bits to its left

```
Example: 4-bit 8-bit 0010 = 0000 0010
1010 = 1111 1010
```

• why is this correct? Prove!

#### MIPS Notes

- lb VS. lbu
  - signed load sign extends to fill 24 left bits
  - unsigned load fills left bits with 0's
- slt & slti
  - compare signed numbers
- sltu & sltiu
  - compare unsigned numbers, i.e., treat both operands as non-negative

#### Two's Complement Addition

- Perform add just as in junior school (carry/borrow 1s)
  - Examples (4-bits):

```
      0101
      0110
      1011
      1001
      1111

      0001
      0101
      0111
      1010
      1110
```

Do these sums **now**!! Remember all registers are 4-bit including result register! So you have to **throw away** the carry-out from the msb!!

- Have to beware of overflow: if the fixed number of bits (4, 8, 16, 32, etc.) in a register cannot represent the result of the operation
  - terminology alert: overflow does not mean there was a carry-out from the msb that we lost (though it sounds like that!) – it means simply that the result in the fixed-sized register is incorrect
    - as can be seen from the above examples there are cases when the result is correct even after losing the carry-out from the msb

# Two's Complement Addition: Verifying Carry/Borrow method

Iwo (n+1)-bit integers: $X = x_n X'$ , $Y = y_n Y'$			
Carry/borrow	$0 \le X' + Y' < 2^n$	$2^n \le X' + Y' < 2^{n+1} -$	
add X + Y	(no CarryIn to last bit)	1	
$x_n = 0, y_n = 0$	ok	(CarryIn to last bit) not ok(overflow!)	
$x_n = 1, y_n = 0$	ok	ok	
$x_n = 0, y_n = 1$	ok	ok	
$x_n = 1, y_n = 1$	not ok(overflow!)	ok	

- Prove the cases above!
- Prove if there is *one more bit* (total n+2 then) available for the result then there is no problem with overflow in add!

## Two's Complement Operations

- Now verify the negation shortcut!
  - consider X + (X +1) = (X + X) + 1:
     associative law but what if there is overflow in one of the adds on either side, i.e., the result is wrong...!
  - think minint!
  - Examples.
    - -0101 = 1010 + 1 = 1011
    - -1100 = 0011 + 1 = 0100
    - $-1000 \neq 0111 + 1 = 1000$

#### **Detecting Overflow**

- No overflow when adding a positive and a negative number
- No overflow when subtracting numbers with the same sign
- Overflow occurs when the result has "wrong" sign (verify!):

Operation	Operand A	Operand B	Result Indicating Overflow
A + B	≥ 0	≥ 0	< 0
A + B	< 0	< 0	≥ 0
А — В	≥ 0	< 0	< 0
А — В	< 0	≥ 0	≥ 0

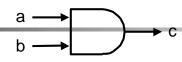
- Consider the operations A + B, and A − B
  - can overflow occur if B is 0?
  - can overflow occur if A is 0?

#### **Effects of Overflow**

- If an exception (interrupt) occurs
  - control jumps to predefined address for exception
  - interrupted address is saved for possible resumption
- Details based on software system/language
  - SPIM: see the EPC and Cause registers
- Don't always want to cause exception on overflow
  - add, addi, sub cause exceptions on overflow
  - addu, addiu, subu do not cause exceptions on overflow

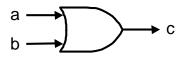
#### Review: Basic Hardware

1. AND gate (c = a . b)



а	b	c = a . b	
0	0	0	
0	1	0	
1	0	0	
1	1	1	

2. OR gate (c = a + b)



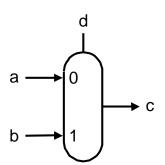
а	b	c = a + b
0	0	0
0	1	1
1	0	1
1	1	1

3. Inverter (c = a)



а	$c = \bar{a}$
0	1
1	0

4. Multiplexor (if d = 0, c = a; else c = b)



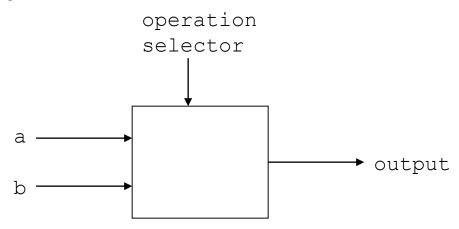
d	С
0	а
1	b

## Review: Boolean Algebra & Gates

- Problem: Consider logic functions with three inputs: A, B, C.
  - output D is true if at least one input is true
  - output E is true if exactly two inputs are true
  - output F is true only if all three inputs are true
- Show the truth table for these three functions
- Show the Boolean equations for these three functions
- Show an implementation consisting of inverters,
   AND, and OR gates.

## A Simple Multi-Function Logic Unit

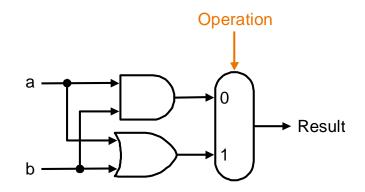
- To warm up let's build a logic unit to support the and or instructions for MIPS (32-bit registers)
  - we'll just build a 1-bit unit and use 32 of them



Possible implementation using a multiplexor :

# Implementation with a Multiplexor

 Selects one of the inputs to be the output based on a control input

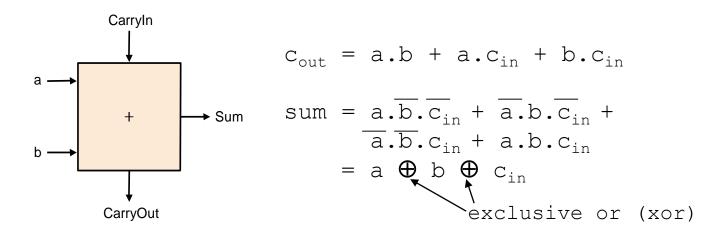


32 units

Lets build our ALU using a MUX (multiplexor):

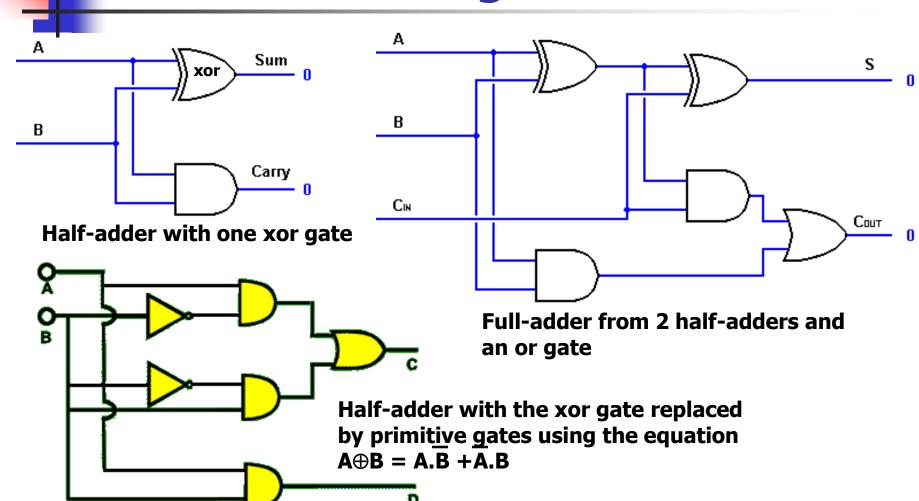
#### **Implementations**

- Not easy to decide the best way to implement something
  - do not want too many inputs to a single gate
  - do not want to have to go through too many gates (= levels)
  - for our purposes, ease of comprehension is important
- Let's look at a 1-bit ALU for addition:

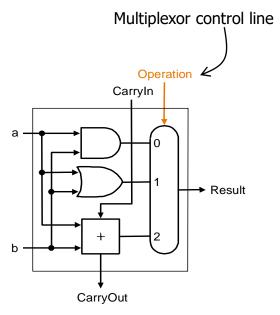


- How could we build a 1-bit ALU for add, and, and or?
- *How could we build a 32-bit ALU?*

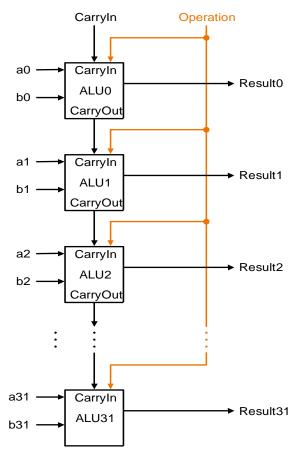
## 1-bit Adder Logic



#### Building a 32-bit ALU



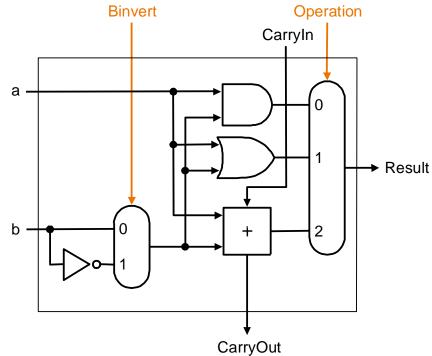
1-bit ALU for AND, OR and add



Ripple-Carry Logic for 32-bit ALU

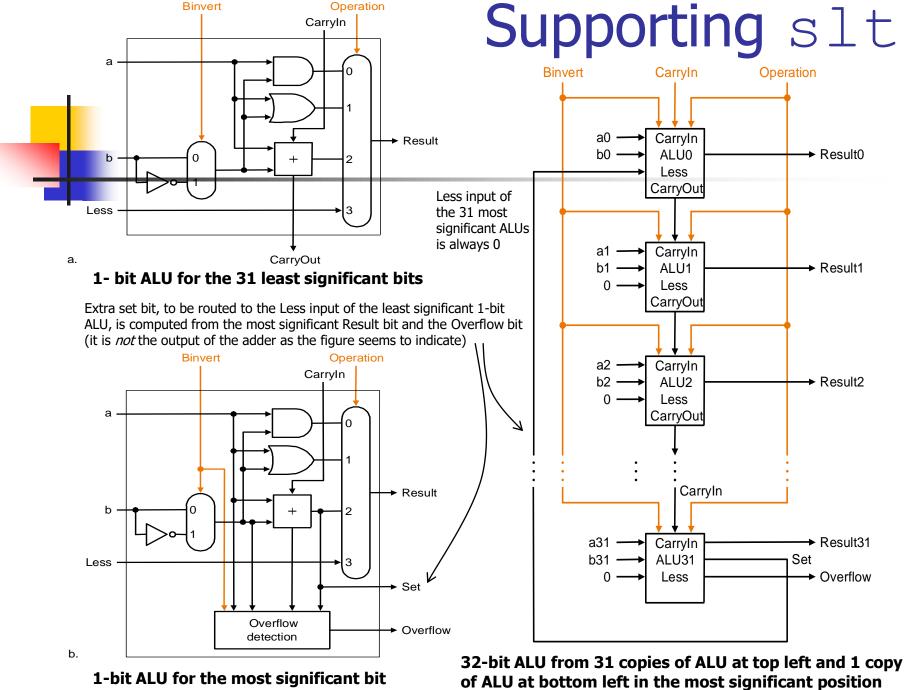
#### What about Subtraction (a - b)?

- Two's complement approach: just negate b and add.
- How do we negate?
  - recall negation shortcut: invert each bit of b and set CarryIn to least significant bit (ALU0) to 1



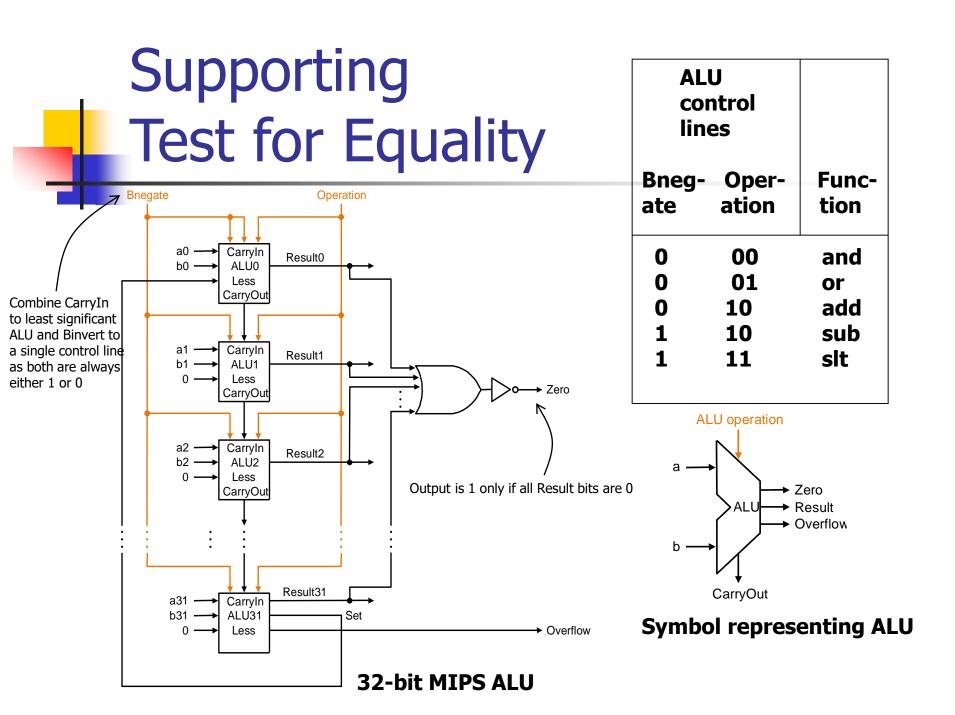
## Tailoring the ALU to MIPS: Test for Less-than and Equality

- Need to support the set-on-less-than instruction
  - e.g., slt \$t0, \$t3, \$t4
  - remember: slt is an R-type instruction that produces 1 if rs < rt and 0 otherwise
  - idea is to use subtraction: rs < rt ⇔ rs − rt < 0. Recall msb of negative number is 1
  - two cases after subtraction rs rt:
    - if no overflow then rs < rt  $\Leftrightarrow$  most significant bit of rs rt = 1
    - if overflow then  $rs < rt \Leftrightarrow most significant bit of <math>rs rt = 0$
  - why?
  - e.g.,  $5_{ten} 6_{ten} = 0101 0110 = 0101 + 1010 = 1111$  (ok!)  $-7_{ten} 6_{ten} = 1001 0110 = 1001 + 1010 = 0011$  (overflow!)
  - therefore
    - set bit = msb of rs rt  $\oplus$  overflow bit where *set bit,* which is output from ALU31, gives the result of slt
      - Fig. 4.17(lower) indicates set bit is the adder output *not correct*!!
  - set bit is sent from ALU31 to ALU0 as the Less bit at ALU0; all other Less bits are hardwired 0; so Less is the 32-bit result of slt



## Tailoring the ALU to MIPS: Test for Less-than and Equality

- What about logic for the overflow bit?
  - overflow bit = carry in to msb ⊕ carry out of msb
  - verify!
  - logic for overflow detection therefore can be put in to ALU31
- Need to support test for equality
  - **e.g.**, beq \$t5, \$t6, \$t7
  - use subtraction: rs rt =  $0 \Leftrightarrow$  rs = rt
  - do we need to consider overflow?



#### Conclusion

- We can build an ALU to support the MIPS instruction set
  - key idea: use multiplexor to select the output we want
  - we can efficiently perform subtraction using two's complement
  - we can replicate a 1-bit ALU to produce a 32-bit ALU
- Important points about hardware
  - all gates are always working
  - speed of a gate depends number of inputs (fan-in) to the gate
  - speed of a circuit depends on number of gates in series
     (particularly, on the *critical path* to the deepest level of logic)
- Speed of MIPS operations
  - clever changes to organization can improve performance (similar to using better algorithms in software)
  - we'll look at examples for addition, multiplication and division

### Problem: Ripple-carry Adder is Slow

- Is a 32-bit ALU as fast as a 1-bit ALU? Why?
- Is there more than one way to do addition? Yes:
  - one extreme: *ripple-carry* carry ripples through 32 ALUs, slow!
  - other extreme: sum-of-products for each CarryIn bit super fast!
  - CarryIn bits:

$$c_1 = b_0.c_0 + a_0.c_0 + a_0.b_0$$

$$c_2 = b_1 \cdot c_1 + a_1 \cdot c_1 + a_1 \cdot b_1$$

= 
$$a_1 \cdot a_0 \cdot b_0 + a_1 \cdot a_0 \cdot c_0 + a_1 \cdot b_0 \cdot c_0$$
 (substituting for  $c_1$ )  
+  $b_1 \cdot a_0 \cdot b_0 + b_1 \cdot a_0 \cdot c_0 + b_1 \cdot b_0 \cdot c_0 + a_1 \cdot b_1$ 

$$c_3 = b_2.c_2 + a_2.c_2 + a_2.b_2$$
  
= ... = sum of 15 4-term products...

How fast? But not feasible for a 32-bit ALU! Why? Exponential complexity!!

Note: c<sub>i</sub> is CarryIn bit *into i* th ALU;

c<sub>0</sub> is the forced CarryIn into the

least significant ALU

## Two-level Carry-lookahead Adder: First Level

- An approach between our two extremes
- Motivation:
  - if we didn't know the value of a carry-in, what could we do?
  - when would we always generate a carry? (generate)  $g_i = a_i$  .  $b_i$
  - when would we propagate the carry? (propagate)  $p_i = a_i + b_i$
- Express (carry-in equations in terms of generate/propagates)

$$c_1 = g_0 + p_0 \cdot c_0$$
  
 $c_2 = g_1 + p_1 \cdot c_1 = g_1 + p_1 \cdot g_0 + p_1 \cdot p_0 \cdot c_0$   
 $c_3 = g_2 + p_2 \cdot c_2 = g_2 + p_2 \cdot g_1 + p_2 \cdot p_1 \cdot g_0 + p_2 \cdot p_1 \cdot p_0 \cdot c_0$   
 $c_4 = g_3 + p_3 \cdot c_3 = g_3 + p_3 \cdot g_2 + p_3 \cdot p_2 \cdot g_1 + p_3 \cdot p_2 \cdot p_1 \cdot g_0$   
 $+ p_3 \cdot p_2 \cdot p_1 \cdot p_0 \cdot c_0$ 

- Feasible for 4-bit adders with wider adders unacceptable complexity.
  - solution: build a first level using 4-bit adders, then a second level on top

# Two-level Carry-lookahead Adder: Second Level for a 16-bit adder

Propagate signals for each of the four 4-bit adder blocks:

$$P_0 = p_3 \cdot p_2 \cdot p_1 \cdot p_0$$
  
 $P_1 = p_7 \cdot p_6 \cdot p_5 \cdot p_4$   
 $P_2 = p_{11} \cdot p_{10} \cdot p_9 \cdot p_8$   
 $P_3 = p_{15} \cdot p_{14} \cdot p_{13} \cdot p_{12}$ 

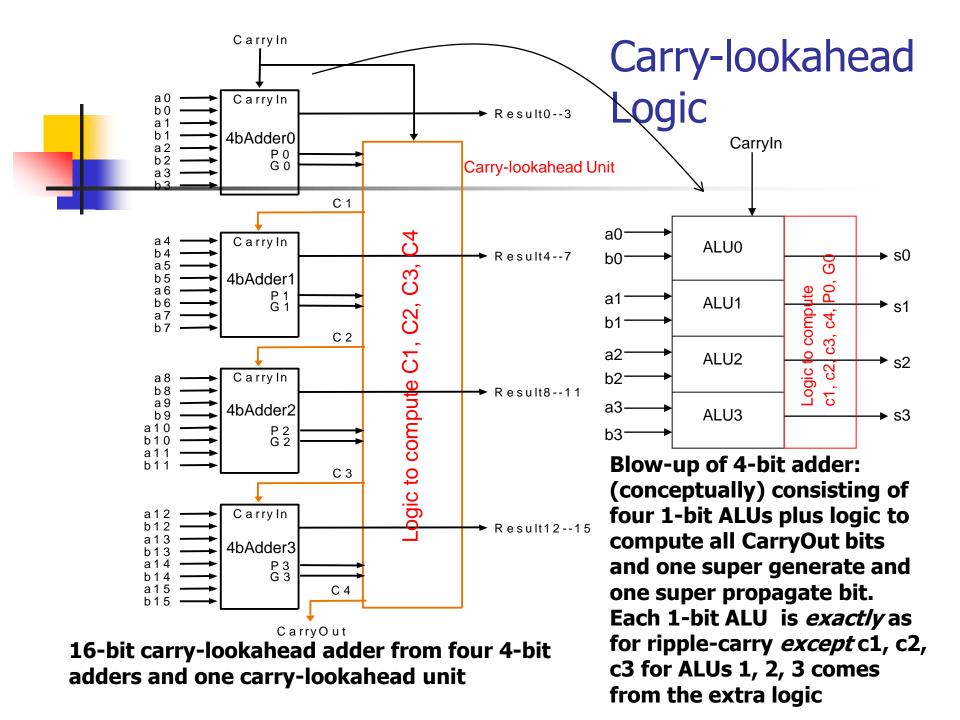
Generate signals for each of the four 4-bit adder blocks:

$$G_0 = g_3 + p_3.g_2 + p_3.p_2.g_1 + p_3.p_2.p_1.g_0$$
  
 $G_1 = g_7 + p_7.g_6 + p_7.p_6.g_5 + p_7.p_6.p_5.g_4$   
 $G_2 = g_{11} + p_{11}.g_{10} + p_{11}.p_{10}.g_9 + p_{11}.p_{10}.p_9.g_8$   
 $G_3 = g_{15} + p_{15}.g_{14} + p_{15}.p_{14}.g_{13} + p_{15}.p_{14}.p_{13}.g_{12}$ 

# Two-level Carry-lookahead Adder: Second Level for a 16-bit adder

 CarryIn signals for each of the four 4-bit adder blocks (see earlier carry-in equations in terms of generate/propagates):

$$C_1 = G_0 + P_0.c_0$$
  
 $C_2 = G_1 + P_1.G_0 + P_1.P_0.c_0$   
 $C_3 = G_2 + P_2.G_1 + P_2.P_1.G_0 + P_2.P_1.P_0.c_0$   
 $C_4 = G_3 + P_3.G_2 + P_3.P_2.G_1 + P_3.P_2.P_1.G_0 + P_3.P_2.P_1.P_0.c_0$ 



# Two-level Carry-lookahead Adder: Second Level for a 16-bit adder

- Two-level carry-lookahead logic steps:
  - 1. compute p<sub>i</sub>'s and g<sub>i</sub>'s at each 1-bit ALU
  - compute  $P_i$ 's and  $G_i$ 's at each 4-bit adder unit
  - compute C; 's in carry-lookahead unit
  - 4. compute  $C_i$ 's at each 4-bit adder unit
  - 5. compute results (sum bits) at each 1-bit ALU
- E.g., add using carry-lookahead logic:
  - 0001 1010 0011 0011
  - **1110 0101 1110 1011**
- Compare times for ripple-carry vs. carry-lookahead for a 16-bit adder assuming unit delay at each gate

### Multiply

Grade school shift-add method:

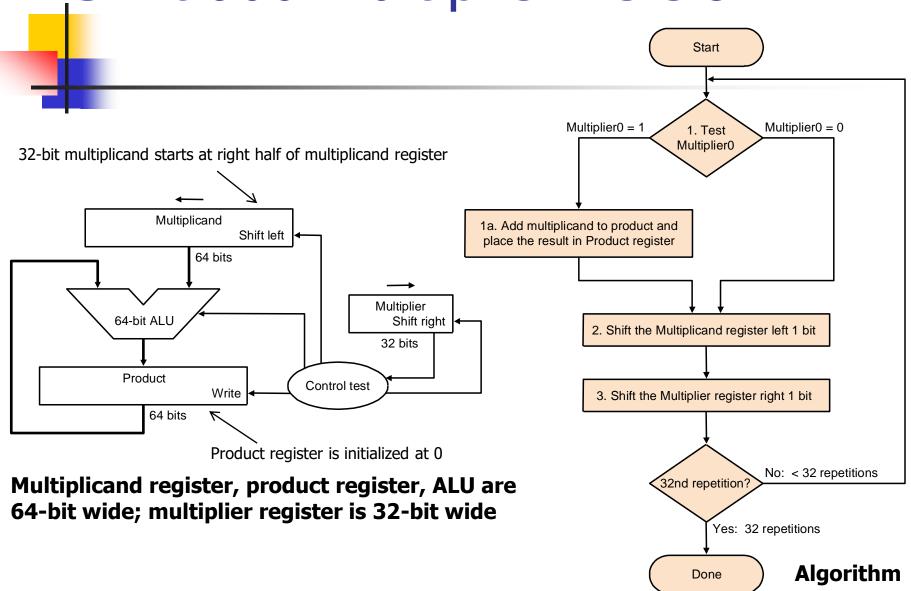
```
Multiplicand 1000

Multiplier x \frac{1001}{1000}
0000
0000
1000

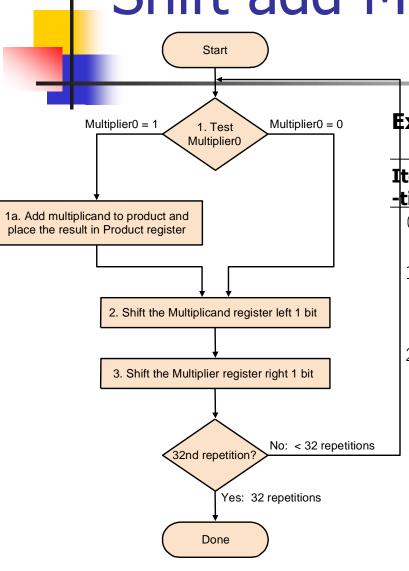
Product 01001000
```

- m bits x n bits = m+n bit product
- Binary makes it easy:
  - multiplier bit 1 => copy multiplicand (1 x multiplicand)
  - multiplier bit 0 => place 0 (0 x multiplicand)
- 3 versions of multiply hardware & algorithm:

#### Shift-add Multiplier Version 1



### Shift-add Multiplier Version1



**Example:** 0010 \* 0011:

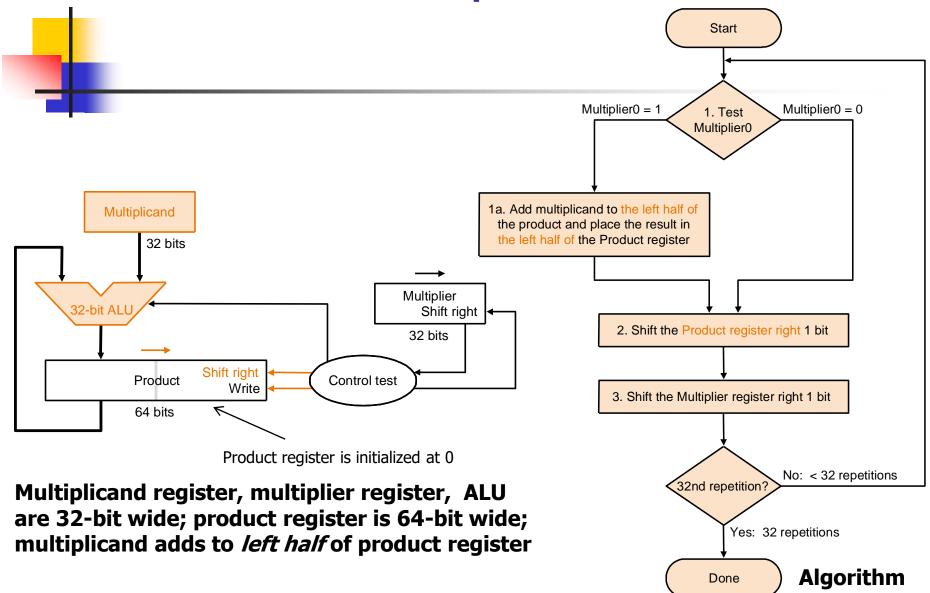
Itera	•	Multiplier	Multiplicand	Product
0	init	0011	0000 0010	0000 0000
	values			
1	1a	0011	0000 0010	0000 0010
	2	0011	0000 0100	0000 0010
	3	0001	0000 0100	0000 0010
2				

**Algorithm** 

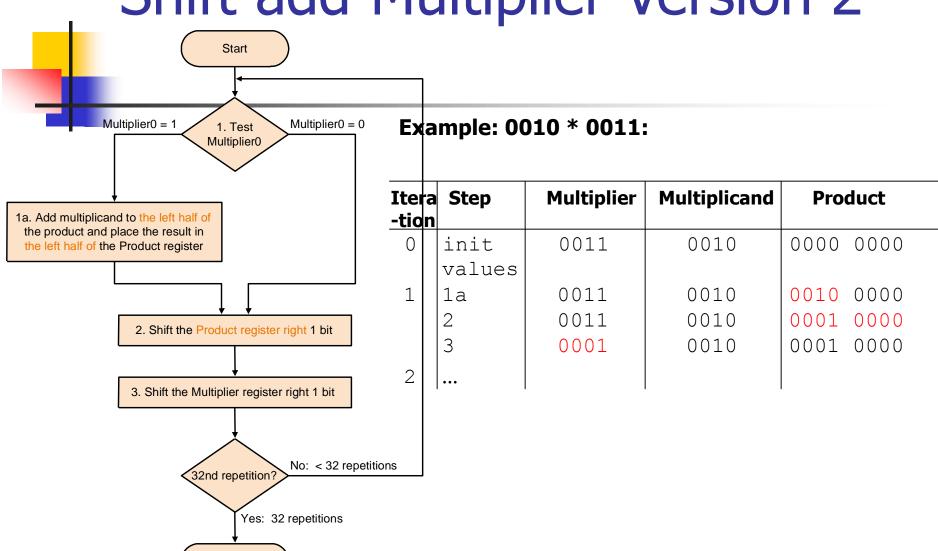
# Observations on Multiply Version 1

- 1 step per clock cycle ⇒ nearly 100 clock cycles to multiply two
   32-bit numbers
- Half the bits in the multiplicand register always 0
  - $\Rightarrow$  64-bit adder is wasted
- 0's inserted to right as multiplicand is shifted left
  - ⇒ least significant bits of product never change once formed
- Intuition: instead of shifting multiplicand to left, shift product to right...

# Shift-add Multiplier Version 2



# Shift-add Multiplier Version 2



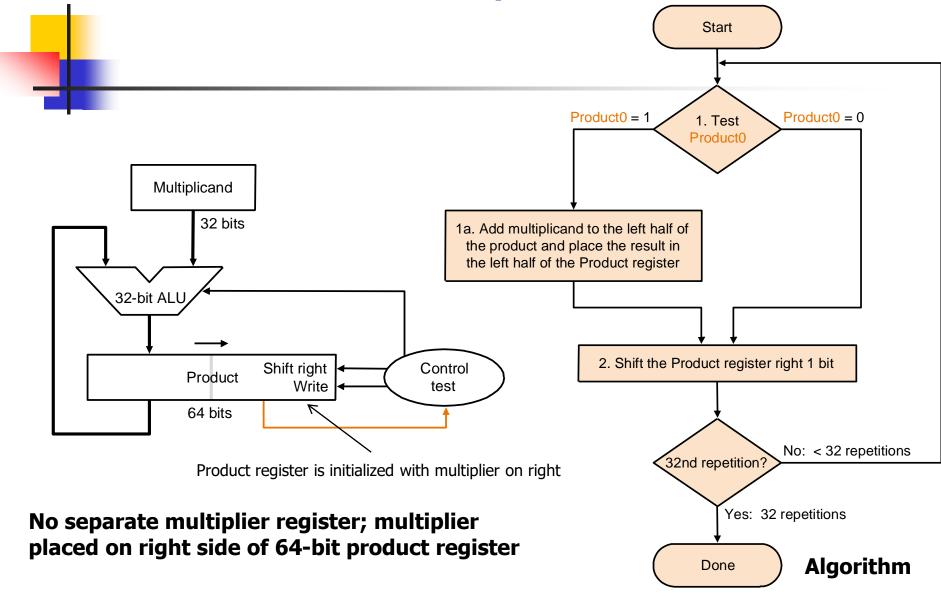
**Algorithm** 

Done

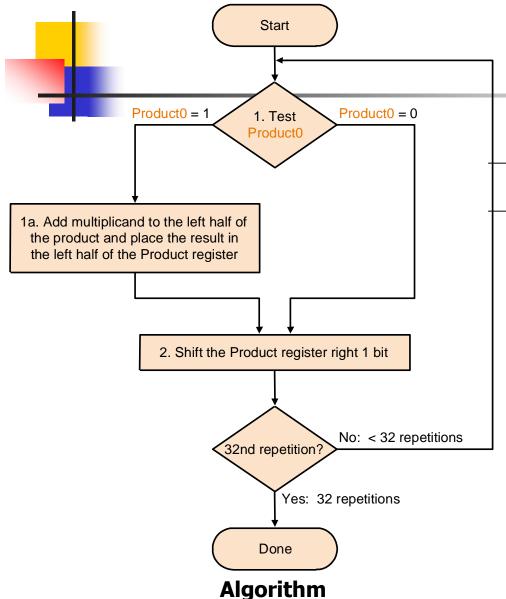
# Observations on Multiply Version 2

- Each step the product register wastes space that exactly matches the current size of the multiplier
- Intuition: combine multiplier register and product register...

# Shift-add Multiplier Version 3



# Shift-add Multiplier Version 3



Example: 0010 \* 0011:

Itera	Step	Multiplicand	Product
-tion		Prancipilicana	Troduct
0	init	0010	0000 0011
	values		
1	1a	0010	0010 0011
	2	0010	0001 0001
2	•••		

# Observations on Multiply Version 3

- 2 steps per bit because multiplier & product combined
- What about signed multiplication?
  - easiest solution is to make both positive and remember whether to negate product when done, i.e., leave out the sign bit, run for 31 steps, then negate if multiplier and multiplicand have opposite signs
- Booth's Algorithm is an elegant way to multiply signed numbers using same hardware – it also often quicker...

### Motivating Booth's algorithm

Example 0010 \* 0110. Traditional:

```
0010

x 0110
0000 shift (0 in multiplier)
0010 add (1 in multiplier)
0010 add (1 in multiplier)
0000 shift (0 in multiplier)
shift (0 in multiplier)
```

Same example. But observe there are two successive 1's in multiplier  $0110 = 2^2 + 2^1 = 2^3 - 2^1$ , so can replace successive 1's by subtract and then add:

```
\begin{array}{c}
0010 \\
\underline{0110} \\
0000
\end{array}

shift (0 in multiplier)

\begin{array}{c}
-0010 \\
0000
\end{array}

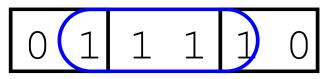
sub (first 1 in multiplier)

\begin{array}{c}
\text{sub} \text{ (middle of string of 1's)} \\
\text{add (previous step had last 1)}
\end{array}
```

## Motivating Booth's Algorithm

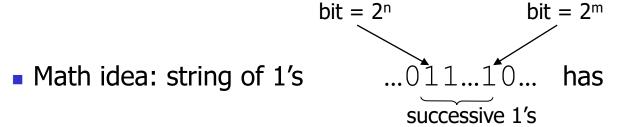
middle of run

end of run



beginning of rul

bit =  $2^{m}$ 



value the sum 
$$2^{n} + 2^{n-1} + ... + 2^{m} = 2^{n+1} - 2^{m}$$

- Replace a string of 1s in multiplier with an initial subtract when we first see a one and then later add after the last one
  - What if the string of 1's started from the left of the (2's complement) number, e.g., 11110001 – would the formula above have to be modified?!

## Booth from Multiply Version 3

Modify Step 1 of the algorithm Multiply Version 3 to consider 2 bits of the multiplier: the current bit and the bit to the right (i.e., the current bit of the previous step). Instead of two outcomes, now there are four:

<u>Case</u>	Current Bit	Bit to the Ri	<u>ght</u> <u>Ex</u> r	<u>planation</u>	<u>Example</u>	<u>Ор</u>
1a	0	0	Middle of	run of 0s	0 <u>00</u> 1111000	none
1b	0	1	End of	run of 1s	00 <u>01</u> 111000	add
1c	1	0	Begins	run of 1s	000111 <u>10</u> 00	sub
1d	1	1	Middle of	run of 1s	00011 <u>11</u> 000	none

- Modify Step 2 of Multiply Version 3 to sign extend when the product is shifted right (arithmetic right shift, rather than logical right shift) because the product is a signed number
- Now draw the flowchart for Booth's algorithm!
- Multiply Version 3 and Booth share the same hardware, except Booth requires one extra flipflop to remember the bit to the right of the current bit in the product register – which is the bit pushed out by the preceding right shift

# Booth Example (2 x 7)

Operation	Multiplicand	Product	next?
0. initial value	e 0010	0000 0111 0	10 -> sub $P = P - M$
1c.	0010	1110 0111 0	shift P (sign ext)
2.	0010	1111 0011 1	11 -> nop
1d.	0010	1111 0011 1	shift P (sign ext)
2.	0010	1111 1001 1	11 -> nop
1d.	0010	1111 1001 1	shift P (sign ext)
2.	0010	1111 1100 1	01 $\rightarrow$ add P = P + M
1b.	0010	0001 1100 1	shift P (sign ext)
2.	0010	0000 1110 0	done

# Booth Algorithm (2 \* -3)

Operation	Multiplicand	Product	next?
O.initial value	0010	0000 1101 0	10 -> sub $P = P - M$
1c.	0010	1110 1101 0	shift P (sign ext)
2.	0010	1111 0110 1	01 -> add $P = P + M$
1b.	0010	0001 0110 1	shift P (sign ext)
2.	0010	0000 1011 0	$10 \rightarrow \text{sub P} = P - M$
1c.	0010	1110 1011 0	shift P
2.	0010	1111 0101 1	11 -> nop
1d.	0010	1111 0101 1	shift P
2.	0010	1111 1010 1	done

# Verifying Booth's Algorithm

- multiplier  $a = a_{31} a_{32}... a_0$ , multiplicand = b
- a<sub>i</sub> a<sub>i-1</sub> Operation
   0 0 nop
   0 1 add b
   1 0 sub b
   1 1 nop
- I.e., if  $a_{i-1} a_i = \begin{cases} 0, & \text{nop} \\ +1, & \text{add b} \\ -1, & \text{sub b} \end{cases}$ 
  - Therefore, Booth computes sum:

$$(a_{-1} - a_0)$$
 \* b \* 2<sup>0</sup>  
+  $(a_0 - a_1)$  \* b \* 2<sup>1</sup>  
+  $(a_1 - a_2)$  \* b \* 2<sup>2</sup>

• • •

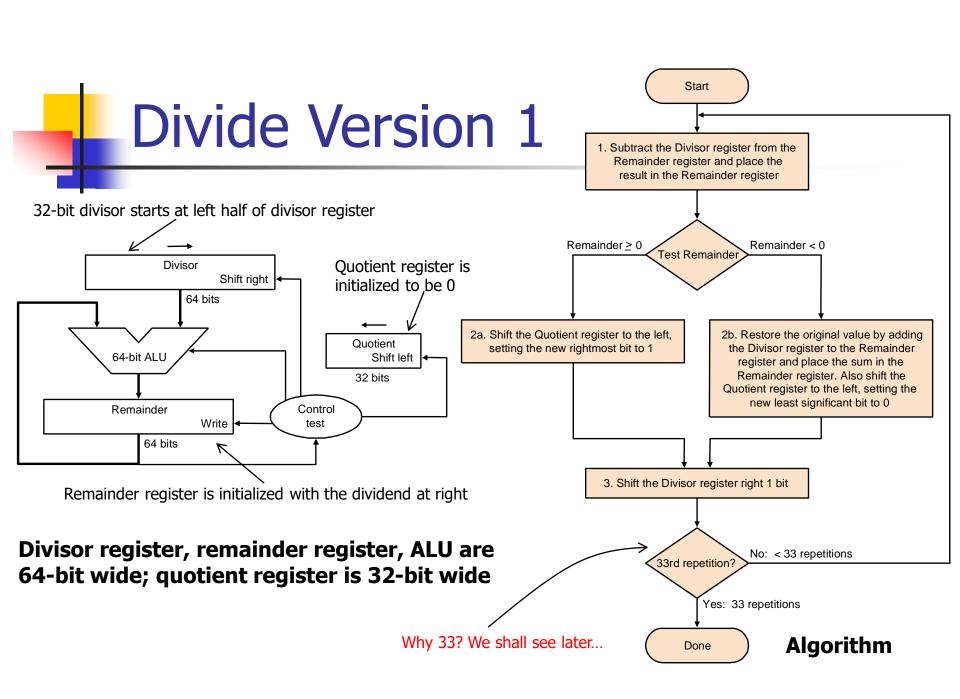
#### **MIPS Notes**

- MIPS provides two 32-bit registers Hi and Lo to hold a 64-bit product
- mult, multu (unsigned) put the product of two 32-bit register operands into Hi and Lo: overflow is ignored by MIPS but can be detected by programmer by examining contents of Hi
- mflo, mfhi moves content of Hi or Lo to a general-purpose register
- Pseudo-instructions mul (without overflow), mulo (with overflow), mulou (unsigned with overflow) take three 32-bit register operands, putting the product of two registers into the third

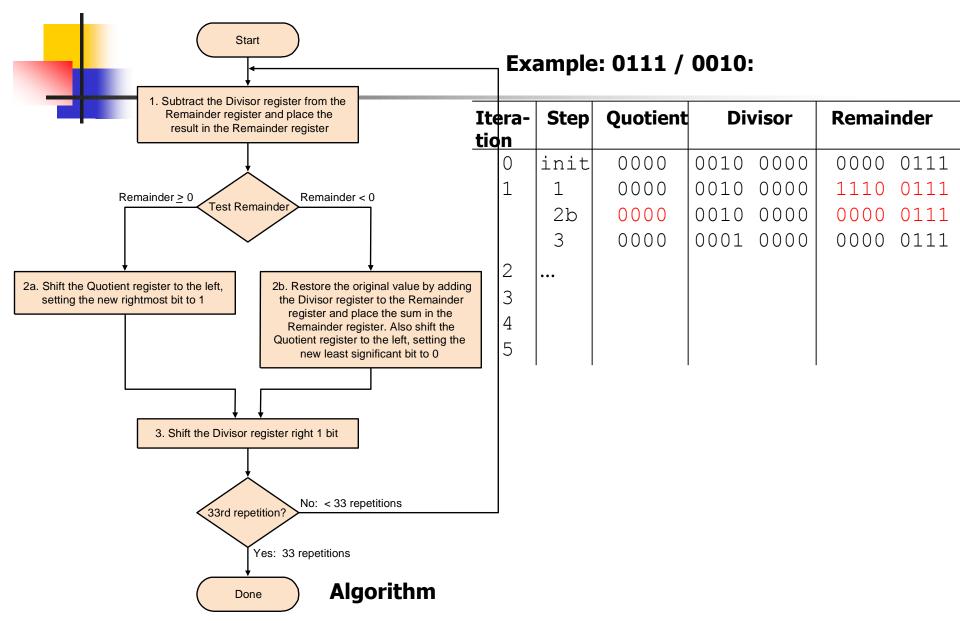
# Divide

```
1001 Quotient
Divisor 1000 1001010 Dividend
-1000
10
101
1010
-1000
10 Remainder
```

- Junior school method: see how big a multiple of the divisor can be subtracted, creating quotient digit at each step
- Binary makes it easy ⇒ first, try 1 \* divisor; if too big, 0 \* divisor
- Dividend = (Quotient \* Divisor) + Remainder
- 3 versions of divide hardware & algorithm:

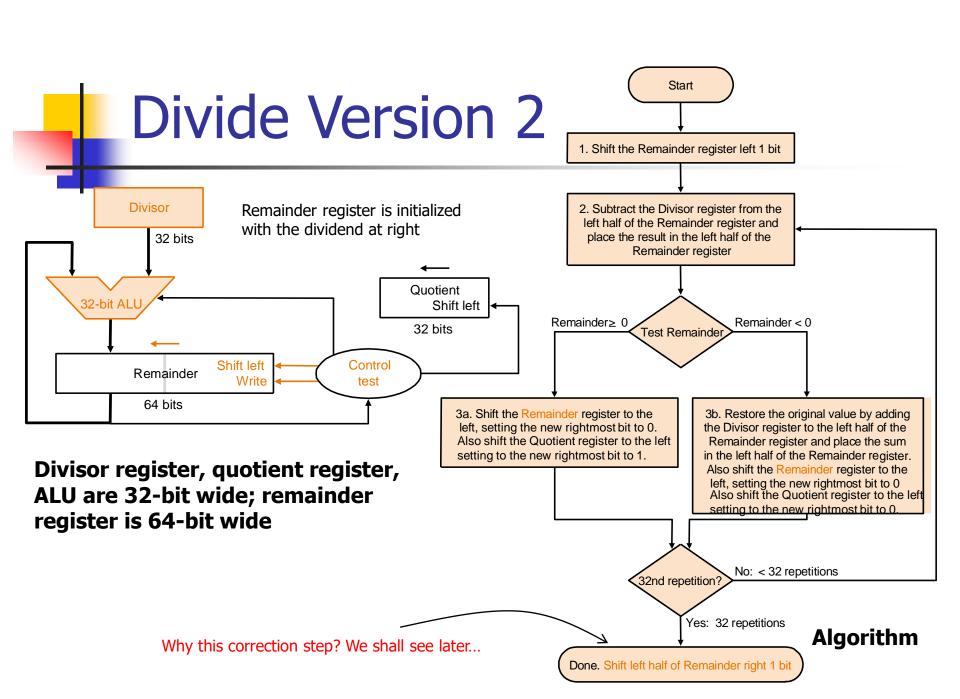


#### Divide Version 1



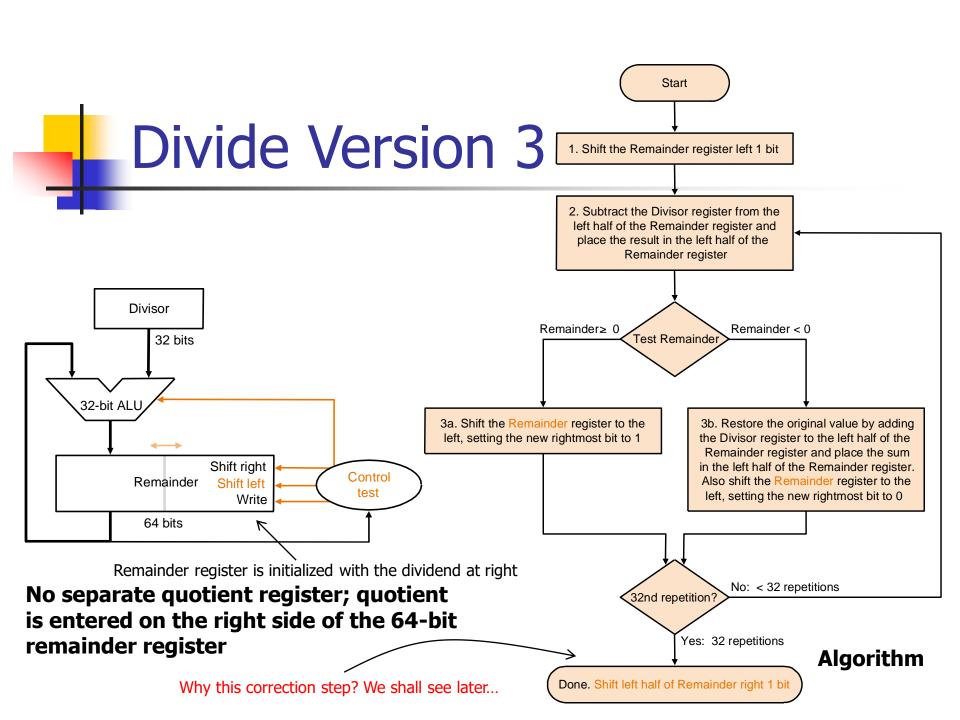
### Observations on Divide Version 1

- Half the bits in divisor always 0
  - $\Rightarrow$  1/2 of 64-bit adder is wasted
  - $\Rightarrow$  1/2 of divisor register is wasted
- Intuition: instead of shifting divisor to right, shift remainder to left...
- Step 1 cannot produce a 1 in quotient bit as all bits corresponding to the divisor in the remainder register are 0 (remember all operands are 32-bit)
- Intuition: switch order to shift first and then subtract can save 1 iteration...

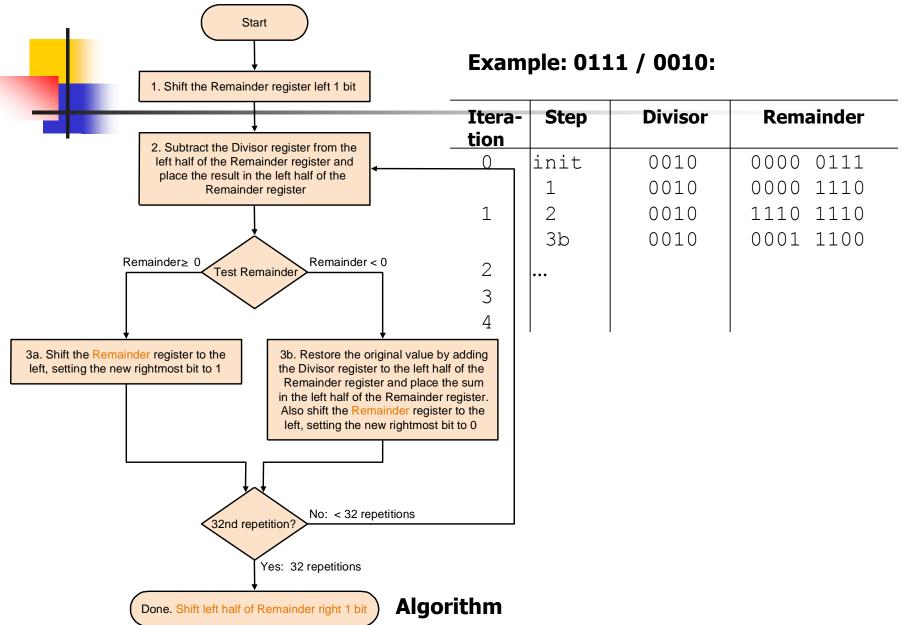


# Observations on Divide Version 2

- Each step the remainder register wastes space that exactly matches the current size of the quotient
- Intuition: combine quotient register and remainder register...



#### Divide Version 3



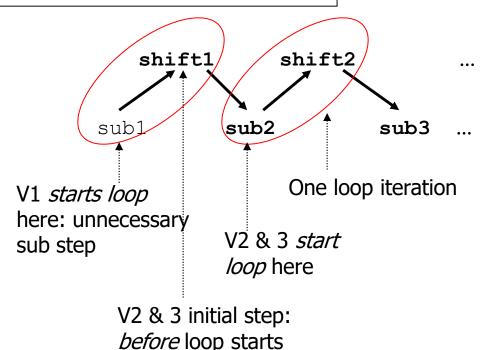
#### Number of Iterations

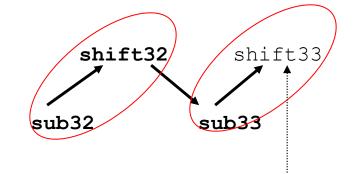
- Why the extra iteration in Version 1?
- Why the final correction step in Versions 2 & 3?

Ovals represent loop iterations

Shift: see the version descriptions for which registers are shifted

Main insight – sub(i+1) must actually follow shifti of the divisor (or remainder, depending on version) and the resulting bit in the quotient appears on shift(i+1)





Critical situation! Only the quotient shift is necessary as it corresponds to the outcome of the previous sub.

So V1 is ok even though the last divisor shift is redundant, as final divisor is ignored any way; V2 & 3 must repair remainder as it has shifted left one time too many

## Observations on Divide Version 3

- Same hardware as Multiply Version 3
- Signed divide:
  - make both divisor and dividend positive and perform division
  - negate the quotient if divisor and dividend were of opposite signs
  - make the sign of the remainder match that of the dividend
  - this ensures always
    - dividend = (quotient \* divisor) + remainder
    - -quotient (x/y) = quotient (-x/y) (e.g. 7 = 3\*2 + 1 & -7 = -3\*2 1)

#### **MIPS Notes**

- div (signed), divu (unsigned), with two 32-bit register operands, divide the contents of the operands and put remainder in Hi register and quotient in Lo; overflow is ignored in both cases
- pseudo-instructions div (signed with overflow), divu (unsigned without overflow) with three 32-bit register operands puts quotients of two registers into third

### Floating Point

- We need a way to represent
  - numbers with fractions, e.g., 3.1416
  - very small numbers (in absolute value), e.g., .00000000023
  - very large numbers (in absolute value) , e.g., −3.15576 \* 10⁴6
- Representation:
  - scientific: sign, exponent, significand form: binary point
     (-1)<sup>sign</sup> \* significand \* 2<sup>exponent</sup>. E.g., -101.001101 \* 2<sup>111001</sup>
  - more bits for significand gives more accuracy
  - more bits for exponent increases range
  - if  $1 \le \text{ significand} < 10_{two} (=2_{ten})$  then number is *normalized*, **except** for number 0 which is normalized to significand 0
    - E.g.,  $-101.001101 * 2^{111001} = -1.01001101 * 2^{111011}$  (normalized)

# **IEEE** 754 Floating-point Standard

- IEEE 754 floating point standard:
  - single precision: one word

31	bits 30 to 23	bits 22 to 0
sign	8-bit exponent	23-bit significand

double precision: two words

31	bits 30 to 20	bits 19 to 0
sign	11-bit exponent	upper 20 bits of 52-bit significand

bits 31 to 0
lower 32 bits of 52-bit significand

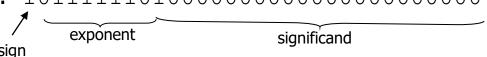
# JEEE 754 Floating-point Standard

- Sign bit is 0 for positive numbers, 1 for negative numbers
- Number is assumed normalized and leading 1 bit of significand left of binary point (for non-zero numbers) is assumed and not shown
  - e.g., significand 1.1001... is represented as 1001...,
  - exception is number 0 which is represented as all 0s (see next slide)
  - for other numbers: value =  $(-1)^{sign} * (1 + significand) * 2^{exponent value}$
- Exponent is biased to make sorting easier
  - all 0s is smallest exponent, all 1s is largest
  - bias of 127 for single precision and 1023 for double precision
  - therefore, for non-0 numbers:
     value = (−1)<sup>sign</sup> \* (1 + significand) \* 2<sup>(exponent bias)</sup>

# **IEEE** 754 Floating-point Standard

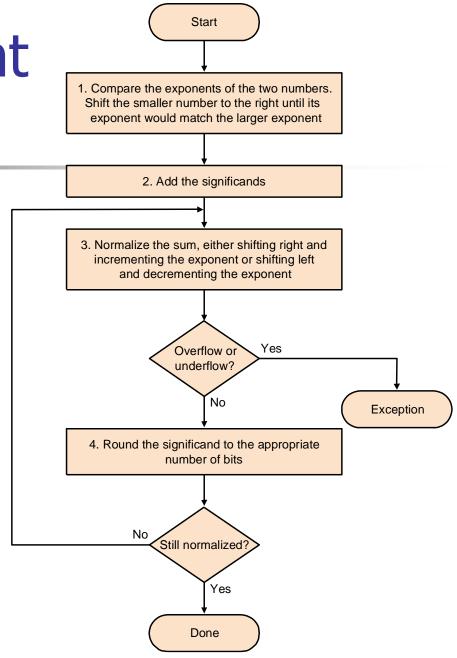
- Special treatment of 0:
  - if exponent is all 0 and significand is all 0, then the value is 0 (sign bit may be 0 or 1)
  - if exponent is all 0 and significand is *not* all 0, then the value is  $(-1)^{sign} * (1 + significand) * 2^{-127}$ 
    - therefore, all 0s is taken to be 0 and not 2<sup>-127</sup> (as would be for a non-zero normalized number); similarly, 1 followed by all 0's is taken to be 0 and not 2<sup>-127</sup>
- Example: Represent -0.75<sub>ten</sub> in IEEE 754 single precision
  - decimal:  $-0.75 = -3/4 = -3/2^2$
  - binary:  $-11/100 = -.11 = -1.1 \times 2^{-1}$
  - IEEE single precision floating point exponent = bias + exponent value

$$= 127 + (-1) = 126_{ten} = 011111110_{two}$$



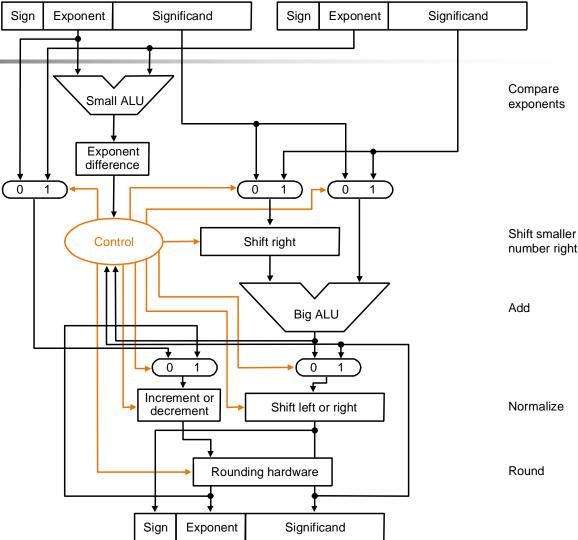
# Floating Point Addition

Algorithm:



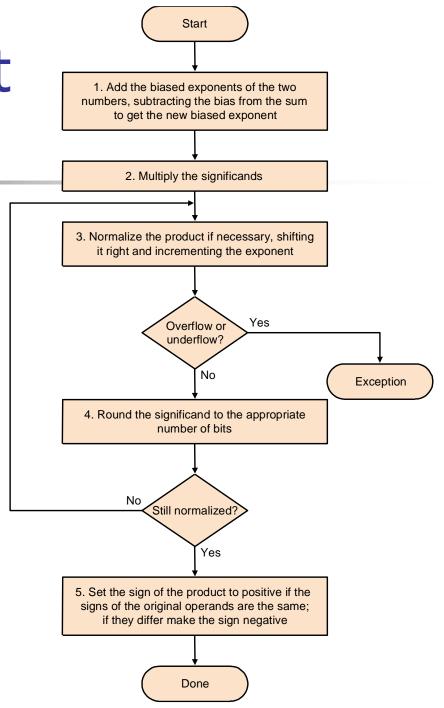
# Floating Point Addition Sign Exponent

Hardware:



# Floating Point Multpication

Algorithm:



### Floating Point Complexities

- In addition to overflow we can have underflow (number too small)
- Accuracy is the problem with both overflow and underflow because we have only a finite number of bits to represent numbers that may actually require arbitrarily many bits
  - limited precision ⇒ rounding ⇒ rounding error
  - IEEE 754 keeps two extra bits, guard and round
  - four rounding modes
  - positive divided by zero yields *infinity*
  - zero divide by zero yields not a number
  - other complexities
- Implementing the standard can be tricky
- Not implementing the standard can be even worse
  - see text for discussion of Pentium bug!

## MIPS Floating Point

- MIPS has a floating point coprocessor (numbered 1, SPIM) with thirty-two 32-bit registers \$f0 - \$f31. Two of these are required to hold doubles. Floating point instructions must use only evennumbered registers (including those operating on single floats). SPIM simulates MIPS floating point.
- Floating point arithmetic. add.s (single addition), add.d (double addition), sub.s, sub.d, mul.s, mul.d, div.s, div.d
- Floating point comparison: c.x.s (single), c.x.d (double), where x may be eq, neq, lt, le, gt, ge
- Other instructions...

### Summary

- Computer arithmetic is constrained by limited precision
- Bit patterns have no inherent meaning but standards do exist:
  - two's complement
  - IEEE 754 floating point
- Computer instructions determine meaning of the bit patterns.
- Performance and accuracy are important so there are many complexities in real machines (i.e., algorithms and implementation)
- Read Computer Arithmetic Algorithms by I. Koren
  - it is easy-to-read and shows new algorithms for arithmetic
  - there will be assignment and projects based on Koren's material