



Computer Organization and Architecture

Instructor: Dr. Rushdi Abu
Zneit

Slide Sources: Patterson &
Hennessy

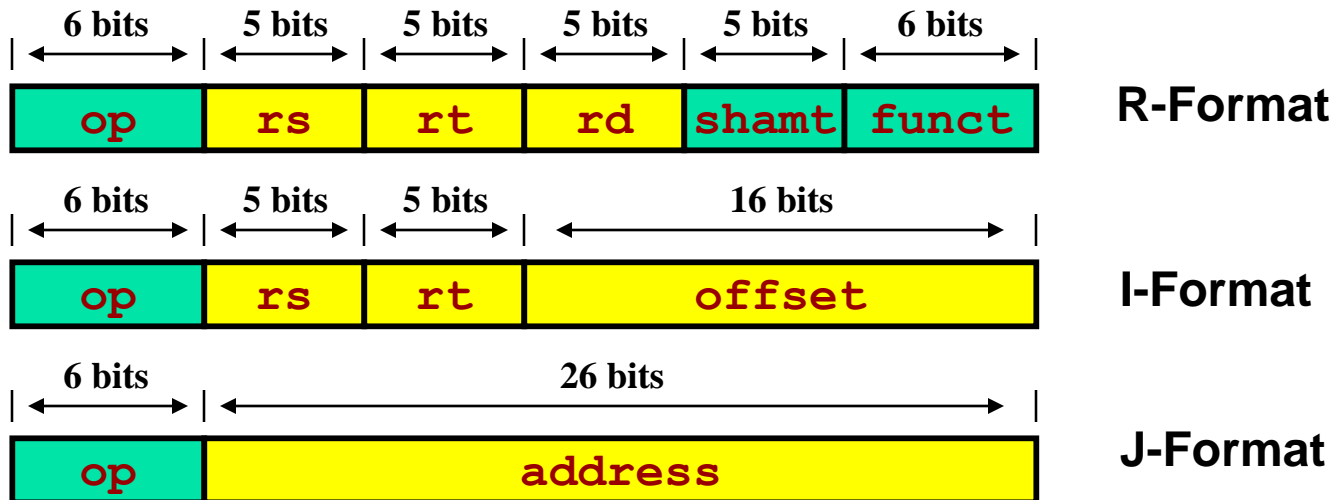
COD Ch. 5



The Processor: Datapath and Control

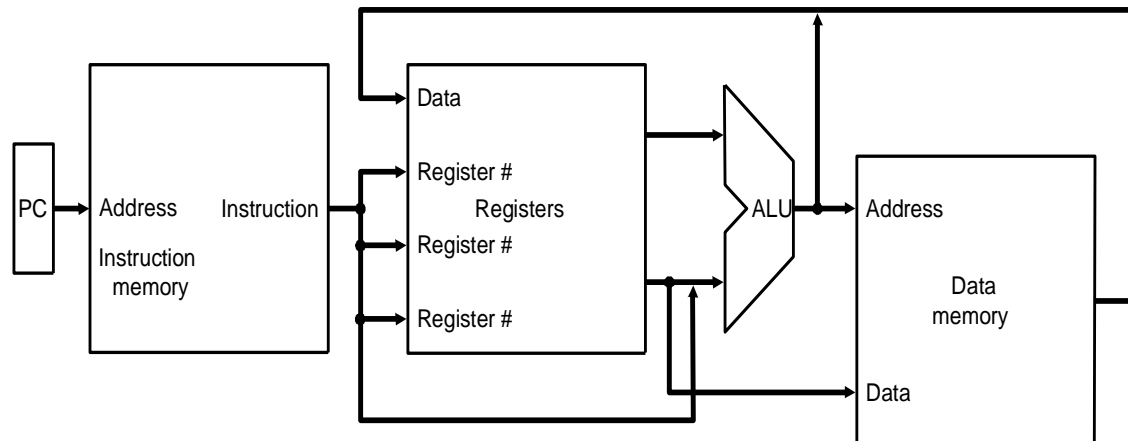
Implementing MIPS

- We're ready to look at an implementation of the MIPS instruction set
- Simplified to contain only
 - arithmetic-logic instructions: `add`, `sub`, `and`, `or`, `slt`
 - memory-reference instructions: `lw`, `sw`
 - control-flow instructions: `beq`, `j`



Implementing MIPS: the Fetch/Execute Cycle

- High-level abstract view of *fetch/execute* implementation
 - use the program counter (PC) to read instruction address
 - *fetch* the instruction from memory and increment PC
 - use fields of the instruction to select registers to read
 - *execute* depending on the instruction
 - repeat...





Overview: Processor Implementation Styles

- **Single Cycle**
 - perform each instruction in 1 clock cycle
 - clock cycle must be long enough for slowest instruction; therefore,
 - disadvantage: only as fast as slowest instruction
- **Multi-Cycle**
 - break fetch/execute cycle into multiple steps
 - perform 1 step in each clock cycle
 - advantage: each instruction uses only as many cycles as it needs
- **Pipelined**
 - execute each instruction in multiple steps
 - perform 1 step / instruction in each clock cycle
 - process multiple instructions in parallel – assembly line




Functional Elements



- Two types of functional elements in the hardware:
 - elements that *operate on data* (called *combinational elements*)
 - elements that *contain data* (called *state or sequential elements*)

Combinational Elements

- Works as an *input* \Rightarrow *output function*, e.g., ALU
- Combinational logic *reads input data from one register and writes output data to another, or same, register*
 - *read/write happens in a single cycle* – combinational element *cannot store data* from one cycle to a future one



Combinational logic hardware units






State Elements

- State elements contain *data* in internal storage, e.g., *registers* and *memory*
- All state elements together *define* the *state of the machine*
 - *What does this mean? Think of shutting down and starting up again...*
- *Flipflops* and *latches* are 1-bit state elements, equivalently, they are *1-bit memories*
- The *output(s)* of a flipflop or latch *always* depends on the bit value stored, i.e., its state, and can be called *1/0* or *high/low* or *true/false*
- The *input* to a flipflop or latch can change its state depending on whether it is clocked or not...


Set-Reset (SR-) latch (unclocked)

Think of \overline{S} as \overline{S} , the inverse of *set* (which sets Q to 1), and \overline{R} as \overline{R} , the inverse of *reset*.



See `sr_latch.v` in Verilog Examples

equivalently with nor gates



A set-reset latch made from two cross-coupled *nand* gates is a basic memory unit.

When both \overline{S} and \overline{R} are 1, then either *one of the following two states is stable*:

- a) $Q = 1$ & $\overline{Q} = 0$
- b) $Q = 0$ & $\overline{Q} = 1$

and the latch will *continue* in the current stable state.


If \overline{S} changes to 0 (while \overline{R} remains at 1), then the latch is forced to the *exactly one* possible stable state (a). If \overline{R} changes to 0 (while \overline{S} remains at 1), the latch is forced to the *exactly one* possible stable state (b).

So, the latch *remembers* which of \overline{S} or \overline{R} was last 0 *during* the time they are both 1.

When both \overline{S} and \overline{R} are 0 the *exactly one* stable state is $Q = \overline{Q} = 1$. However, if after that both \overline{S} and \overline{R} return to 1, the latch must then *jump non-deterministically* to one of stable states (a) or (b), which is undesirable behavior.

Synchronous Logic: Clocked Latches and Flipflops


- Clocks are used in *synchronous* logic to determine *when* a state element is to be updated
 - in *level-triggered* clocking methodology either the state changes only when the clock is high or only when it is low (technology-dependent)



- in *edge-triggered* clocking methodology either the *rising edge* or *falling edge* is active (depending on technology) – i.e., states change only on rising edges or only on falling edge
- Latches are level-triggered
- Flipflops are edge-triggered

Clocked SR-latch


- State can change only when clock is *high*
- *Potential problem* : both inputs $S_{\text{bar}} = 0$ & $R_{\text{bar}} = 0$ will cause non-deterministic behavior




See clockedSr_latch.v in Verilog Examples

Clocked D-latch

- State can change only when clock is *high*
- Only *single* data input (compare SR-latch)
- *No problem* with non-deterministic behavior



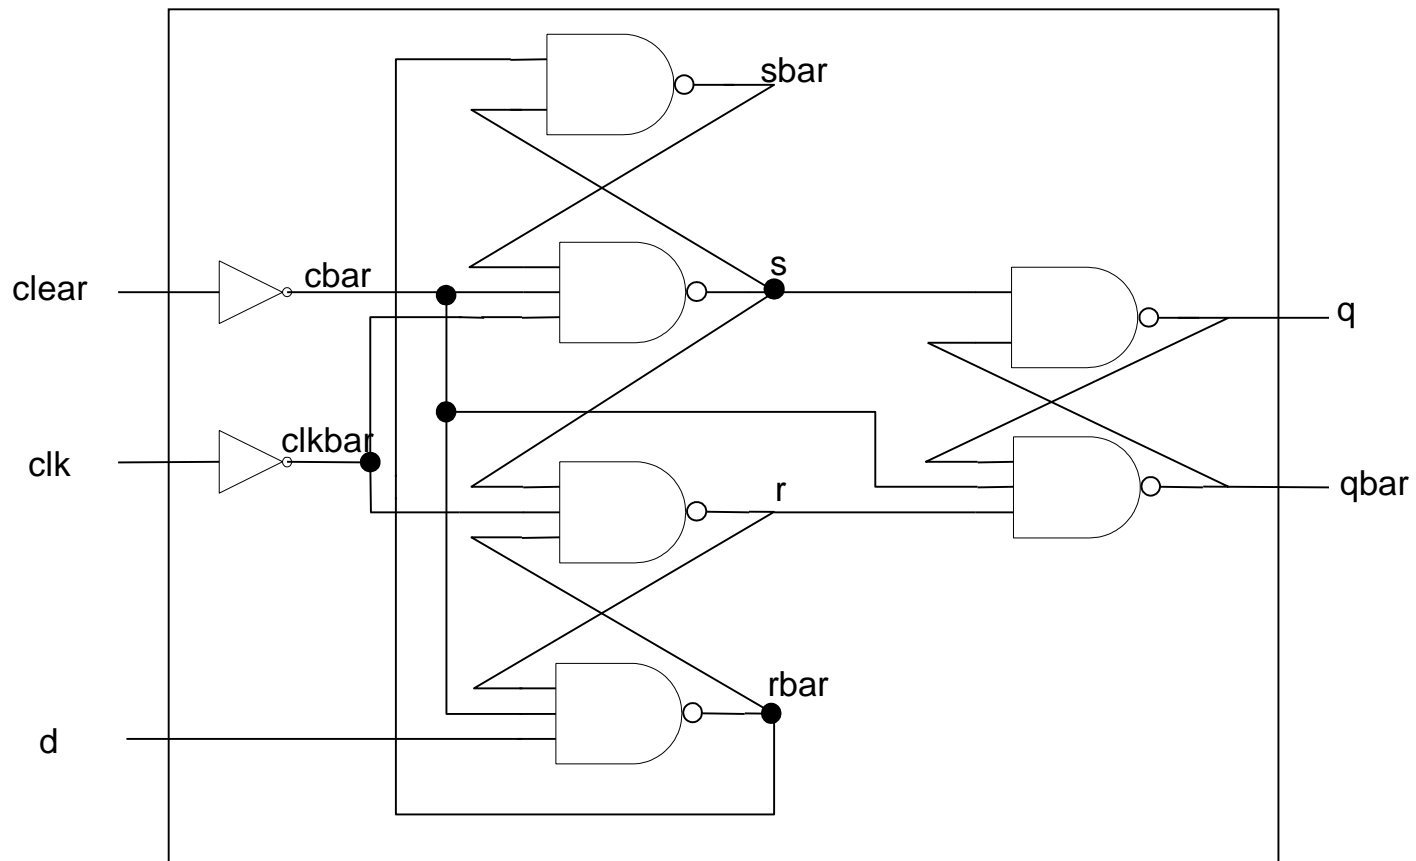
See clockedD_latch.v in Verilog Examples



Timing diagram of D-latch

Clocked D-flipflop

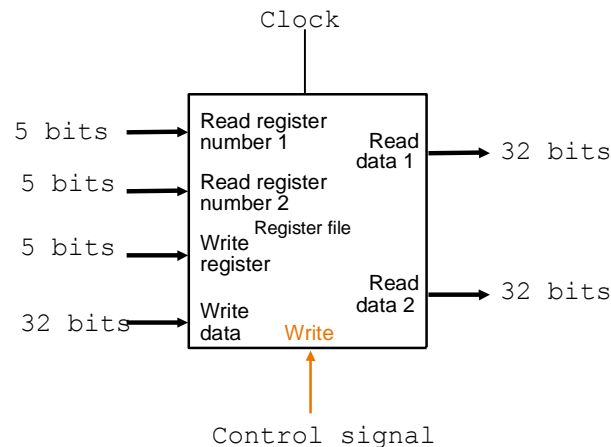
- *Negative edge-triggered*
- Made from *three SR-latches*



See edge_dffGates.v in Verilog Examples

State Elements on the Datapath: Register File

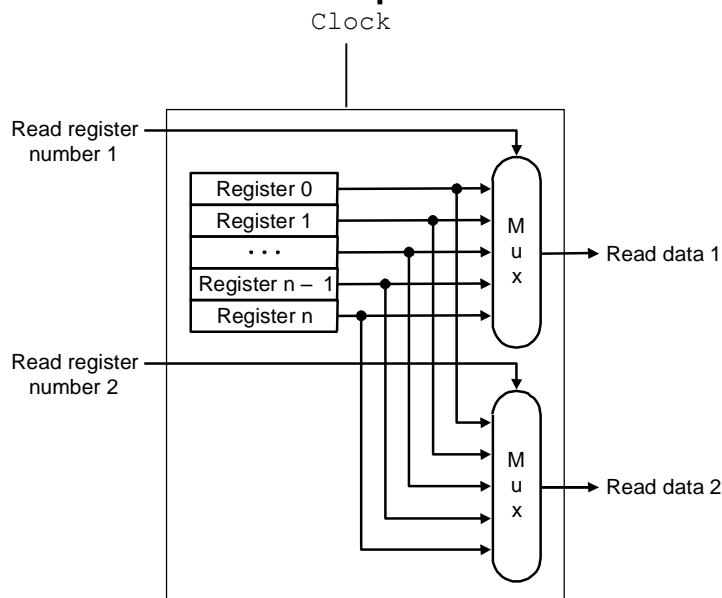
- Registers are implemented with arrays of D-flipflops



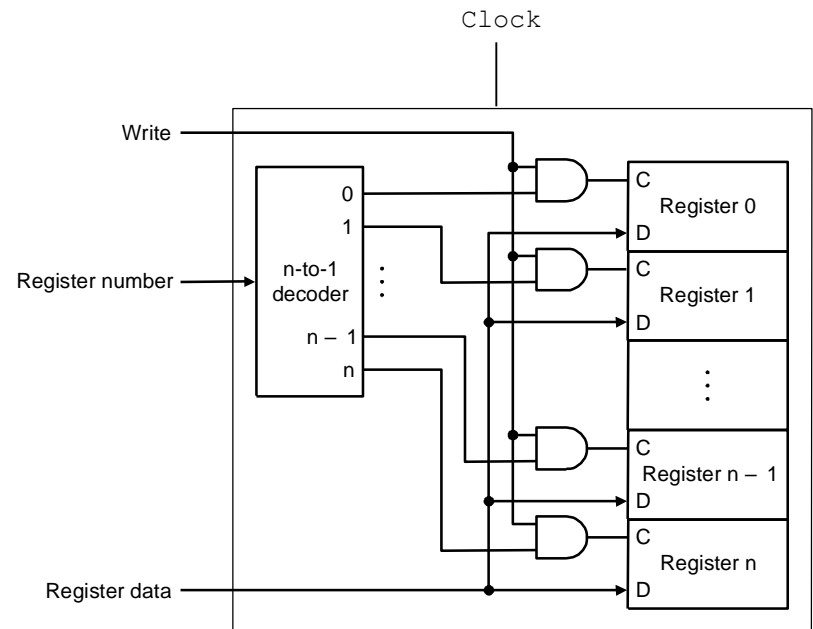
Register file with two read ports and one write port

State Elements on the Datapath: Register File

- Port implementation:



Read ports are implemented with a pair of multiplexors – 5 bit multiplexors for 32 registers



Write port is implemented using a decoder – 5-to-32 decoder for 32 registers. Clock is relevant to write as register state may change only at clock edge



Verilog


- All components that we have discussed – and shall discuss – can be fabricated using Verilog
- Refer to our Verilog slides and examples



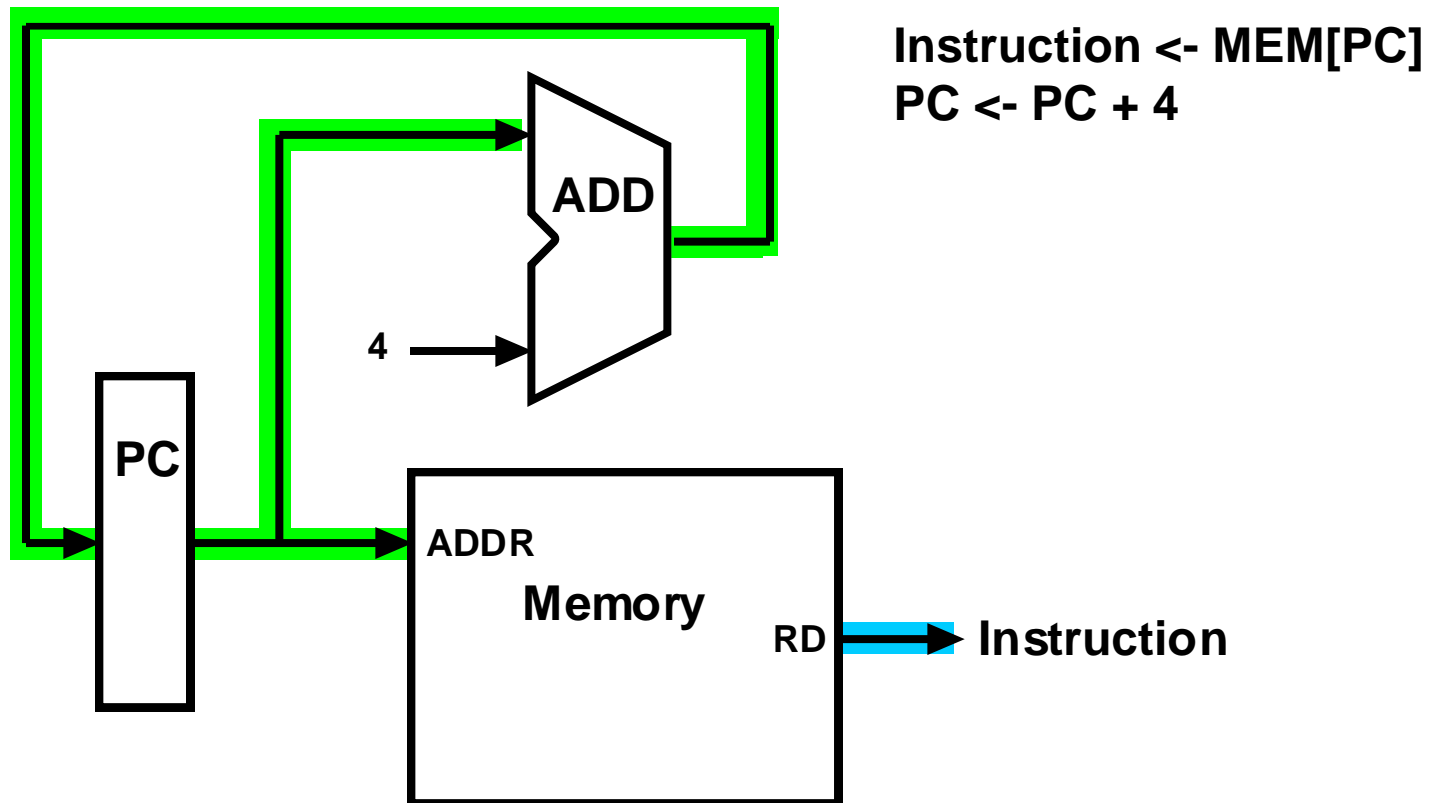
Single-cycle Implementation of MIPS

- Our first implementation of MIPS will use a *single* long clock cycle for every instruction
- Every instruction begins on one up (or, down) clock edge and ends on the next up (or, down) clock edge
- This approach is *not practical* as it is much slower than a *multicycle* implementation where different instruction classes can take different numbers of cycles
 - in a single-cycle implementation every instruction must take the same amount of time as the slowest instruction
 - in a multicycle implementation this problem is avoided by allowing quicker instructions to use fewer cycles
- Even though the single-cycle approach is not practical it is simple and useful to understand first
- *Note* : we shall implement `jump` at the very end

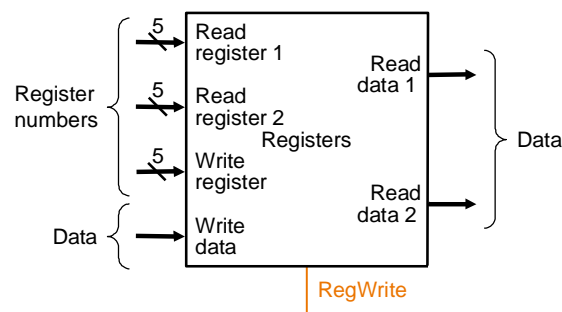
Datapath: Instruction Store/Fetch & PC Increment




Animating the Datapath



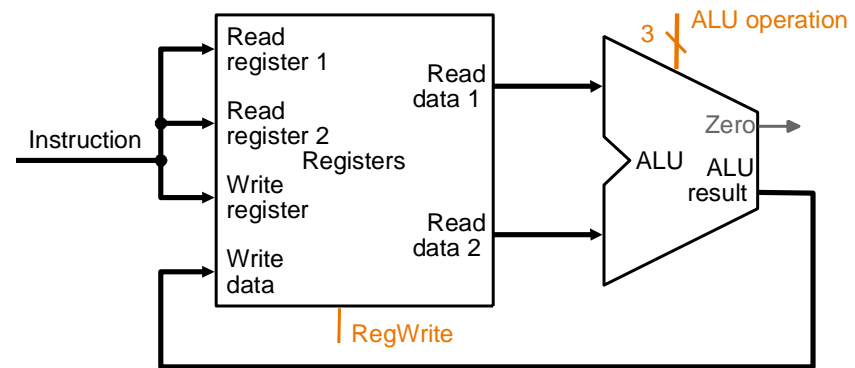
Datapath: R-Type Instruction



a. Registers



b. ALU



Datapath

**Two elements used to implement
R-type instructions**

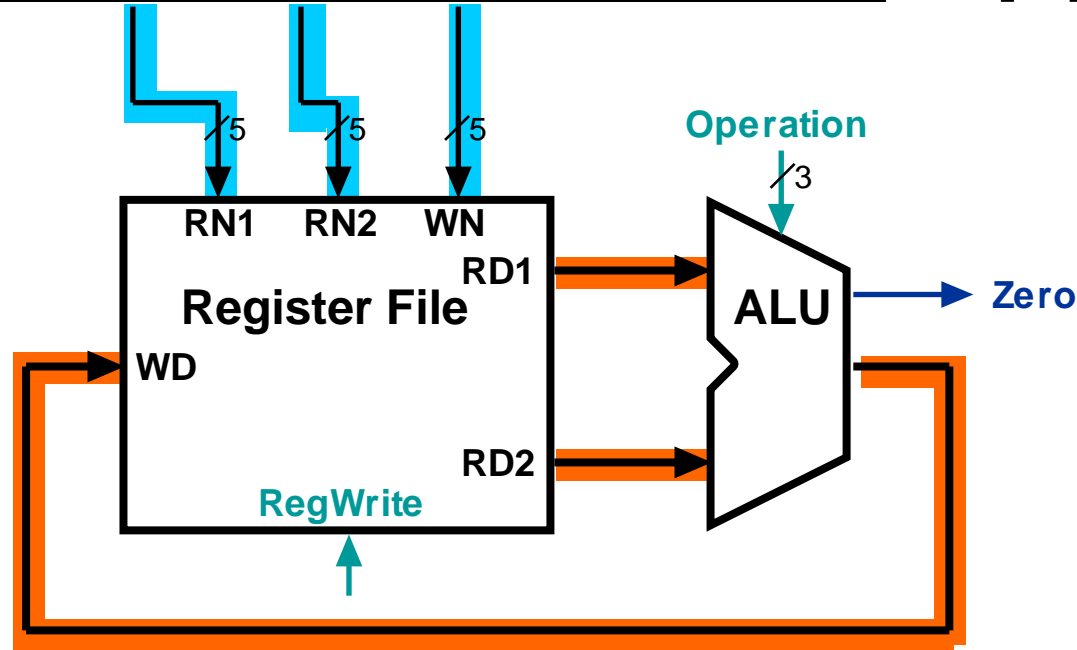
Animating the Datapath

Instruction

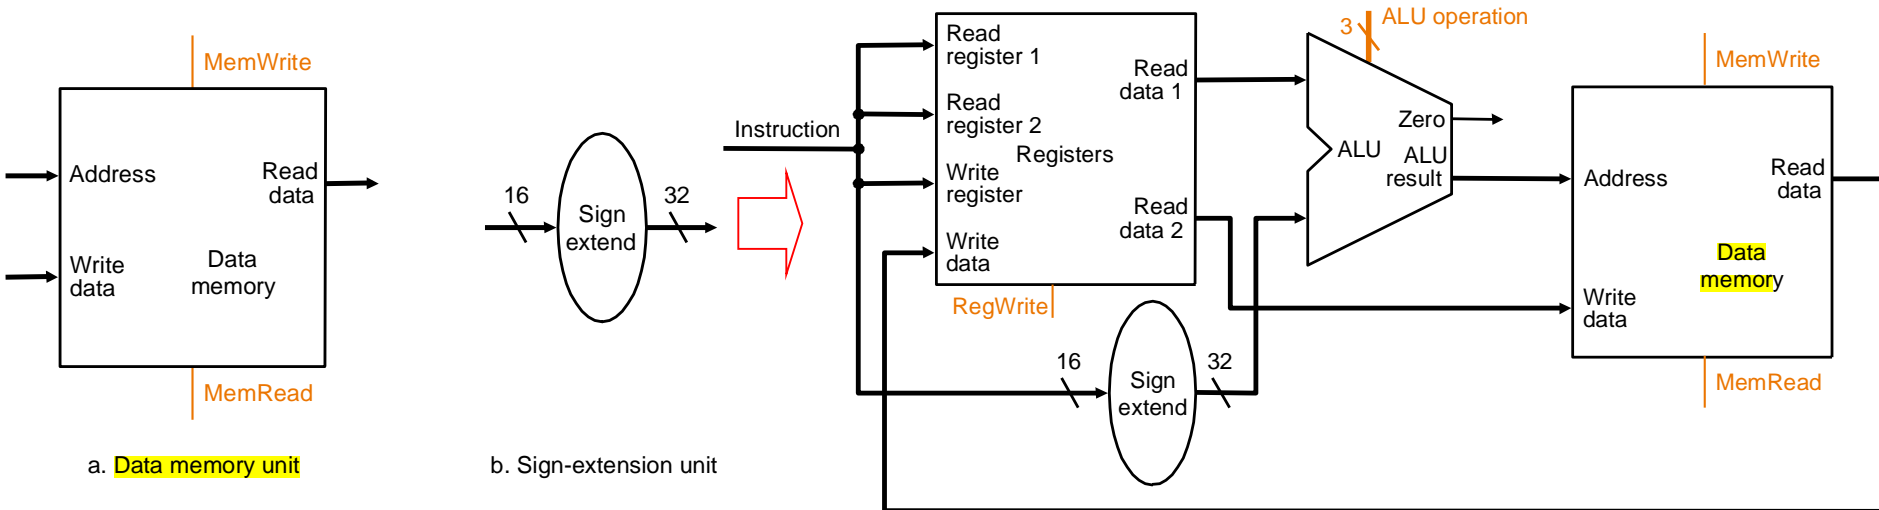


add rd, rs, rt

$R[rd] \leftarrow R[rs] + R[rt];$



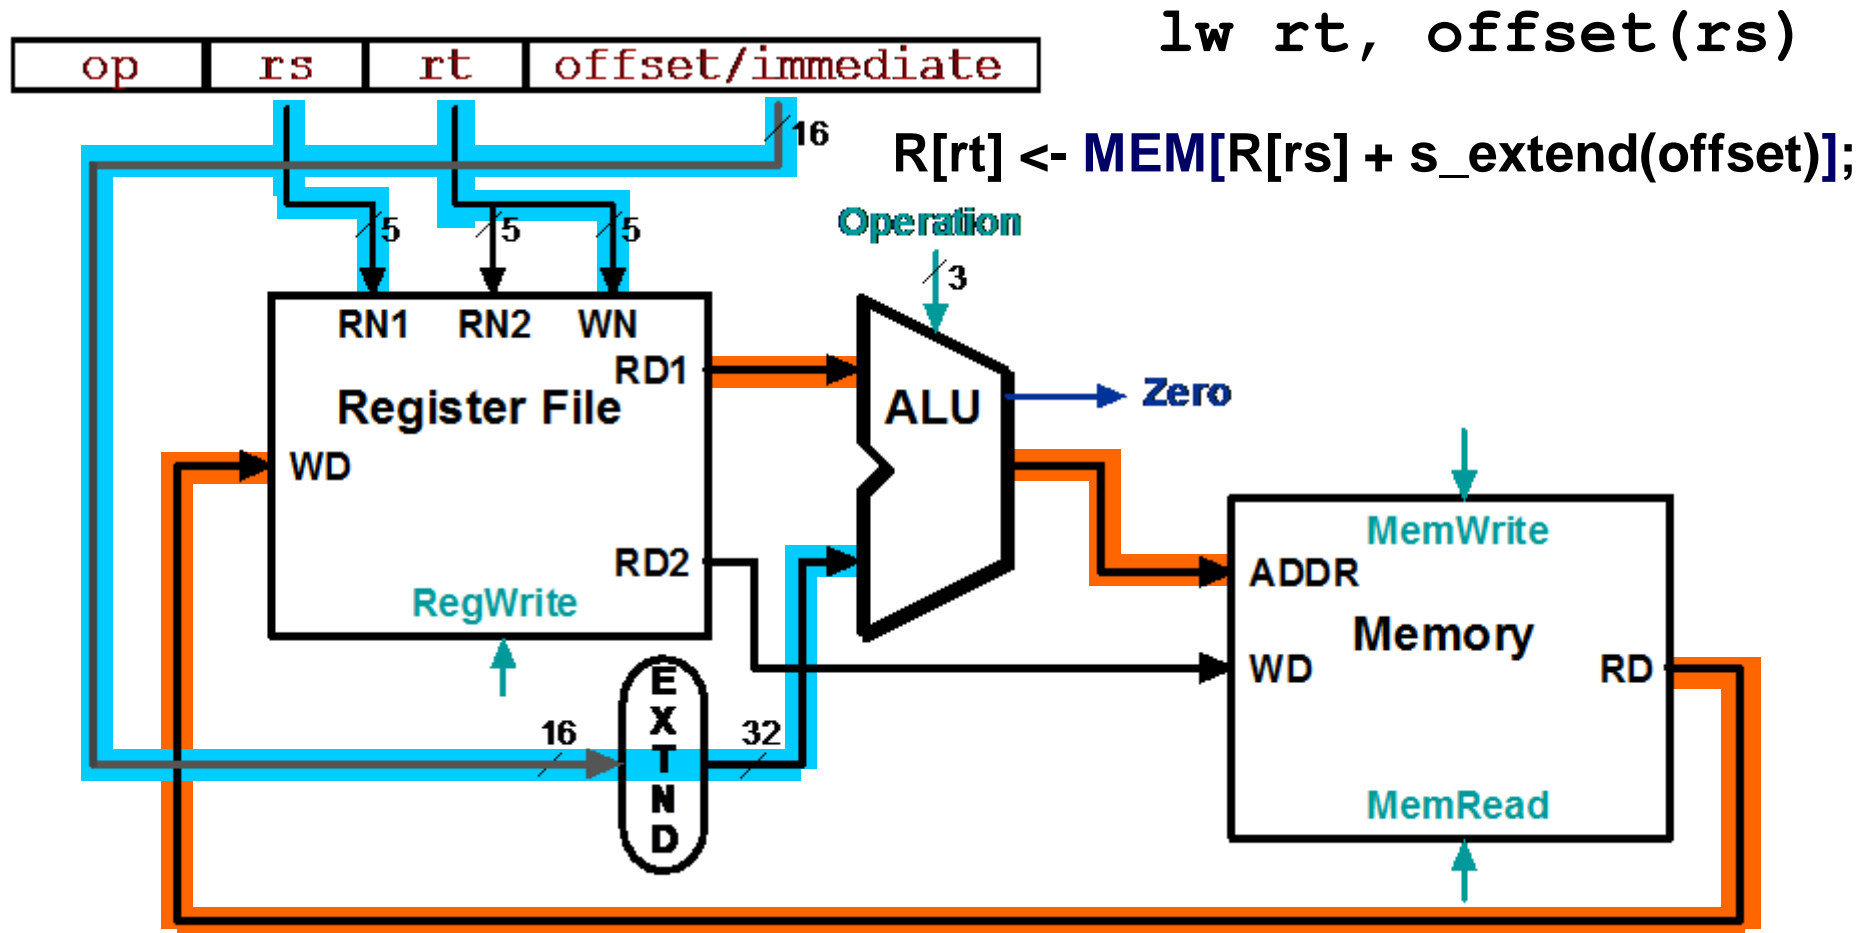
Datapath: Load/Store Instruction



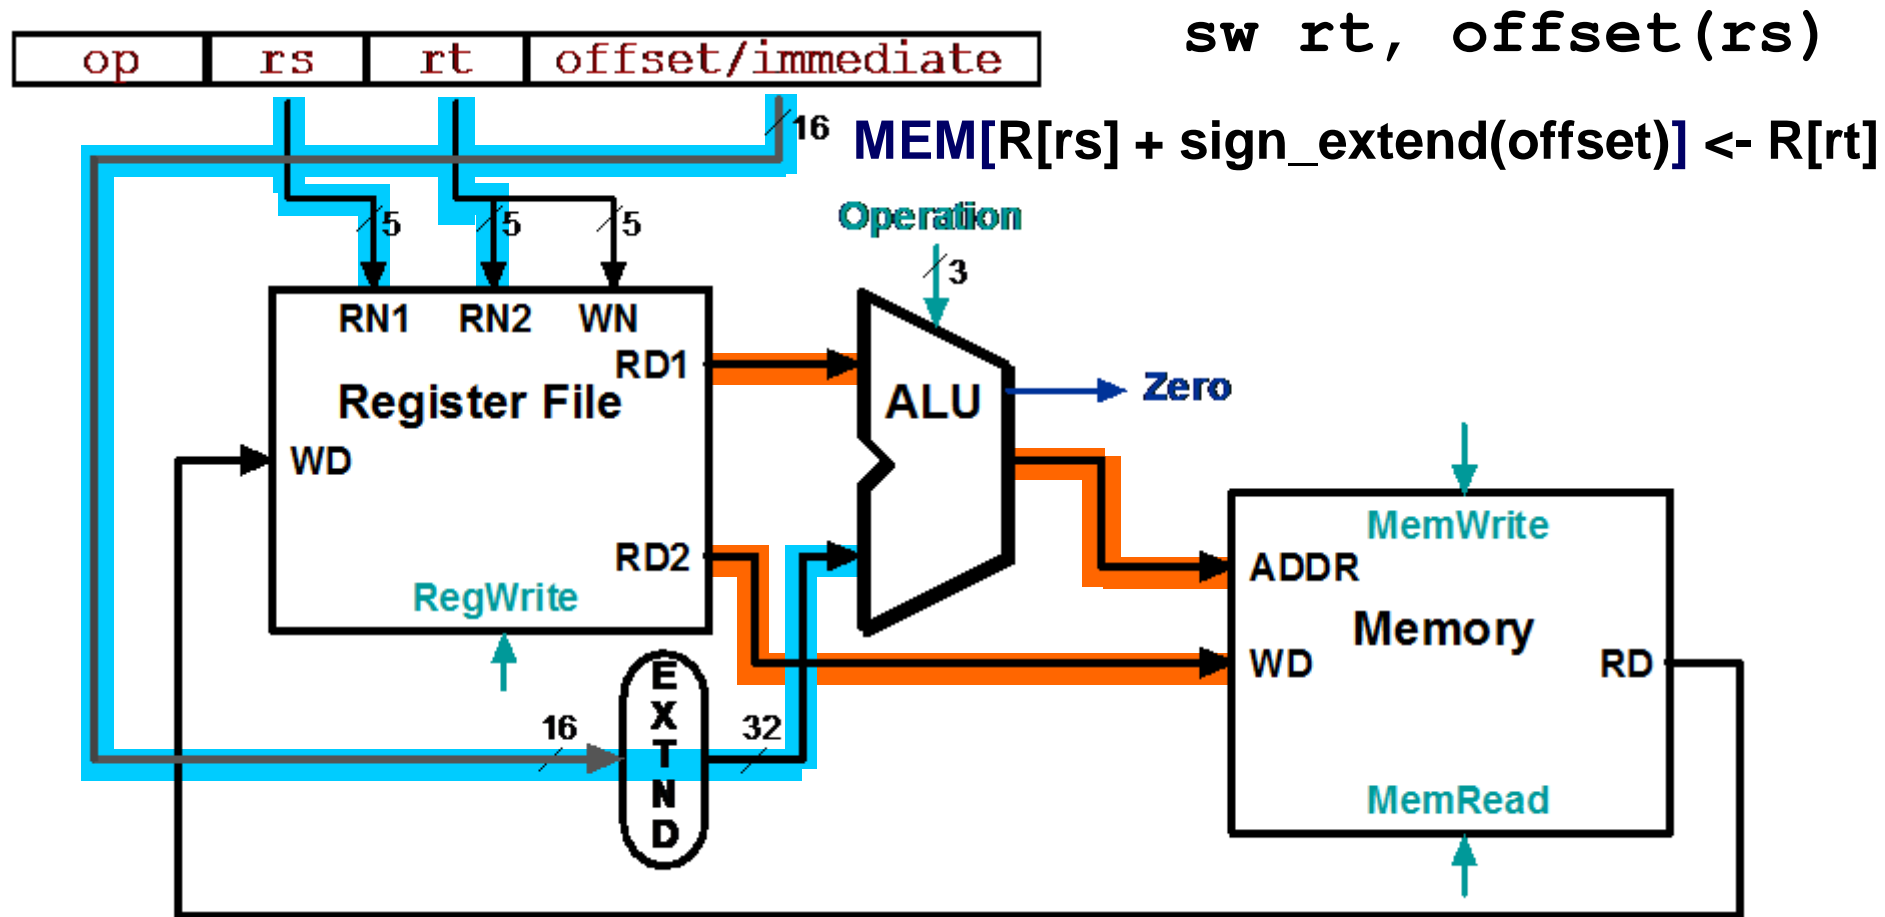
**Two additional elements used
To implement load/stores**

Datapath

Animating the Datapath

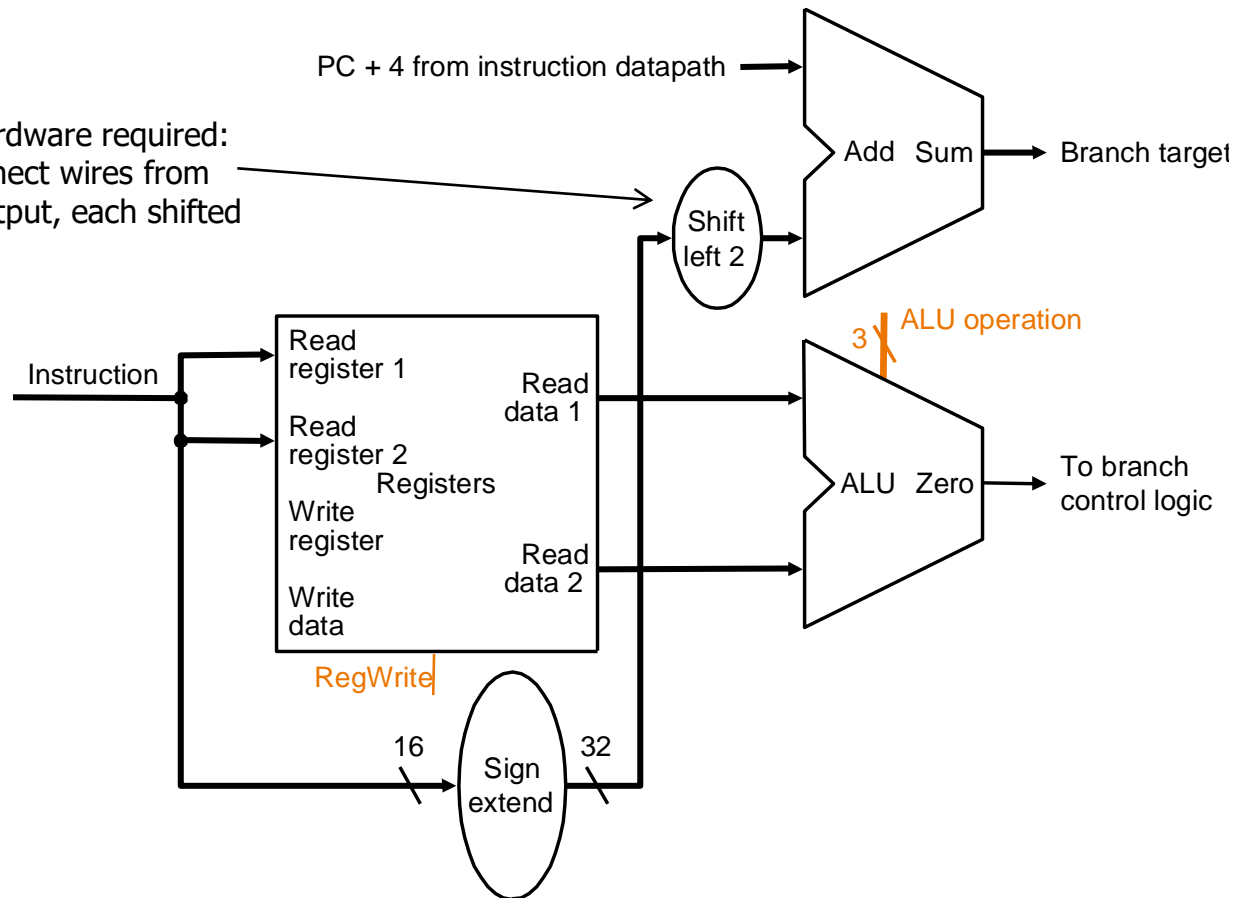


Animating the Datapath



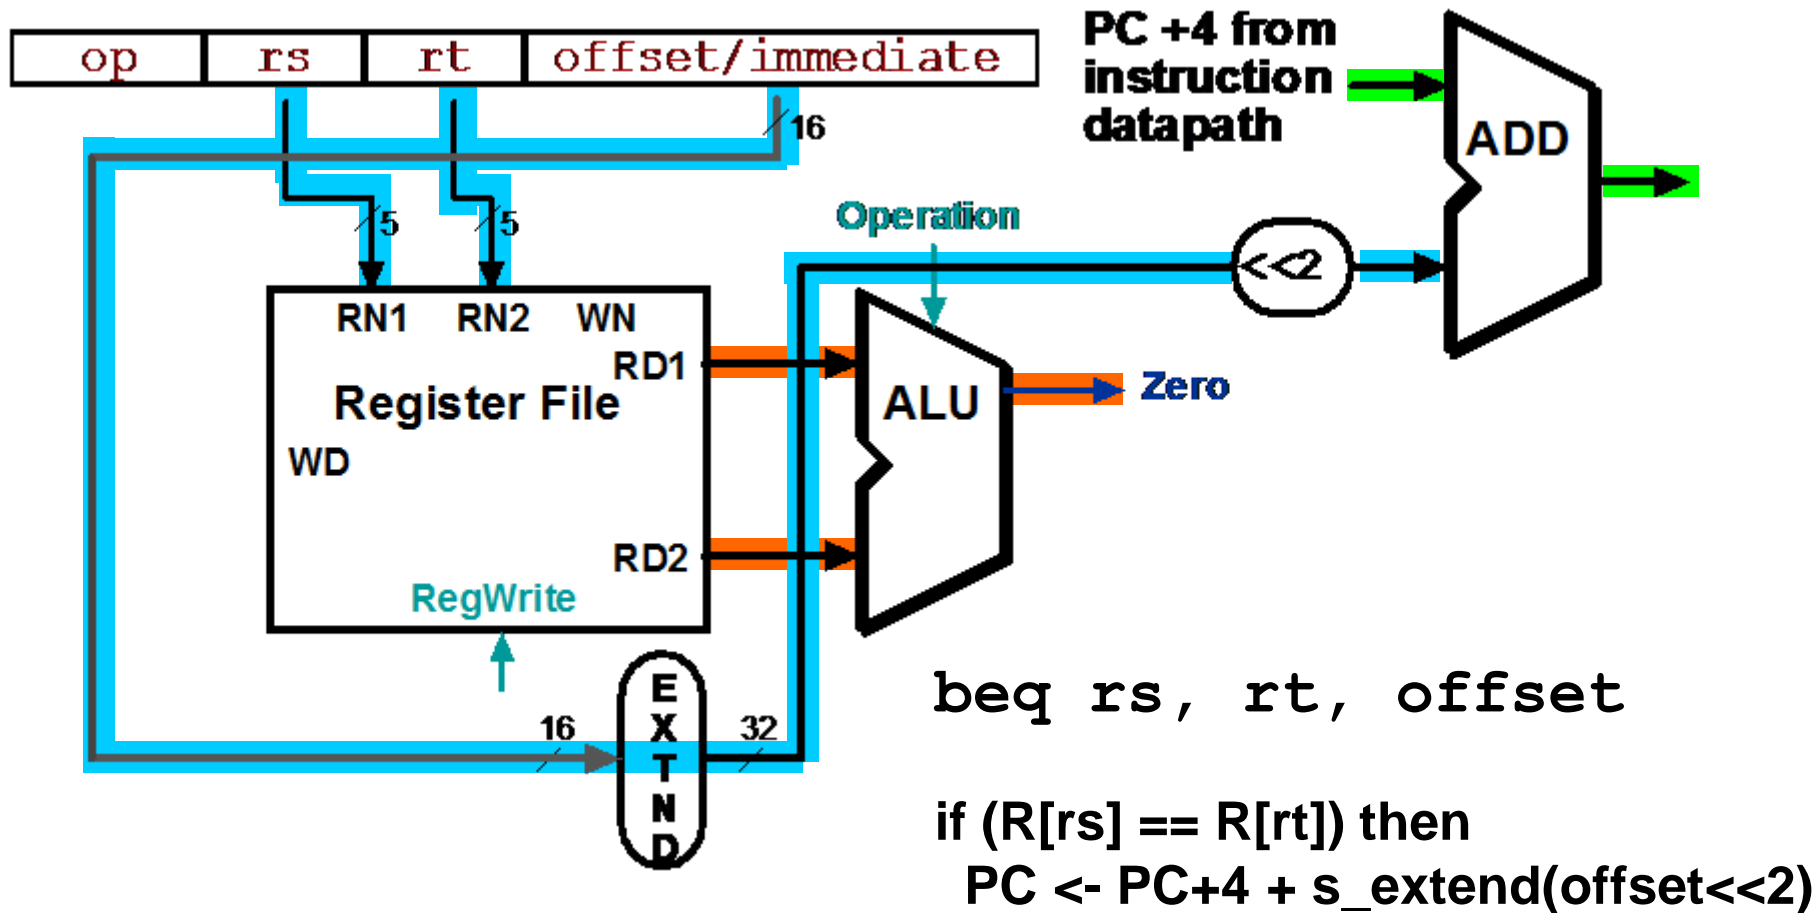
Datapath: Branch Instruction

No shift hardware required:
simply connect wires from
input to output, each shifted
left 2 bits



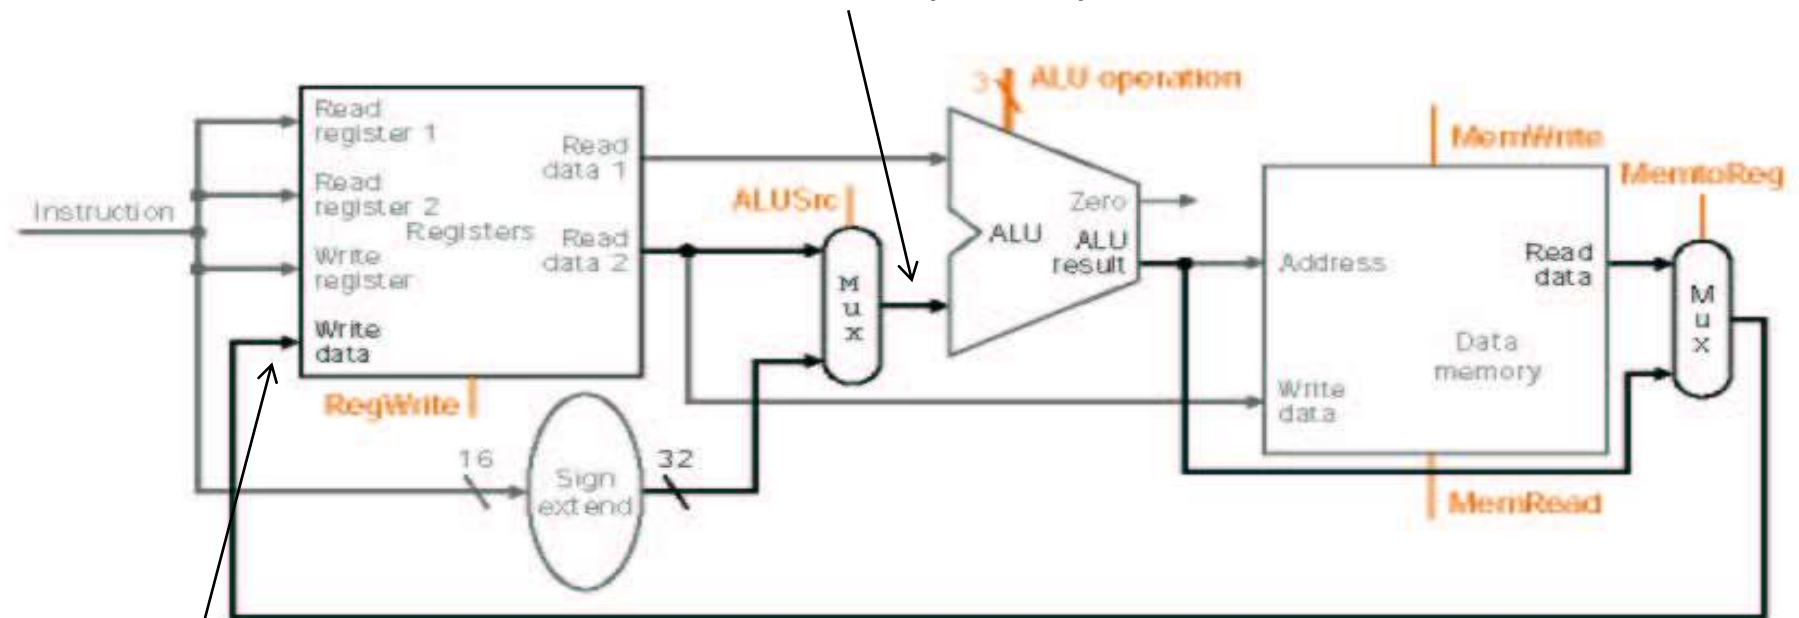
Datapath

Animating the Datapath



MIPS Datapath I: Single-Cycle

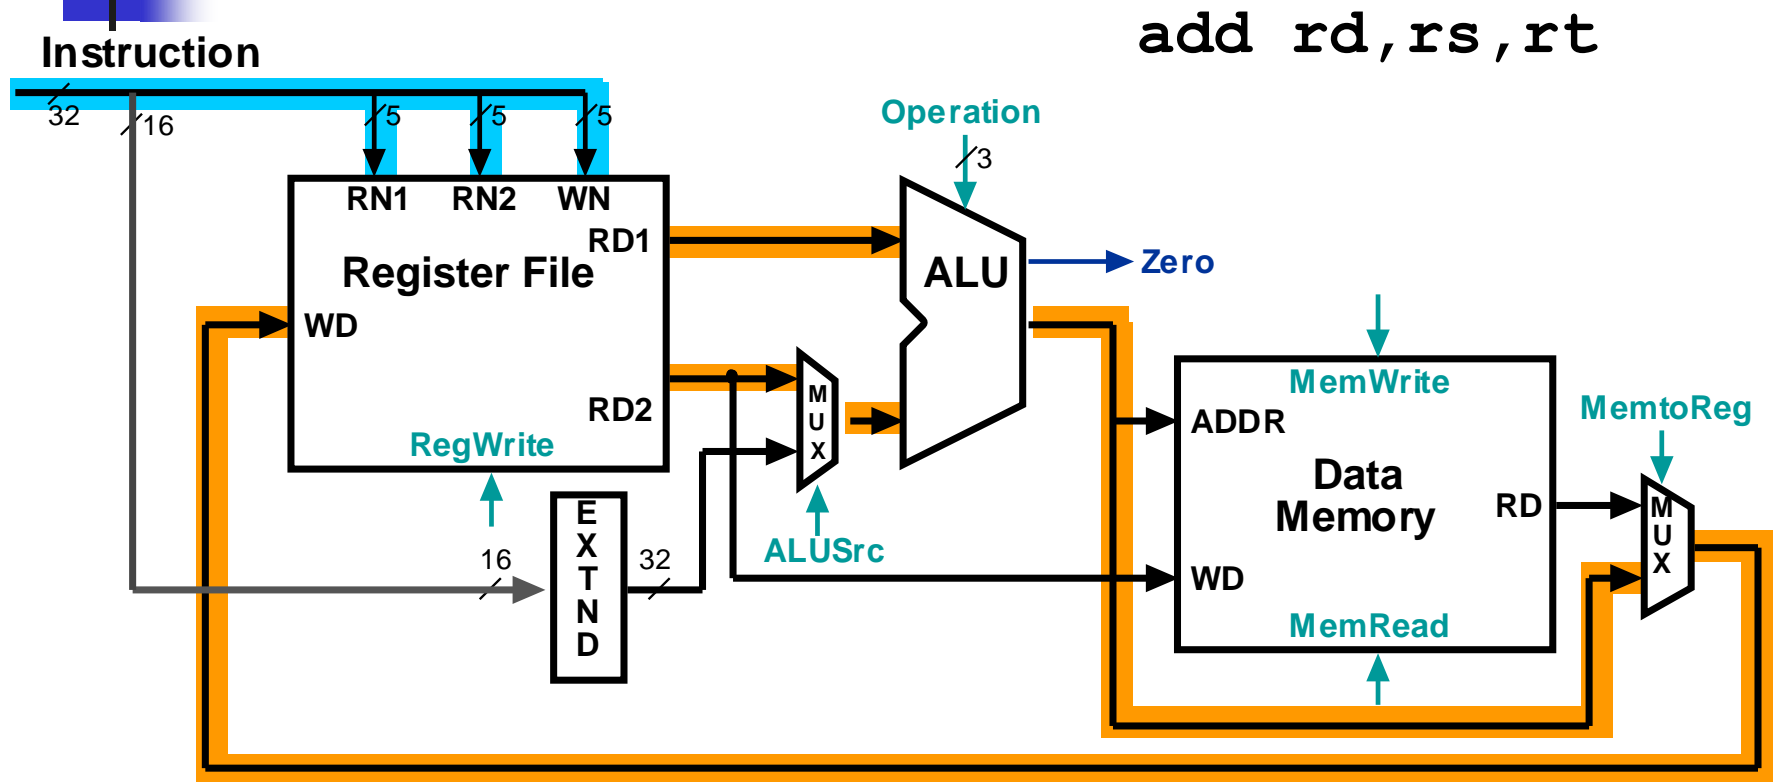
Input is either register (R-type) or sign-extended lower half of instruction (load/store)



Data is either from ALU (R-type) or memory (load)

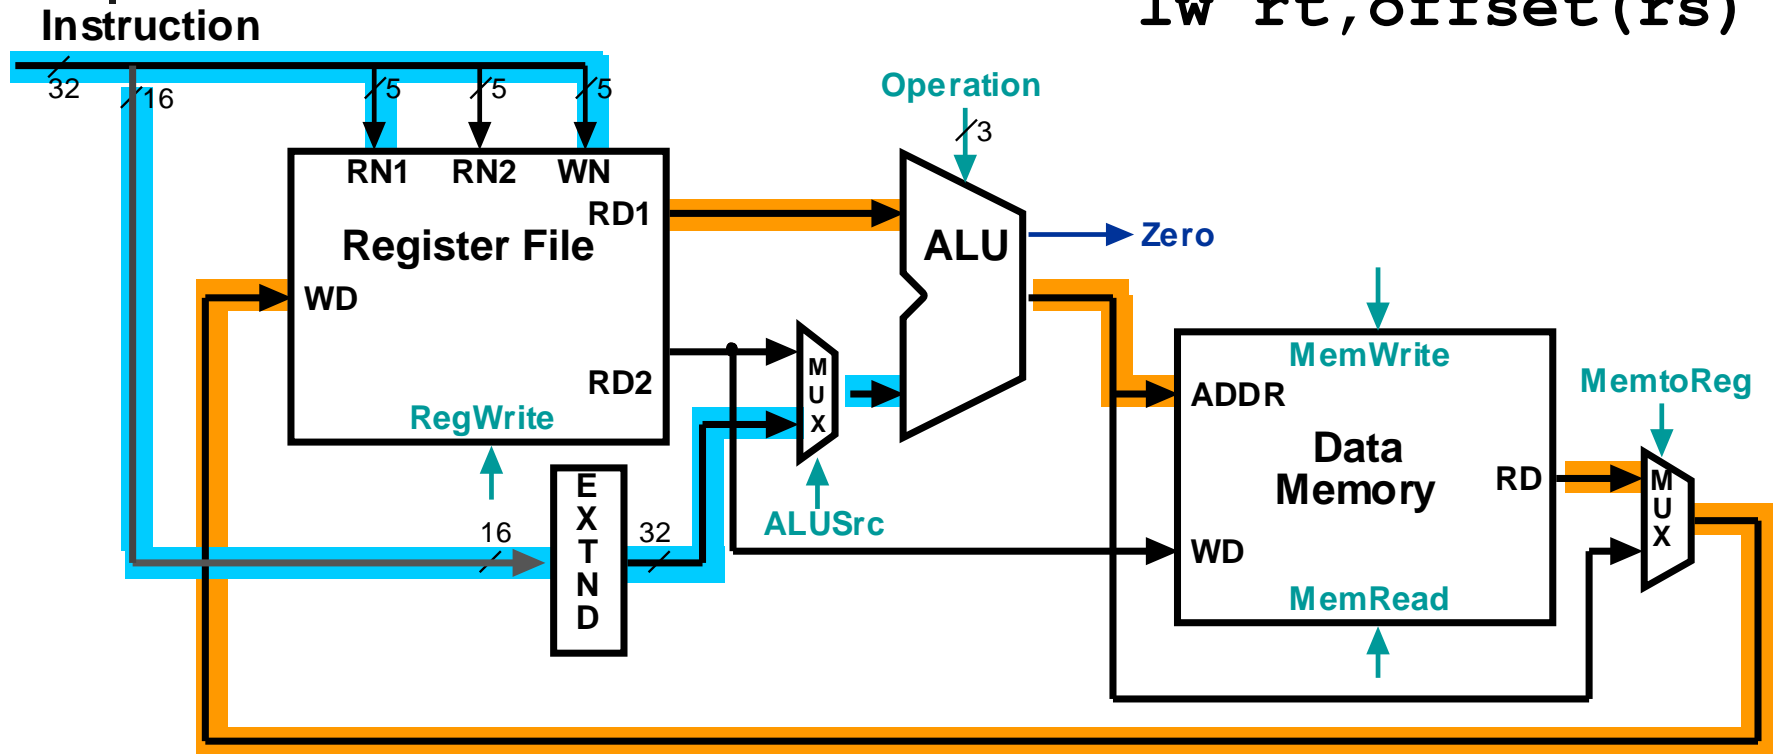
Combining the datapaths for R-type instructions and load/stores using two multiplexors

Animating the Datapath: R-type Instruction

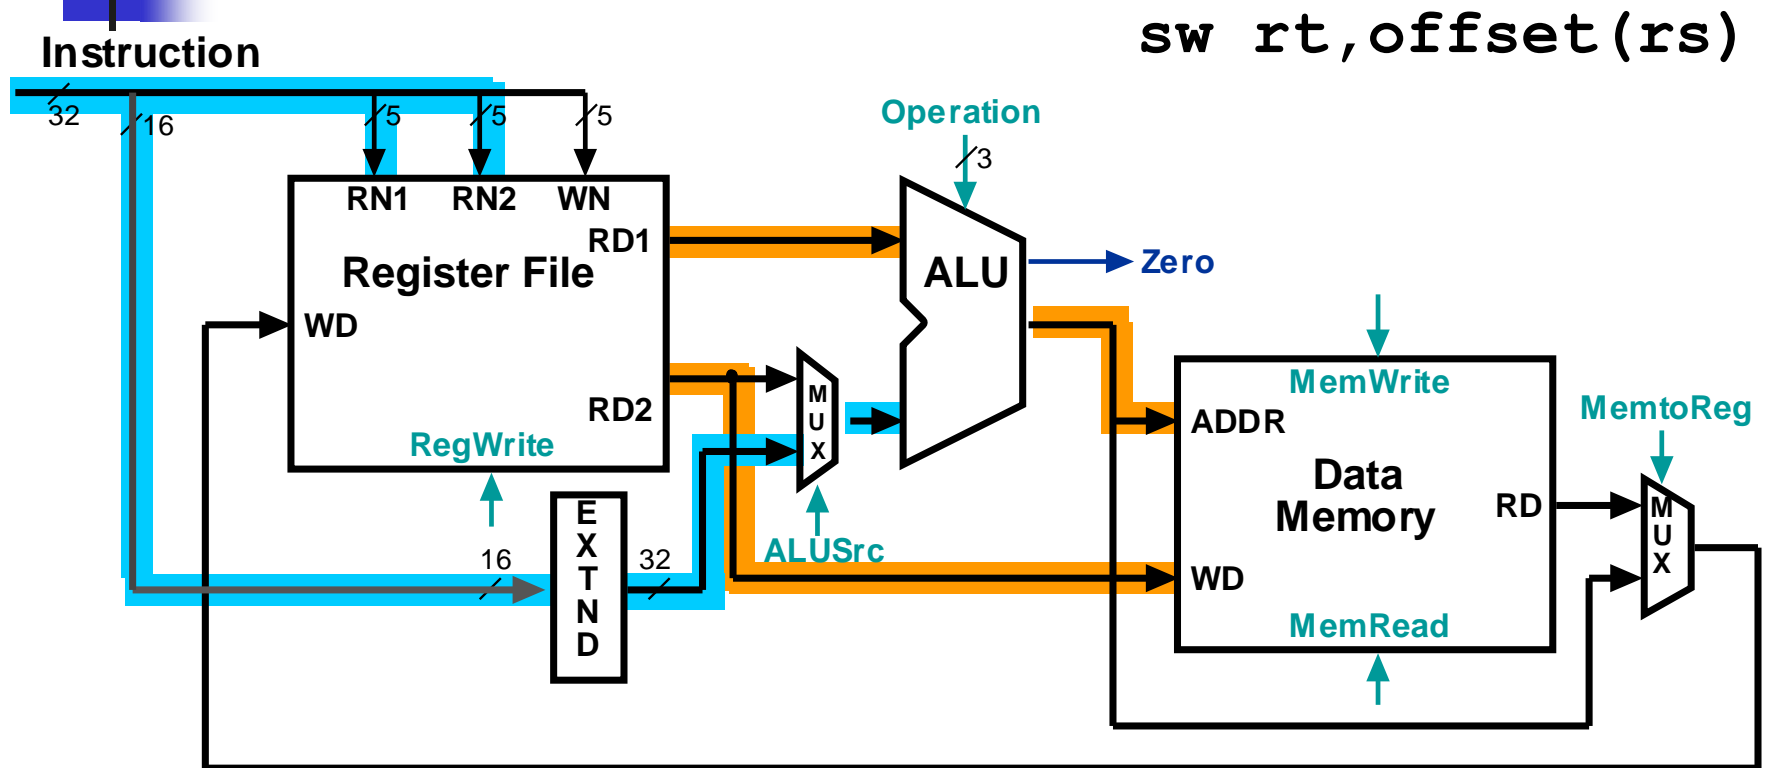


Animating the Datapath: Load Instruction


`lw rt,offset(rs)`



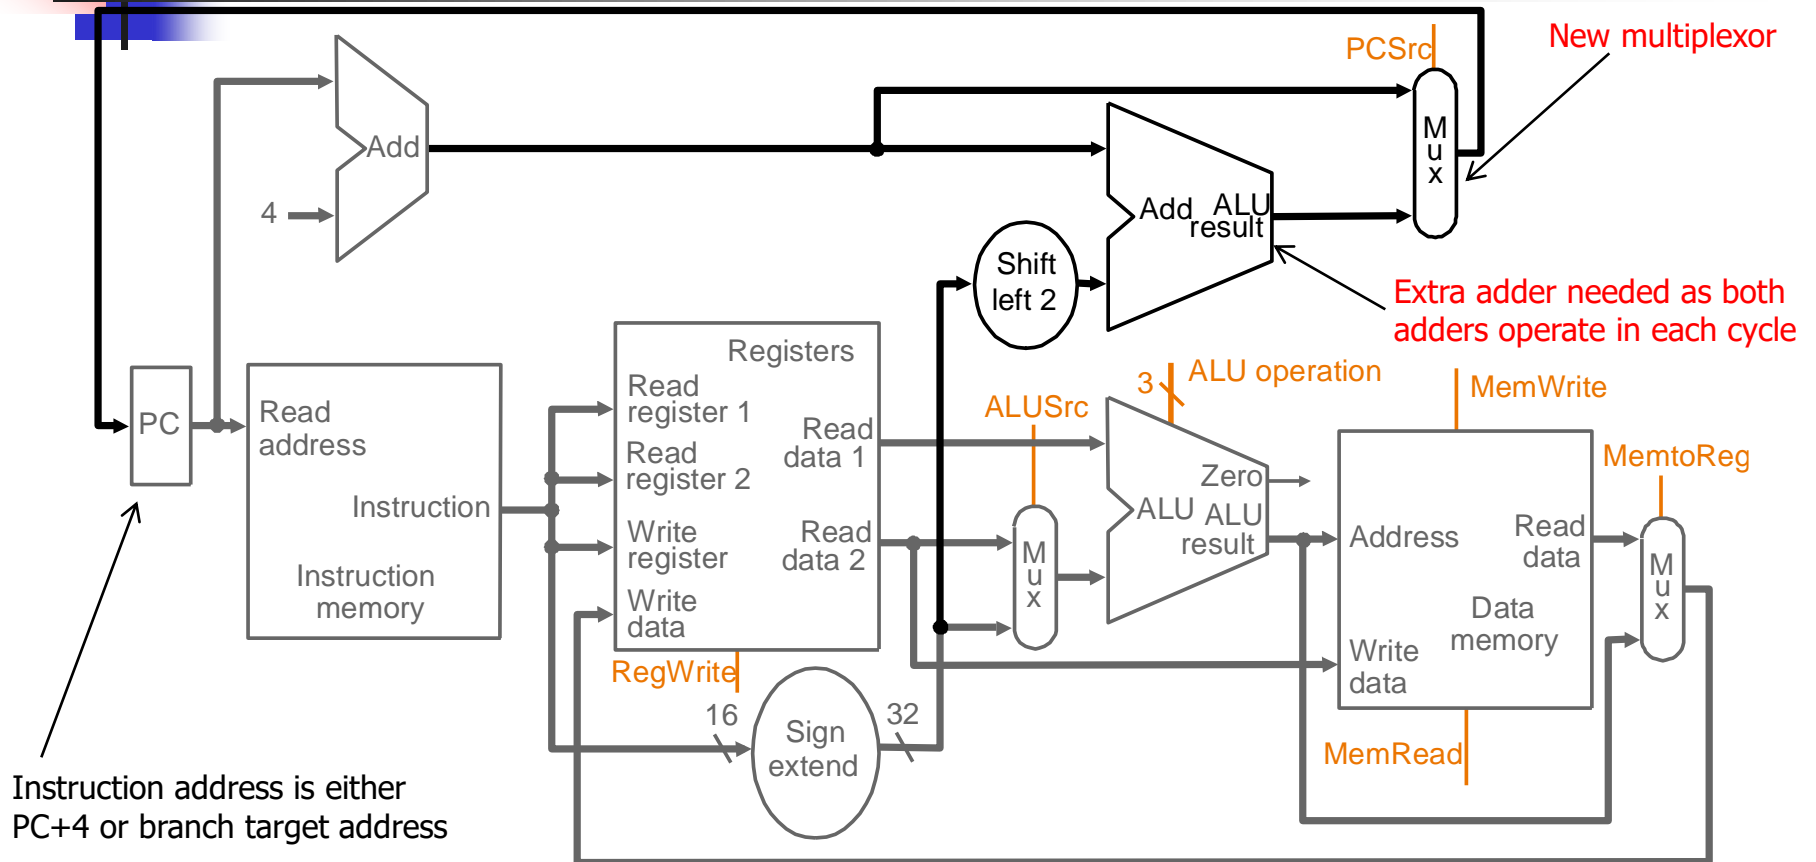
Animating the Datapath: Store Instruction



MIPS Datapath II: Single-Cycle



MIPS Datapath III: Single-Cycle

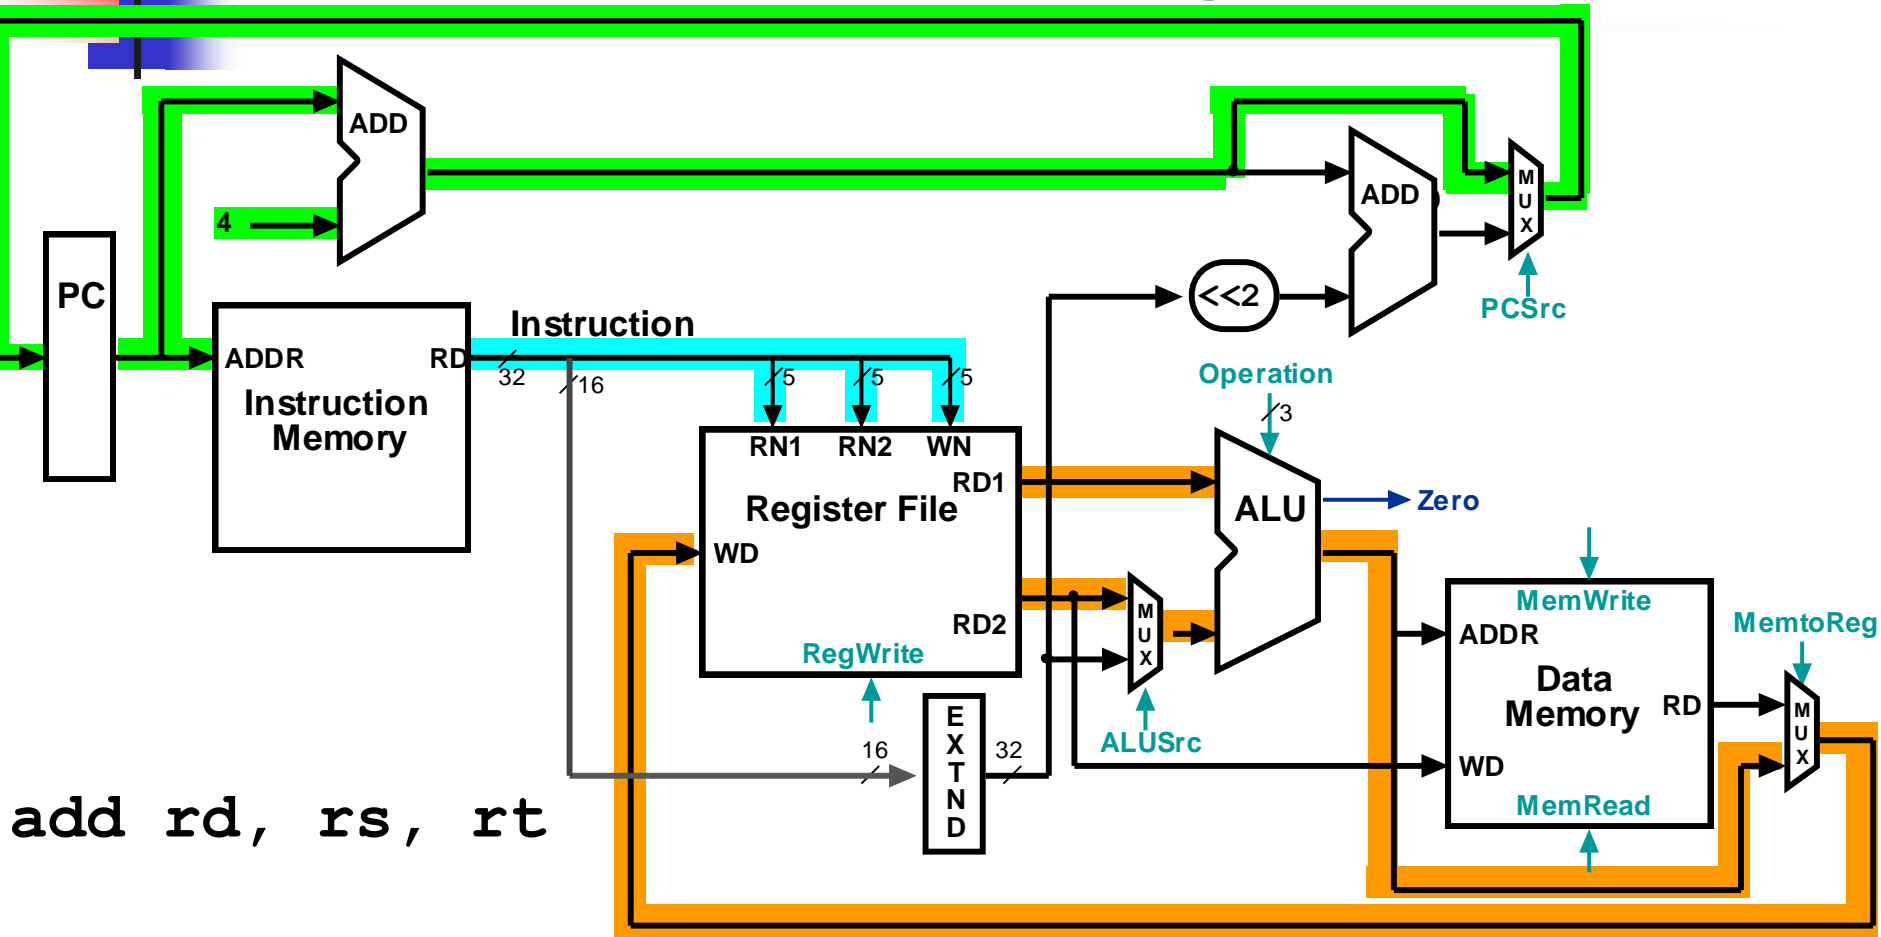


Adding branch capability and another multiplexor

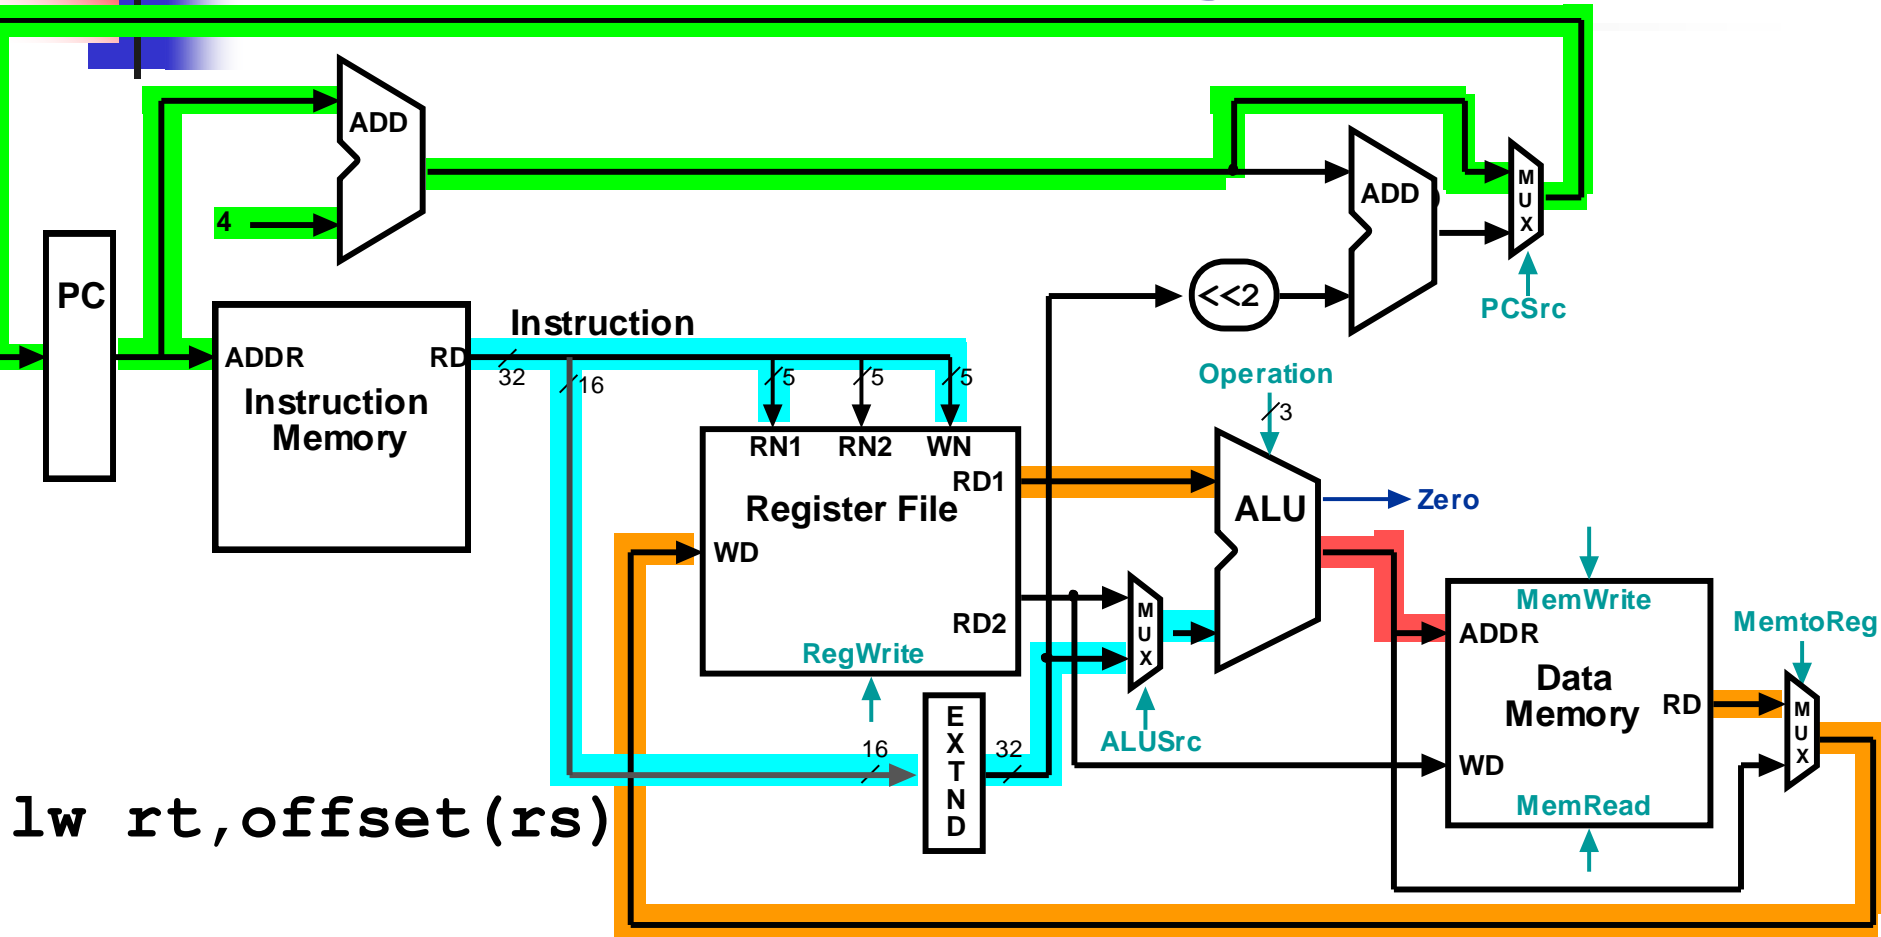
Important note: in a single-cycle implementation data cannot be stored during an instruction – it only moves through combinational logic

Question: is the MemRead signal really needed?! Think of RegWrite...!

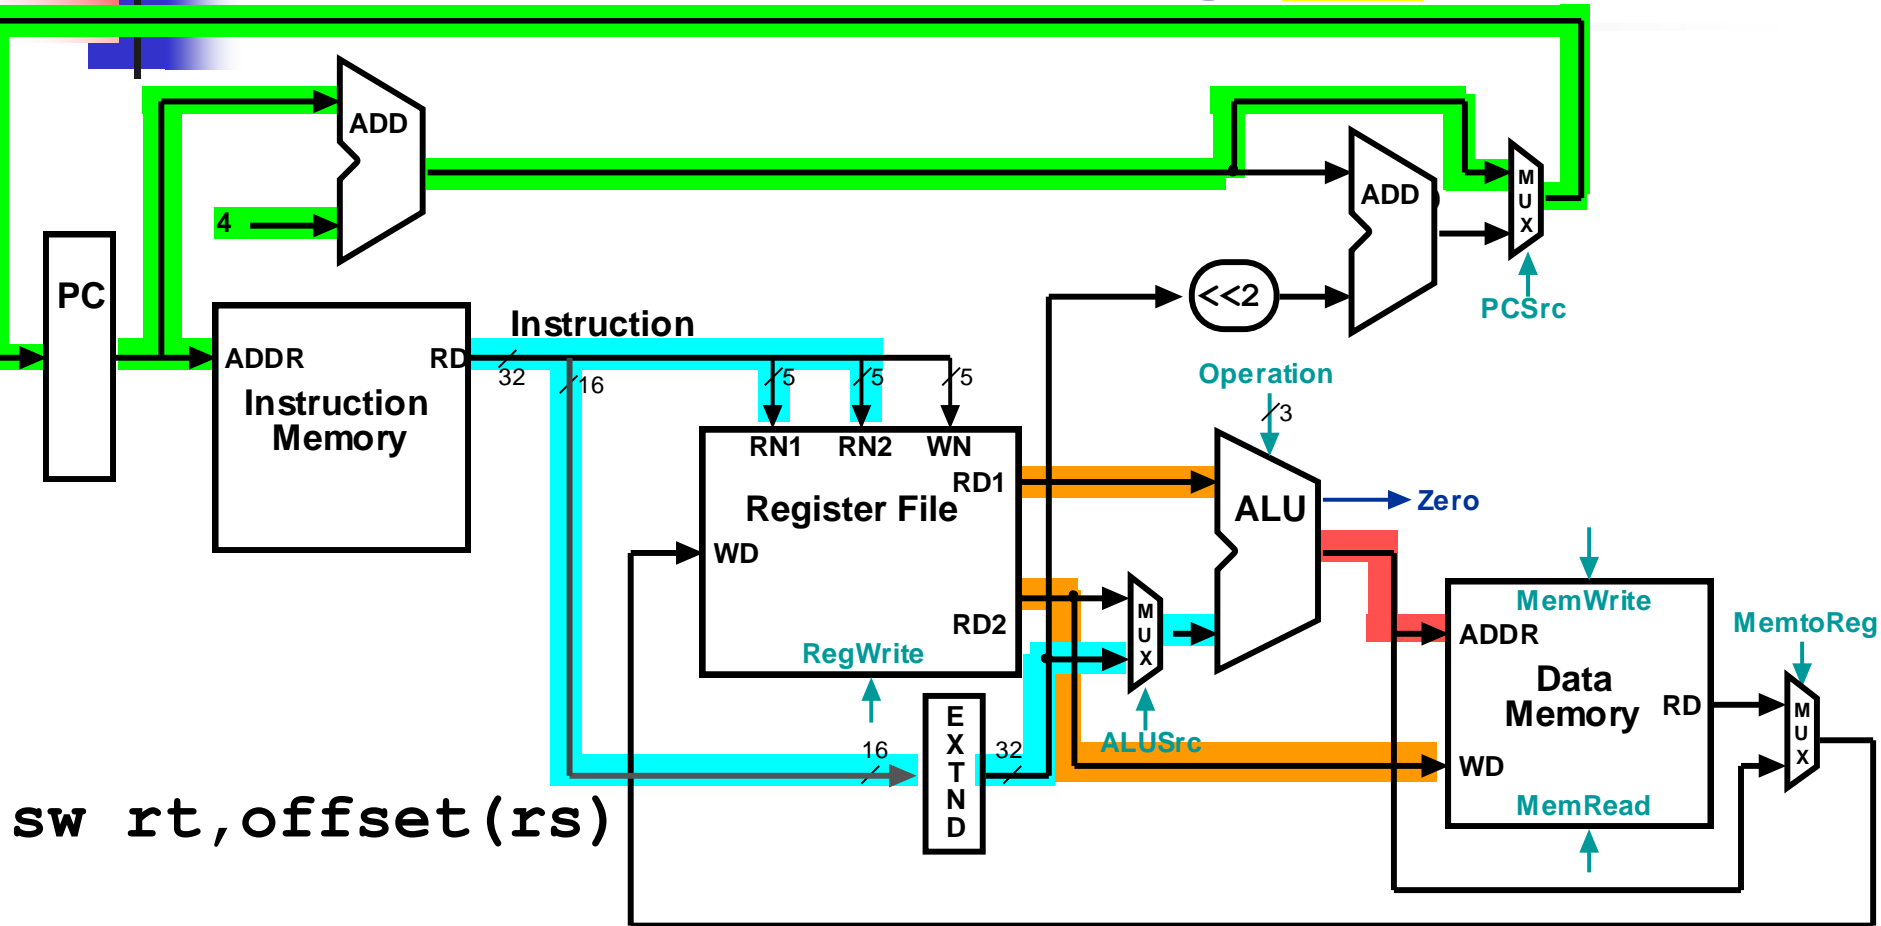
Datapath Executing `add`



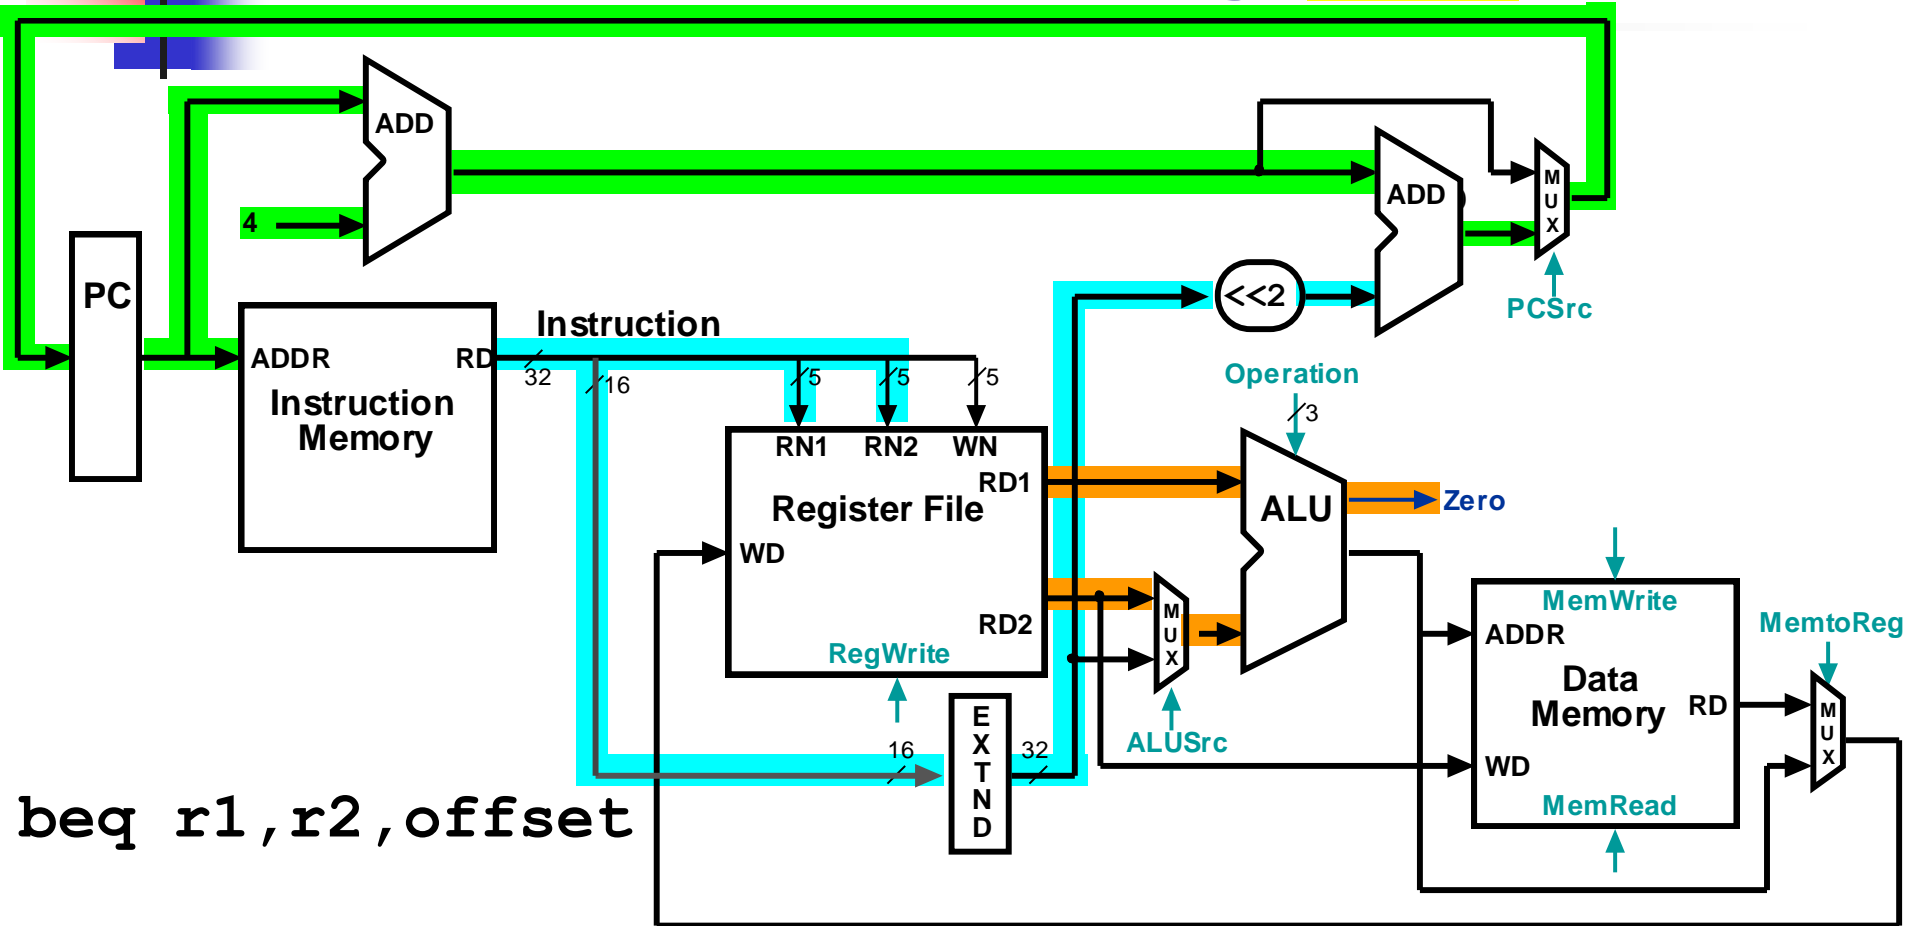
Datapath Executing lw



Datapath Executing **SW**



Datapath Executing `beq`





Control

- Control unit takes input from
 - the instruction opcode bits
- Control unit generates
 - ALU control input
 - write enable (possibly, read enable also) signals for each storage element
 - selector controls for each multiplexor

ALU Control

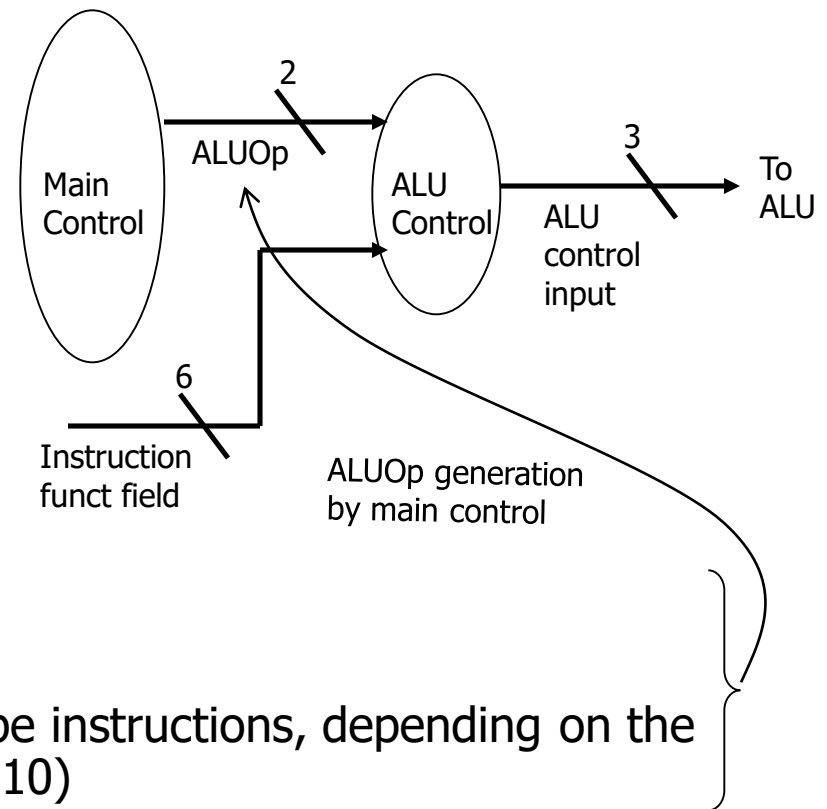
Plan to control ALU: main control sends a 2-bit ALUOp control field to the ALU control. Based on ALUOp and funct field of instruction the ALU control generates the 3-bit ALU control field

Recall from Ch. 4

ALU control field	Function
000	and
001	or
010	add
110	sub
111	slt

■ ALU must perform

- *add* for load/stores (ALUOp 00)
- *sub* for branches (ALUOp 01)
- one of *and*, *or*, *add*, *sub*, *slt* for R-type instructions, depending on the instruction's 6-bit funct field (ALUOp 10)



Setting ALU Control Bits

Instruction opcode	AluOp	Instruction operation	Funct Field	Desired ALU action	ALU control input
LW	00	load word	xxxxxx	add	010
SW	00	store word	xxxxxx	add	010
Branch eq	01	branch eq	xxxxxx	subtract	110
R-type	10	add	100000	add	010
R-type	10	subtract	100010	subtract	110
R-type	10	AND	100100	and	000
R-type	10	OR	100101	or	001
R-type	10	set on less	101010	set on less	111

ALUOp		Funct field						Operation
ALUOp1	ALUOp0	F5	F4	F3	F2	F1	F0	
0	0	X	X	X	X	X	X	010
0*	1	X	X	X	X	X	X	110
1	X	X	X	0	0	0	0	010
1	X	X	X	0	0	1	0	110
1	X	X	X	0	1	0	0	000
1	X	X	X	0	1	0	1	001
1	X	X	X	1	0	1	0	111

*Typo in text
Fig. 5.15: if it is X
then there is potential
conflict between
line 2 and lines 3-7!

Truth table for ALU control bits



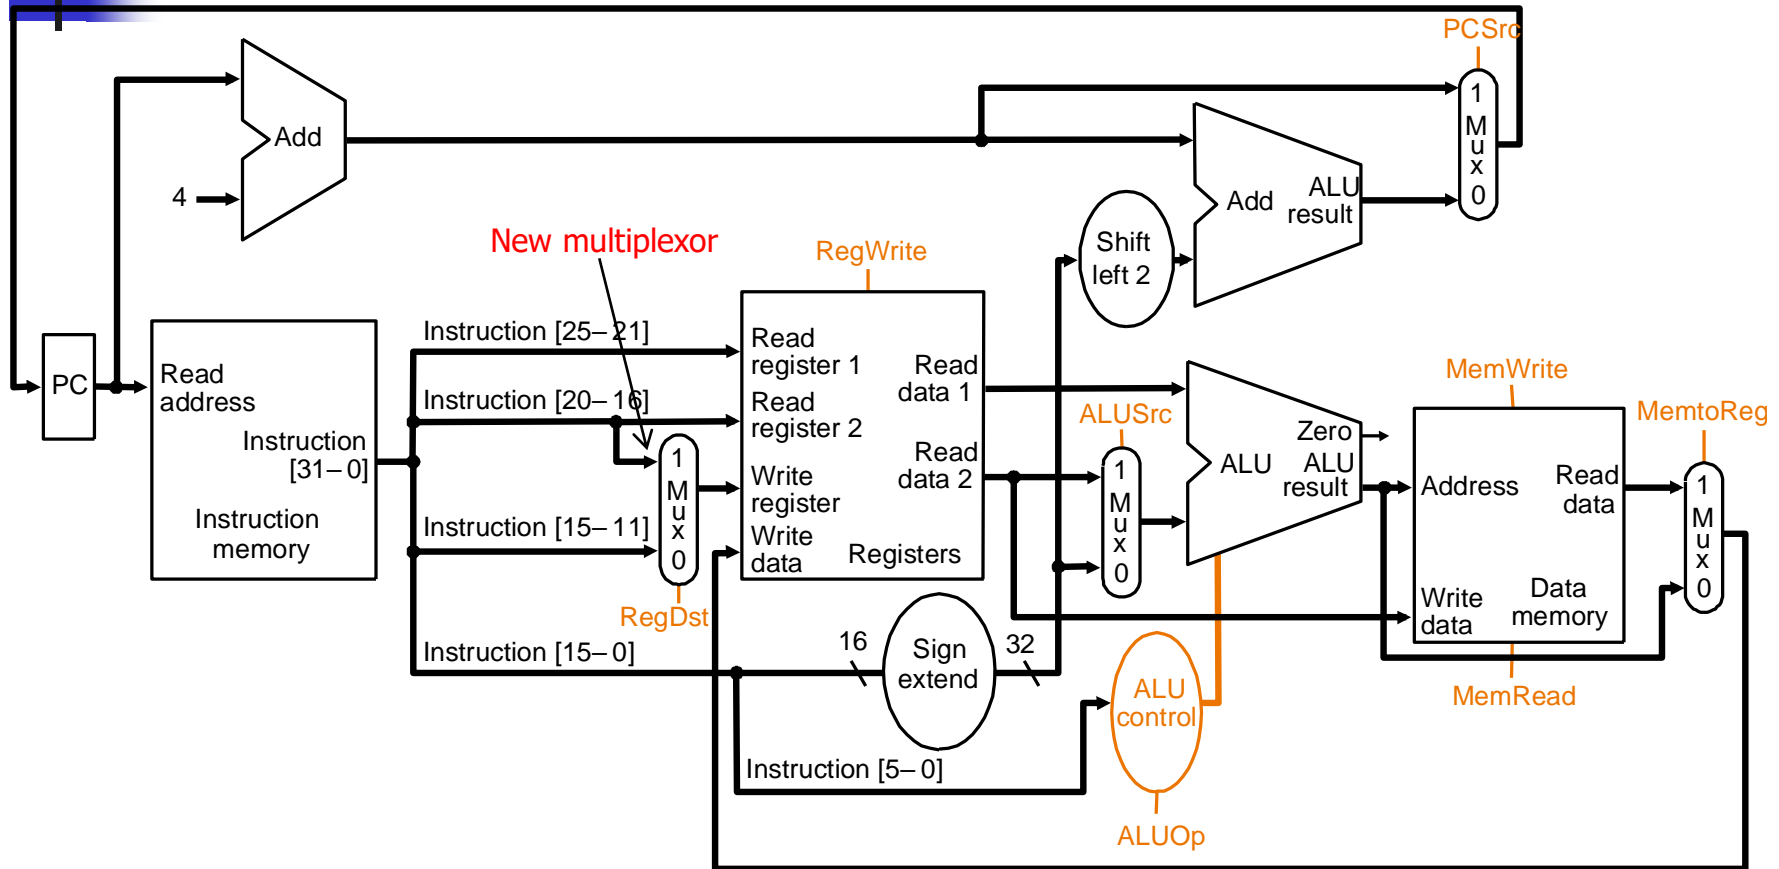
Designing the Main Control

R-type	opcode	rs	rt	rd	shamt	funct
	31-26	25-21	20-16	15-11	10-6	5-0

Load/store or branch	opcode	rs	rt	address
	31-26	25-21	20-16	15-0

- Observations about MIPS instruction format
 - opcode is always in bits 31-26
 - two registers to be read are always rs (bits 25-21) and rt (bits 20-16)
 - base register for load/stores is always rs (bits 25-21)
 - 16-bit offset for branch equal and load/store is always bits 15-0
 - destination register for loads is in bits 20-16 (rt) while for R-type instructions it is in bits 15-11 (rd) (*will require multiplexor to select*)

Datapath with Control I



Adding control to the MIPS Datapath III (and a new multiplexor to select field to specify destination register): **what are the functions of the 9 control signals?**

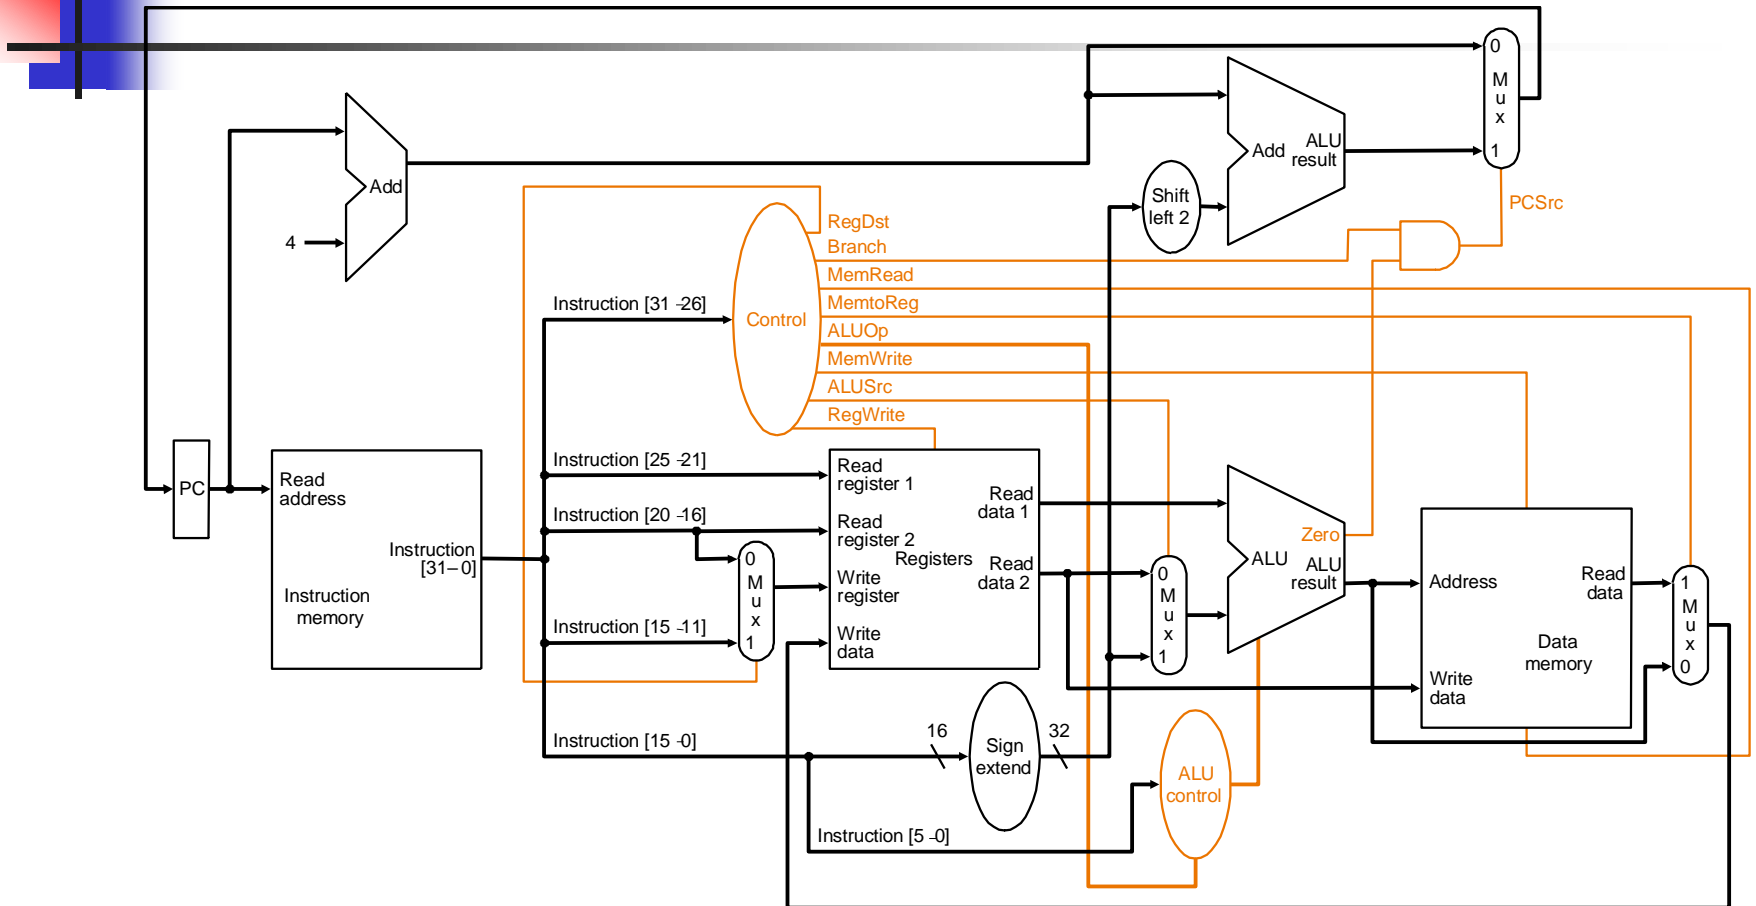


Control Signals

Signal Name	Effect when deasserted	Effect when asserted
RegDst	The register destination number for the Write register comes from the rt field (bits 20-16)	The register destination number for the Write register comes from the rd field (bits 15-11)
RegWrite	None	The register on the Write register input is written with the value on the Write data input
ALUSrc	The second ALU operand comes from the second register file output (Read data 2)	The second ALU operand is the sign-extended, lower 16 bits of the instruction
PCSrc	The PC is replaced by the output of the adder that computes the value of PC + 4	The PC is replaced by the output of the adder that computes the branch target
MemRead	None	Data memory contents designated by the address input are put on the first Read data output
MemWrite	None	Data memory contents designated by the address input are replaced by the value of the Write data input
MemtoReg	The value fed to the register Write data input comes from the ALU	The value fed to the register Write data input comes from the data memory

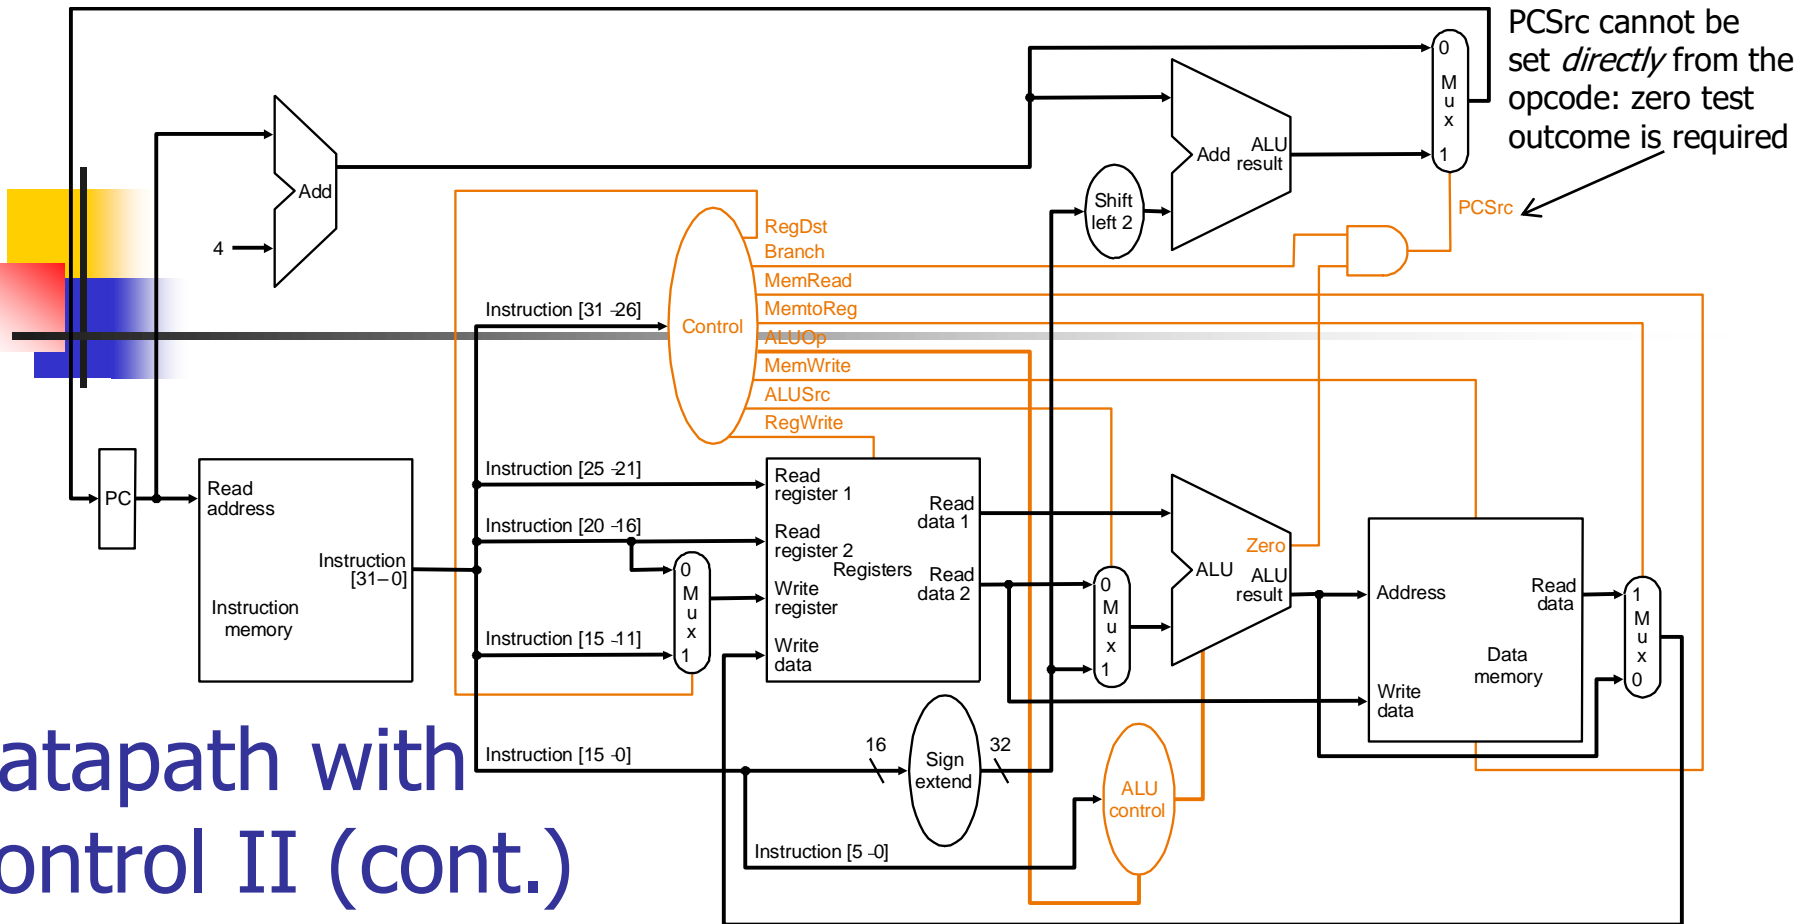
Effects of the seven control signals

Datapath with Control II



MIPS datapath with the control unit: input to control is the 6-bit instruction opcode field, output is seven 1-bit signals and the 2-bit ALUOp signal

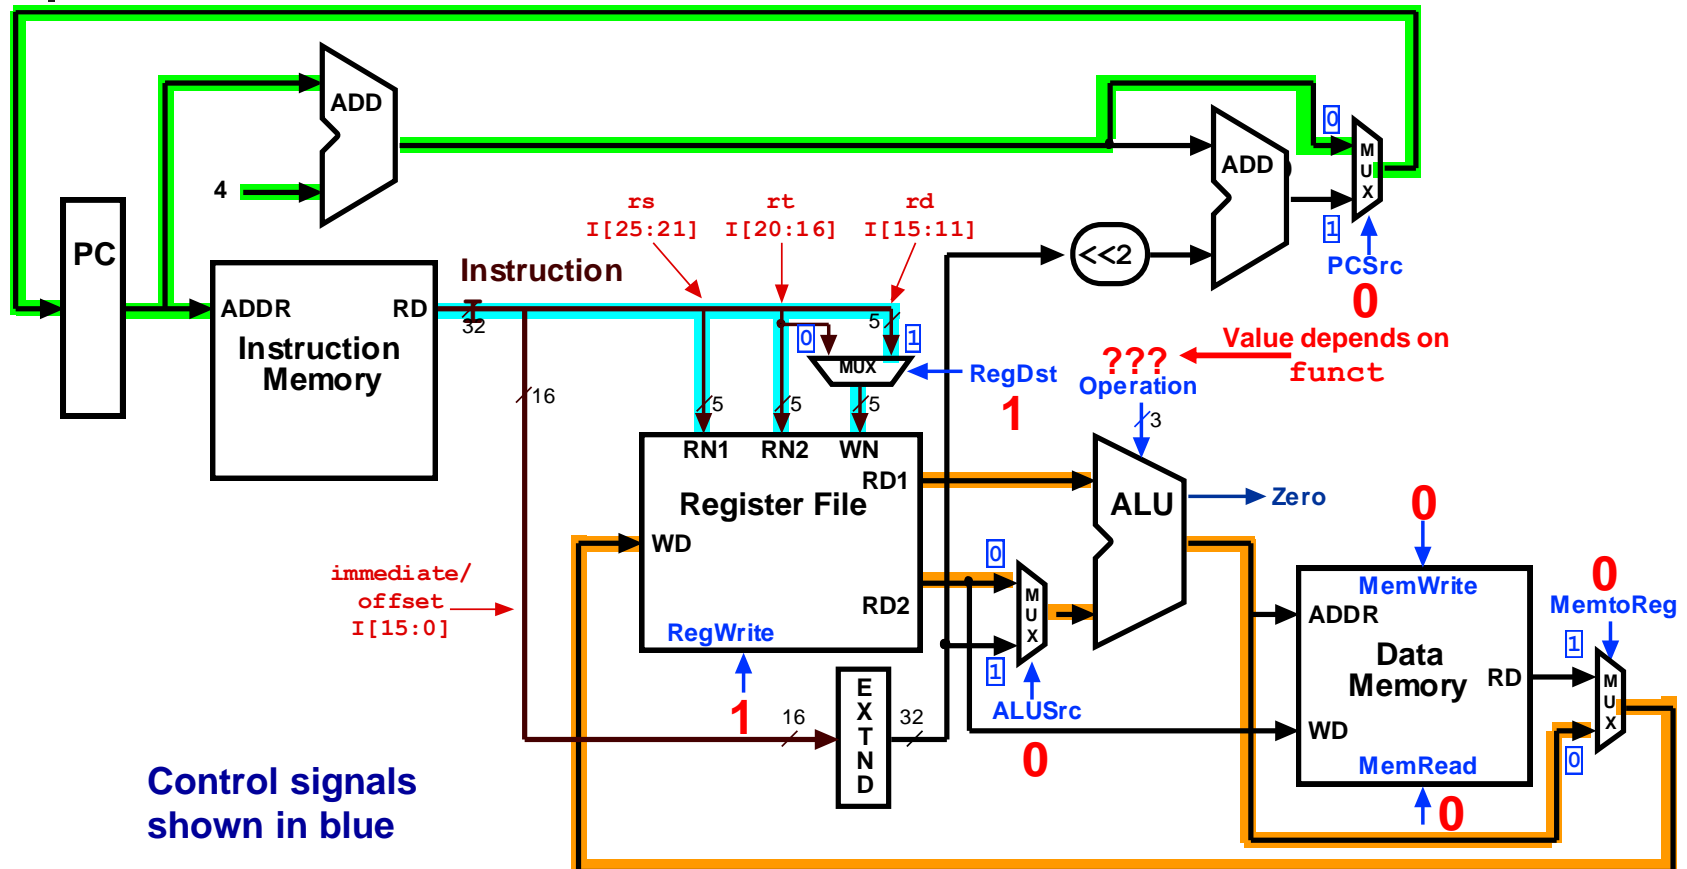
Datapath with Control II (cont.)



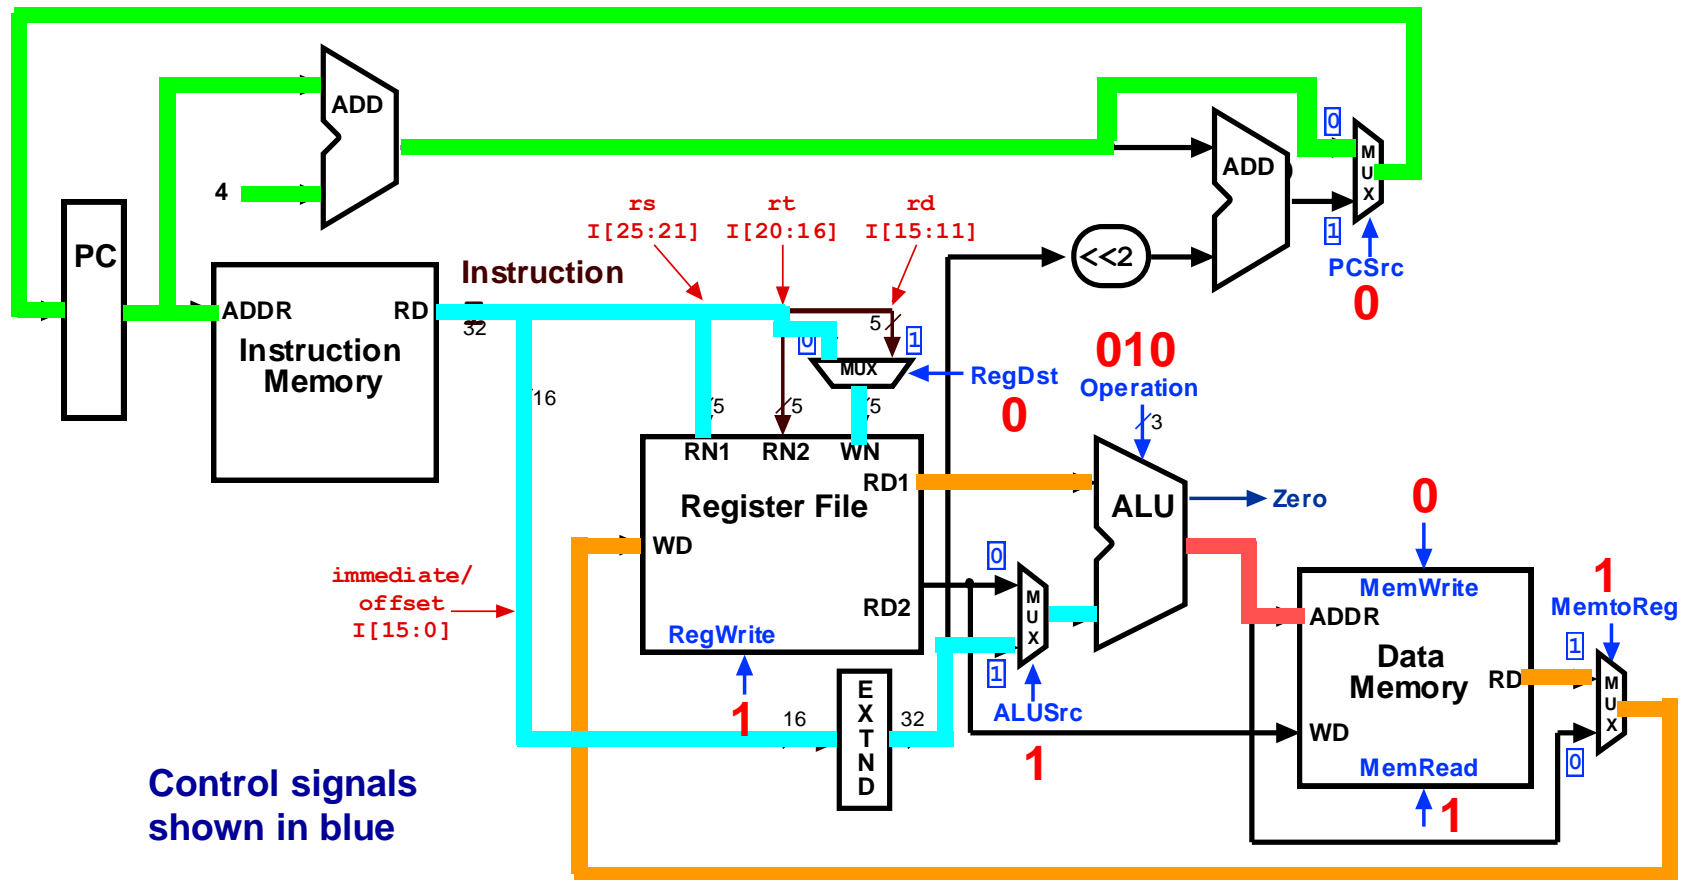
Determining control signals for the MIPS datapath based on instruction opcode

Instruction	RegDst	ALUSrc	Memto-Reg	Reg Write	Mem Read	Mem Write	Branch	ALUOp1	ALUp0
R-format	1	0	0	1	0	0	0	1	0
lw	0	1	1	1	1	0	0	0	0
sw	X	1	X	0	0	1	0	0	0
beq	X	0	X	0	0	0	1	0	1

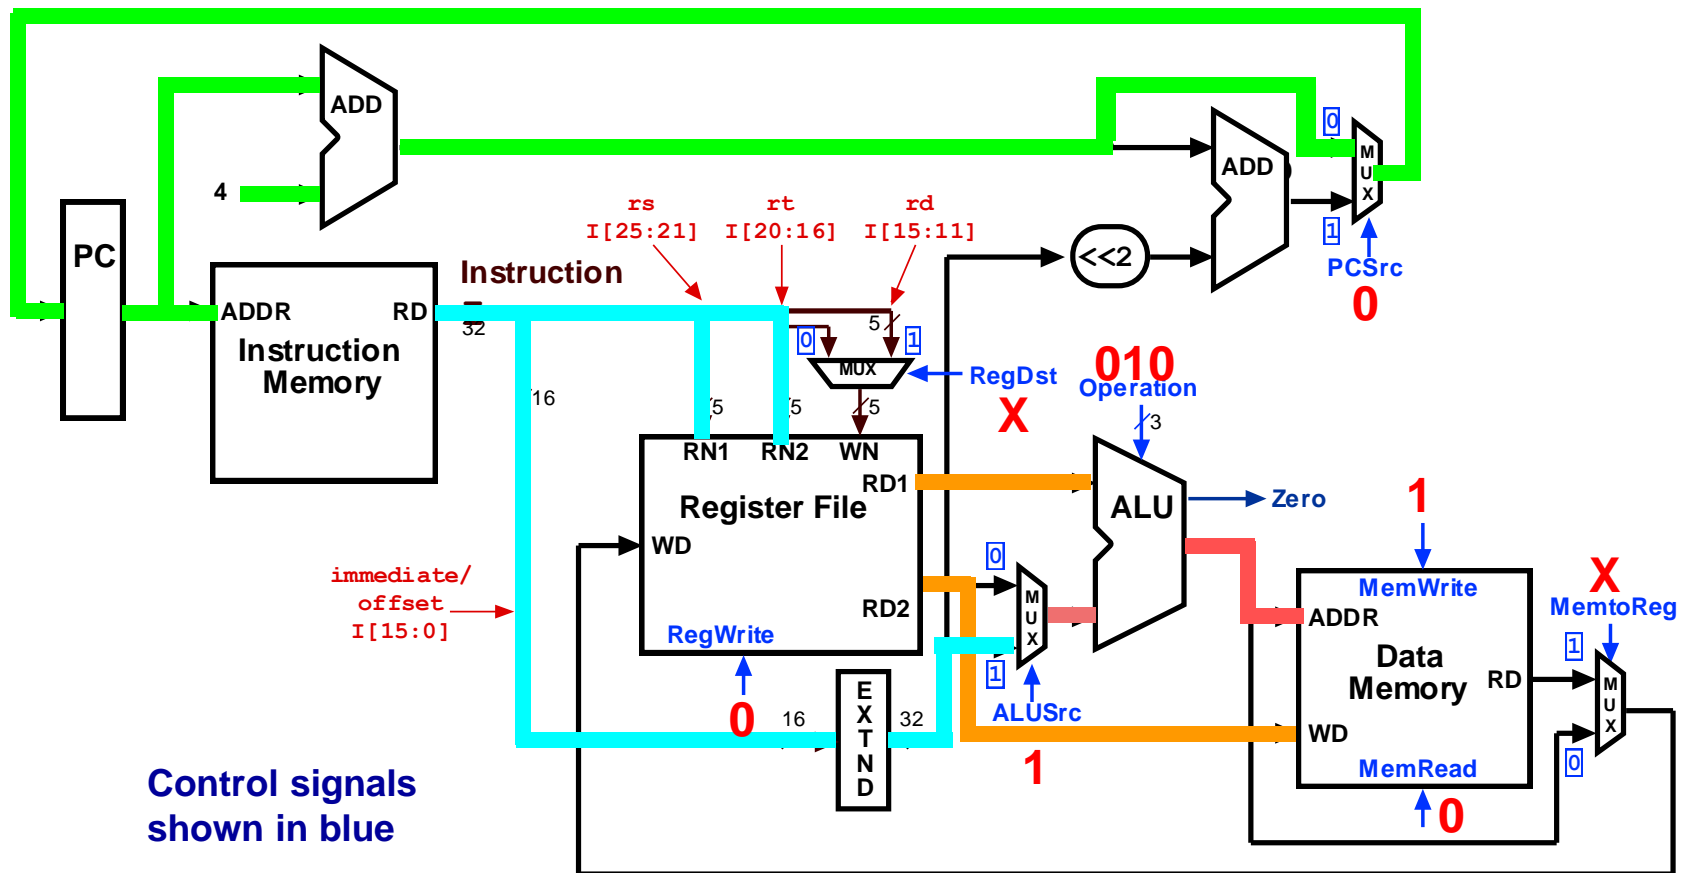
Control Signals: R-Type Instruction



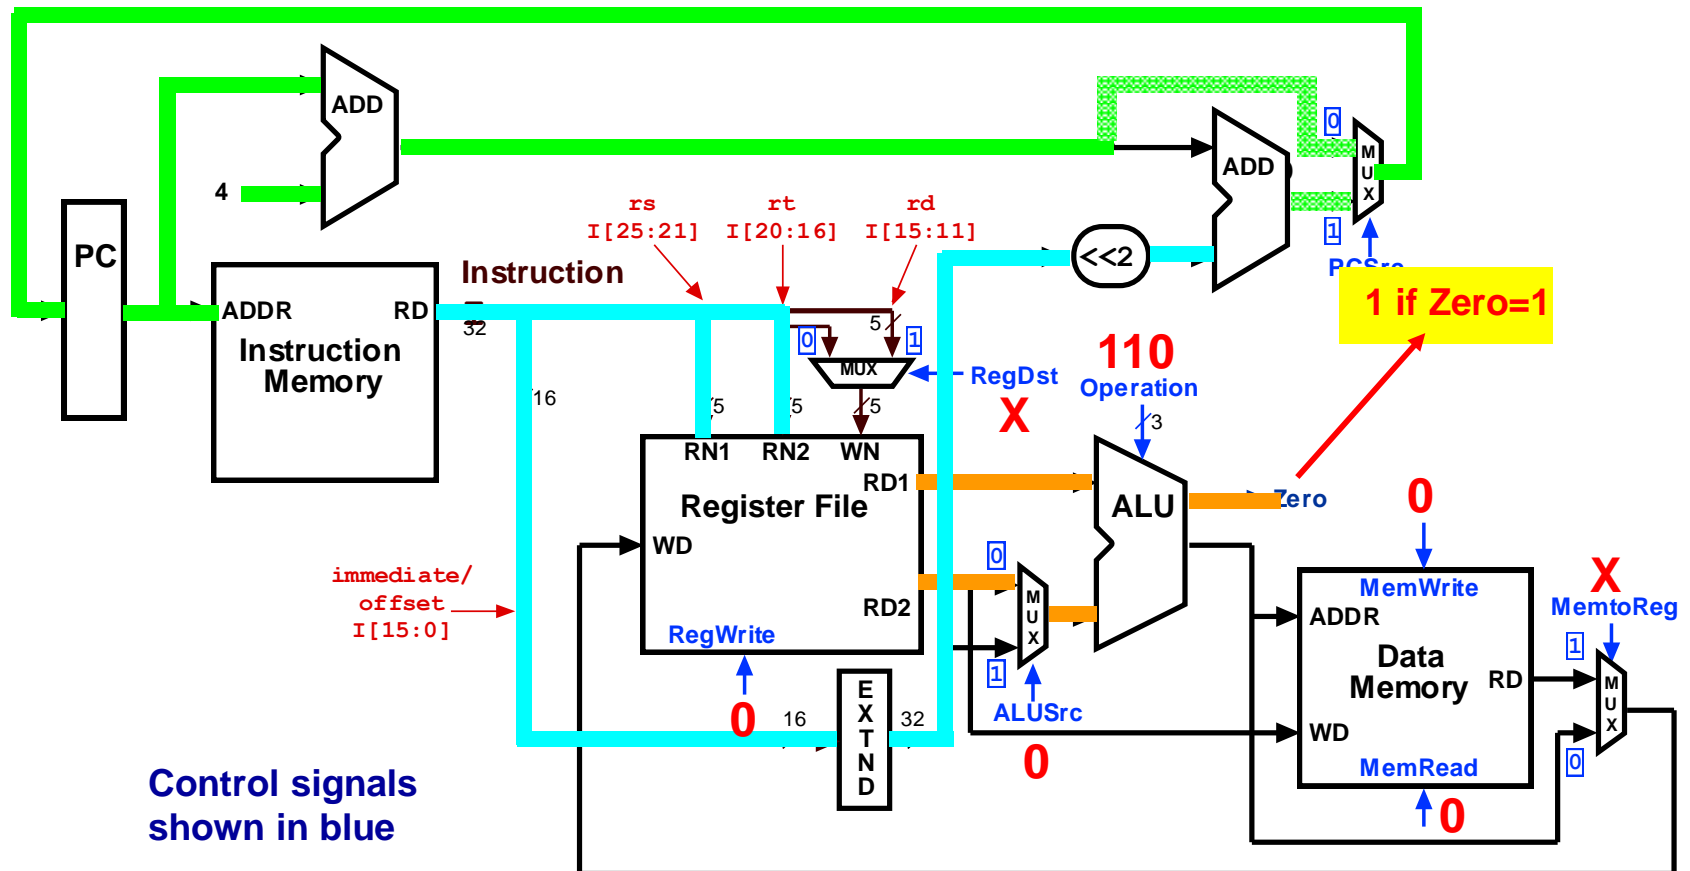
Control Signals: lw Instruction



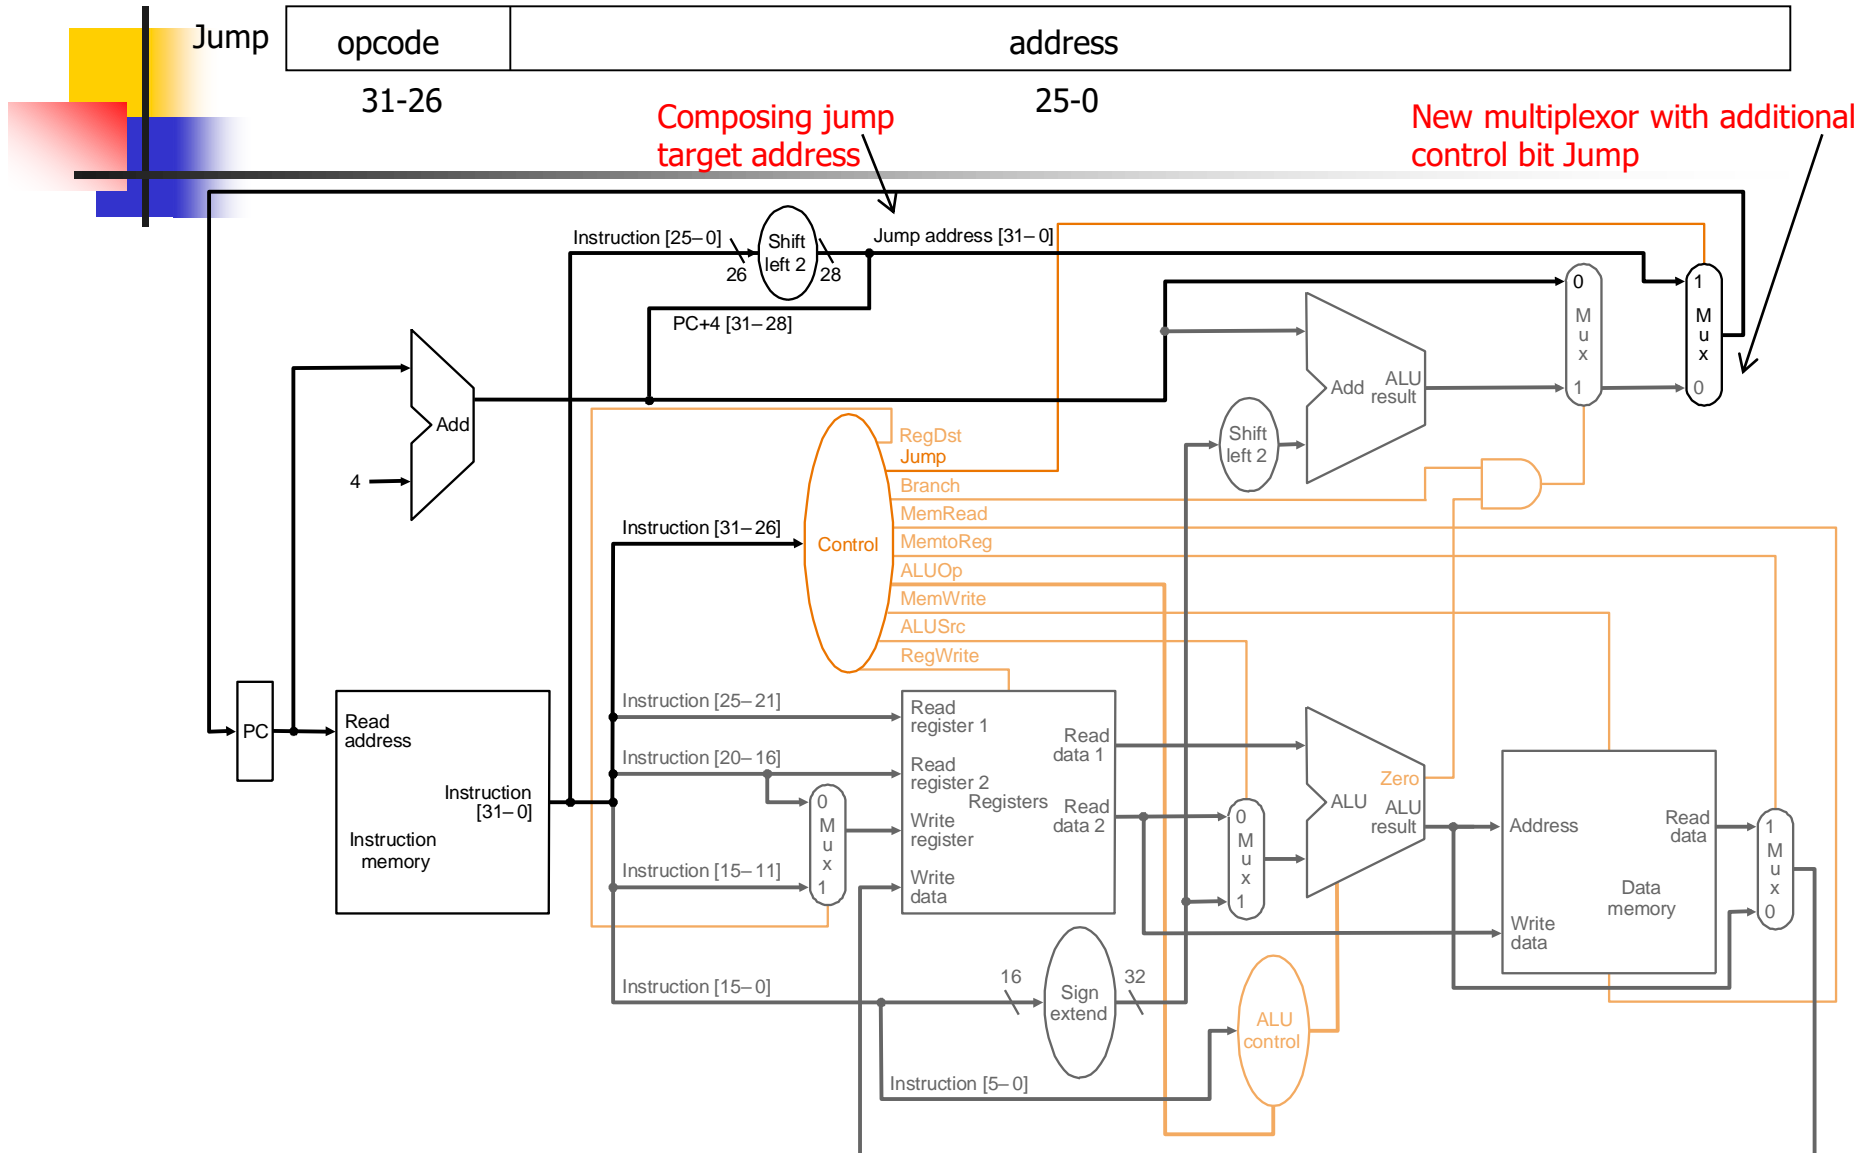
Control Signals: SW Instruction



Control Signals: beq Instruction

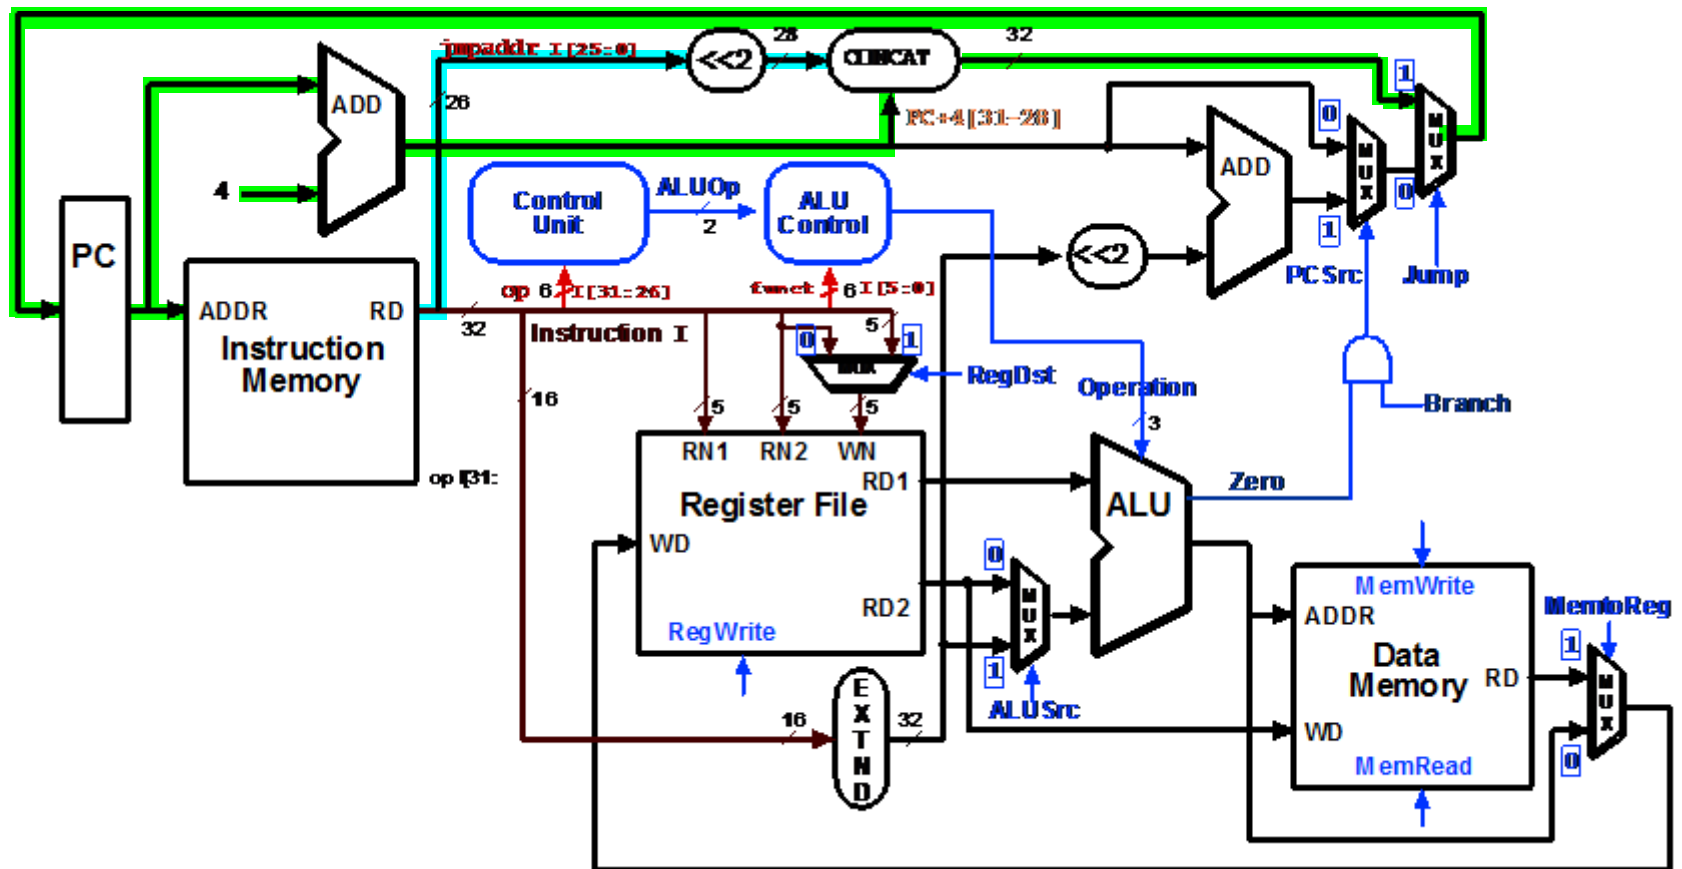


Datapath with Control III



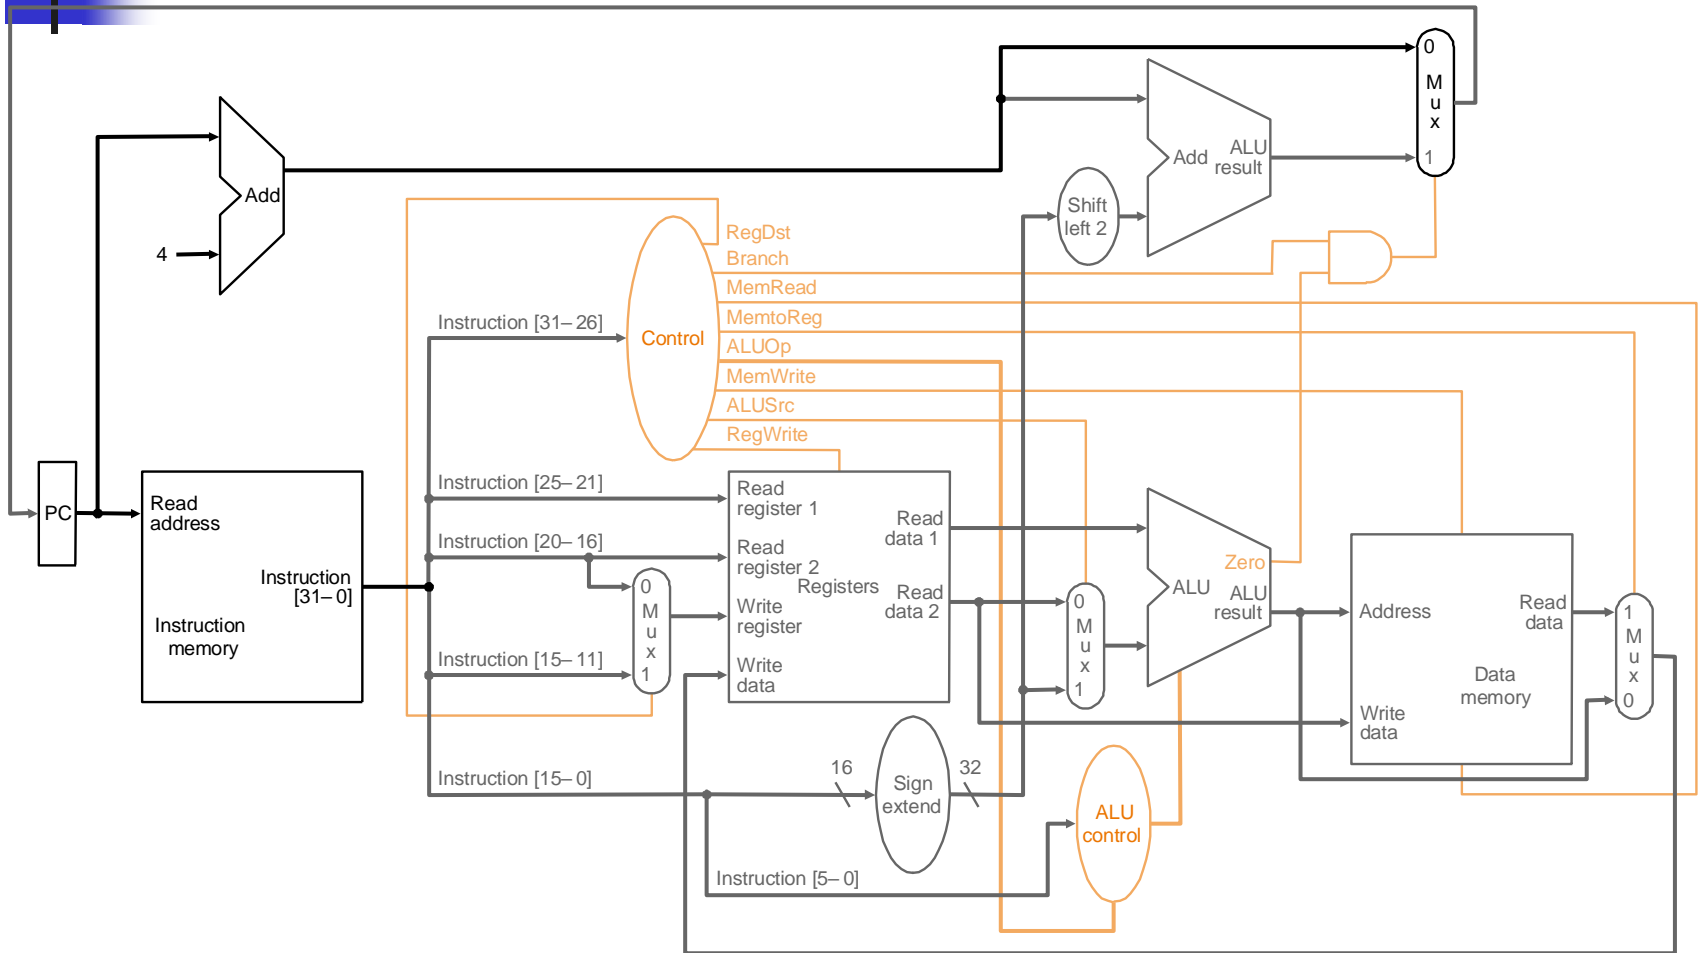
MIPS datapath extended to jumps: control unit generates new Jump control bit

Datapath Executing j



R-type Instruction: Step 1

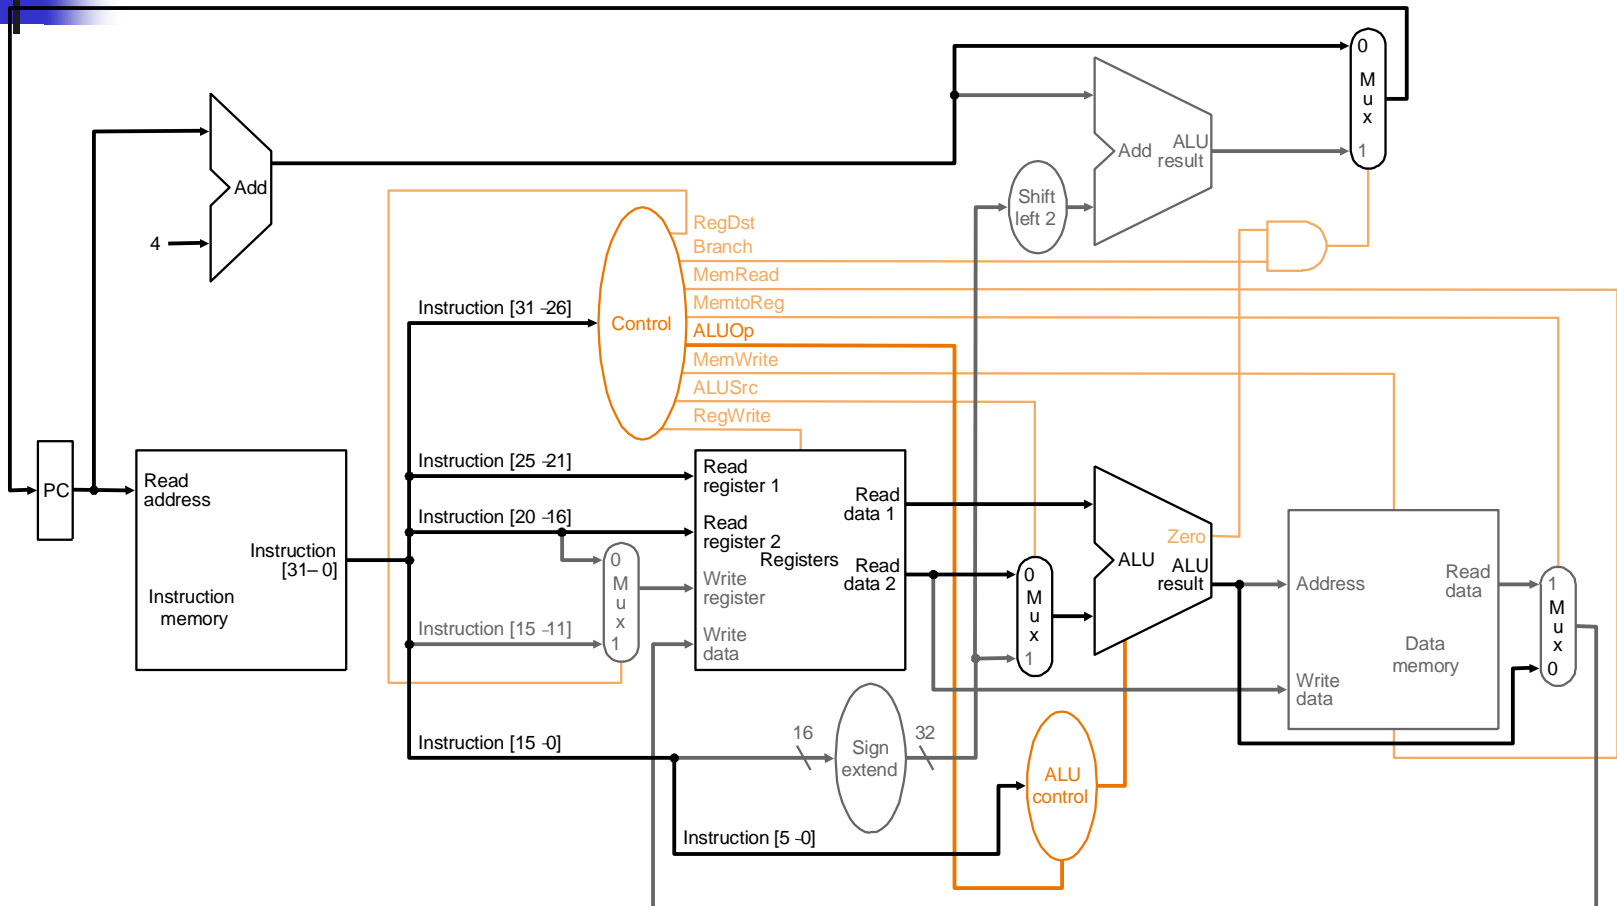
add \$t1, \$t2, \$t3 (active = bold)



Fetch instruction and increment PC count

R-type Instruction: Step 3

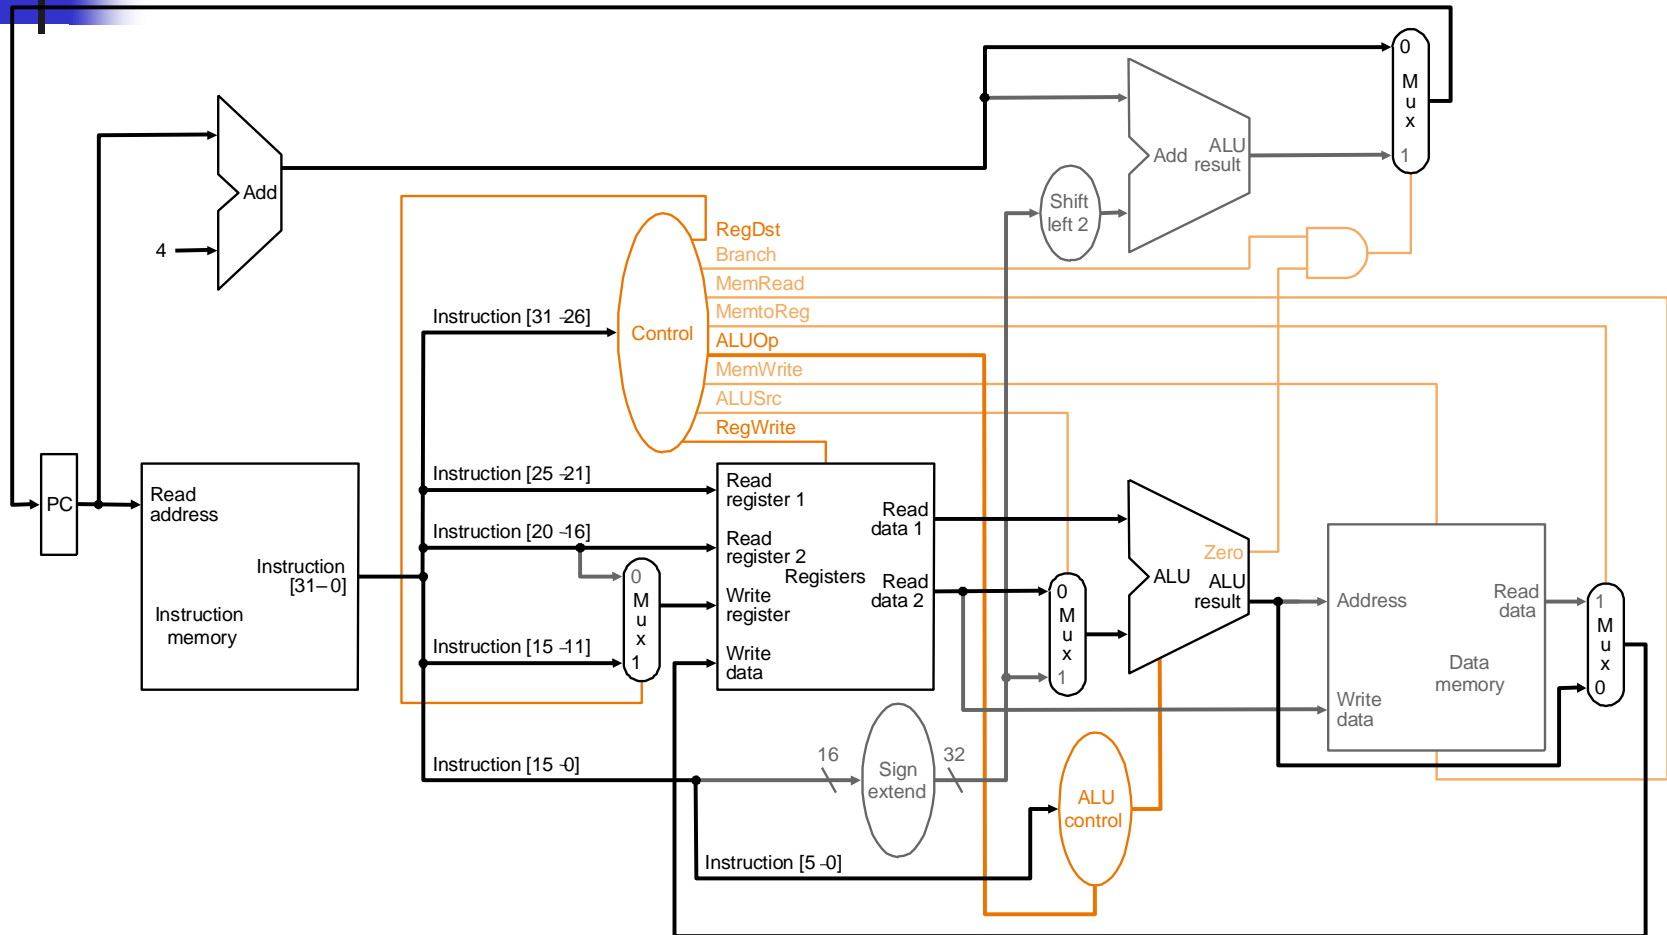
add \$t1, \$t2, \$t3 (active = bold)



ALU operates on the two register operands

R-type Instruction: Step 4

add \$t1, \$t2, \$t3 (active = bold)



Write result to register



Single-cycle Implementation Notes

- *The steps are not really distinct* as each instruction completes in exactly one clock cycle – they simply indicate the sequence of data flowing through the datapath
- *The operation of the datapath during a cycle is purely combinational* – nothing is stored during a clock cycle
- Therefore, the machine is stable in a particular state at the start of a cycle and reaches a new stable state only at the end of the cycle
- *Very important for understanding single-cycle computing:*
See our simple Verilog single-cycle computer in the folder SimpleSingleCycleComputer in Verilog/Examples



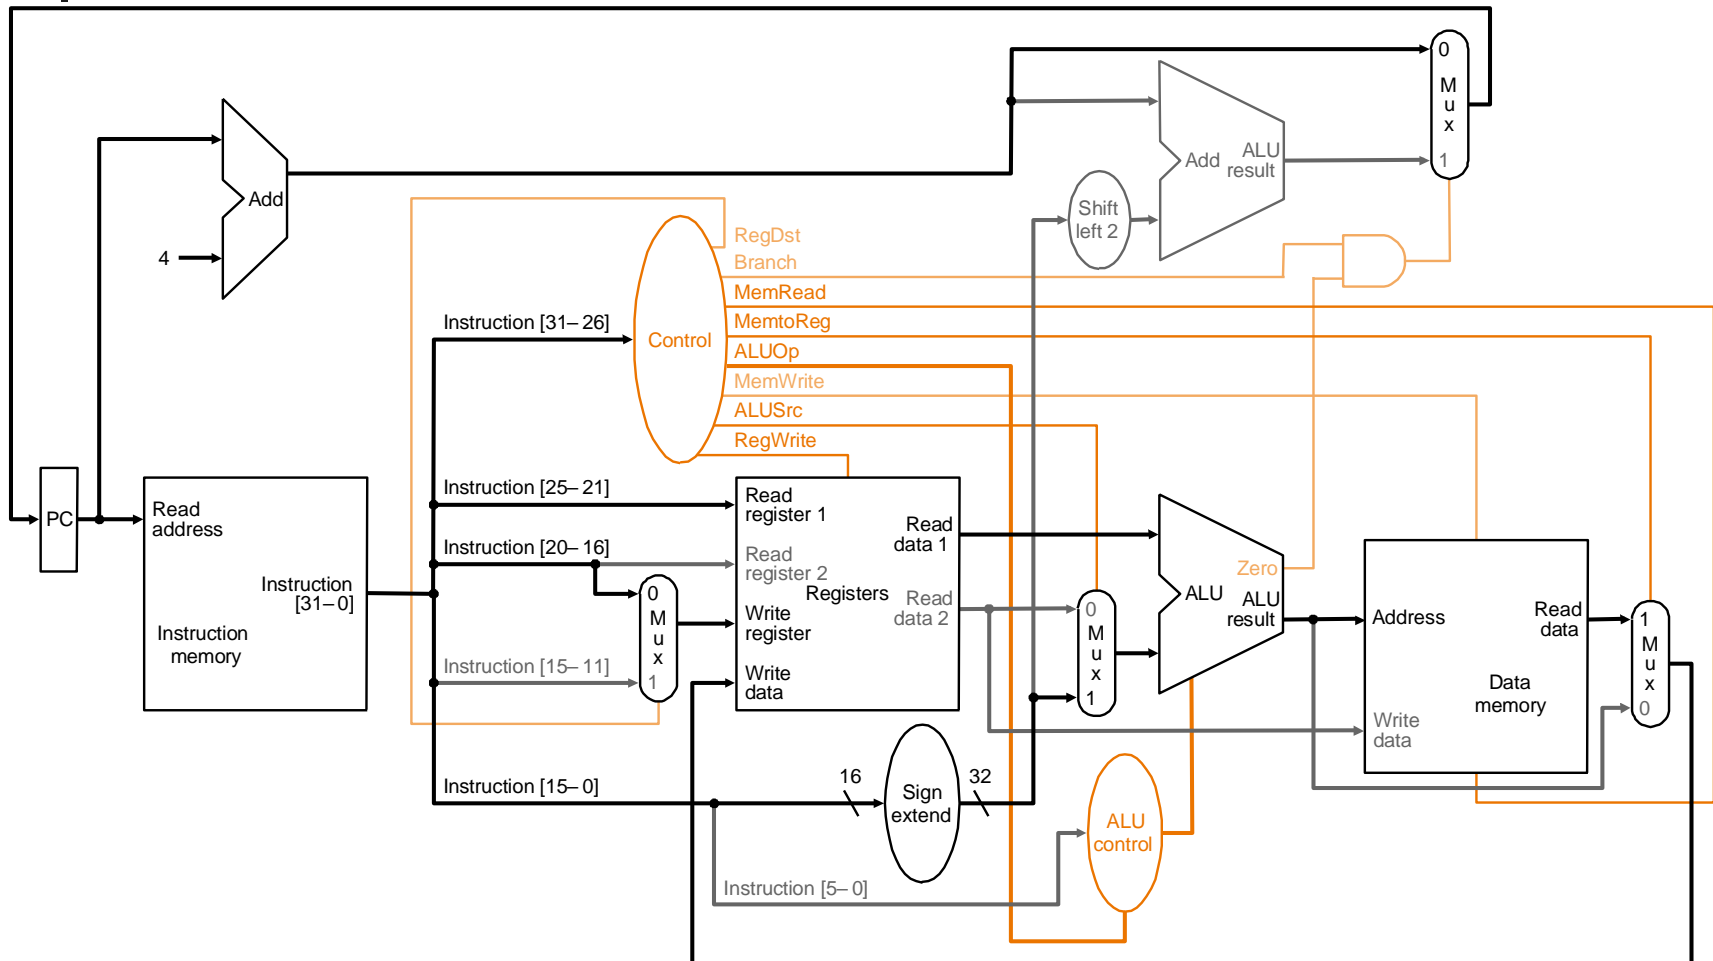
Load Instruction Steps

lw \$t1, offset(\$t2)

1. Fetch instruction and increment PC
2. Read base register from the register file: the base register (\$t2) is given by bits 25-21 of the instruction
3. ALU computes sum of value read from the register file and the sign-extended lower 16 bits (offset) of the instruction
4. The sum from the ALU is used as the address for the data memory
5. The data from the memory unit is written into the register file: the destination register (\$t1) is given by bits 20-16 of the instruction

Load Instruction

lw \$t1, offset(\$t2)





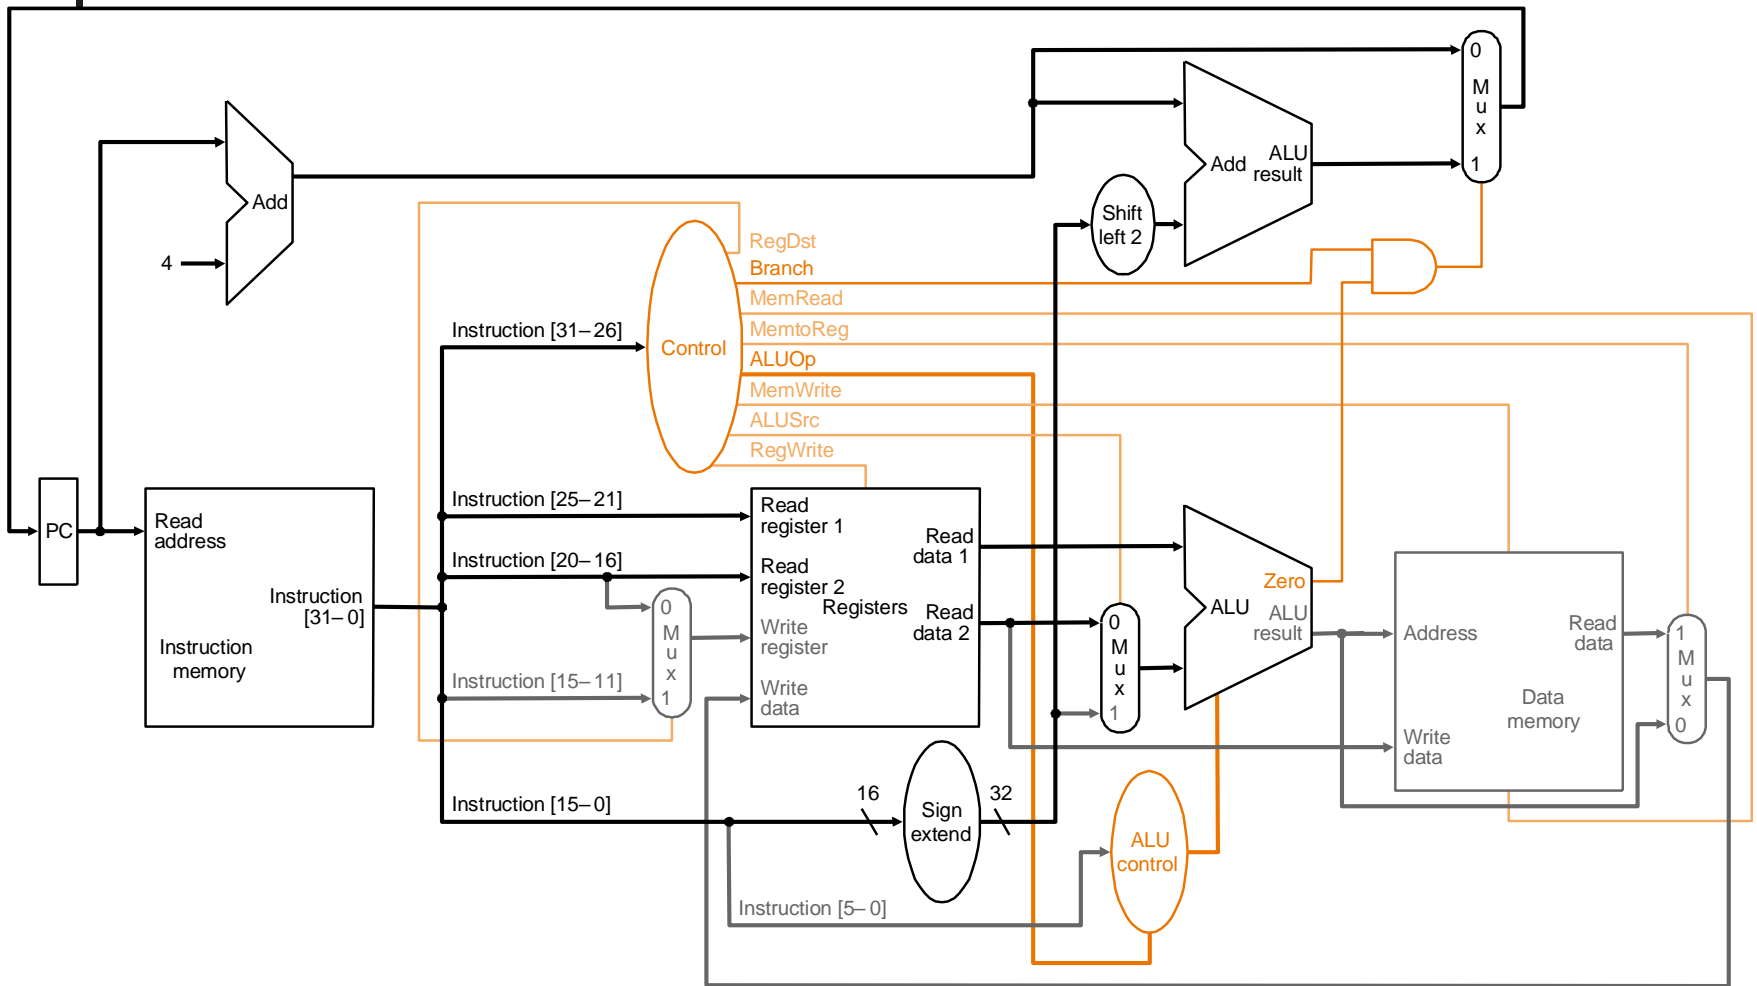
Branch Instruction Steps

beq \$t1, \$t2, offset

1. Fetch instruction and increment PC
2. Read two register (\$t1 and \$t2) from the register file
3. ALU performs a subtract on the data values from the register file; the value of PC+4 is added to the sign-extended lower 16 bits (offset) of the instruction shifted left by two to give the branch target address
4. The Zero result from the ALU is used to decide which adder result (from step 1 or 3) to store in the PC

Branch Instruction

beq \$t1, \$t2, offset

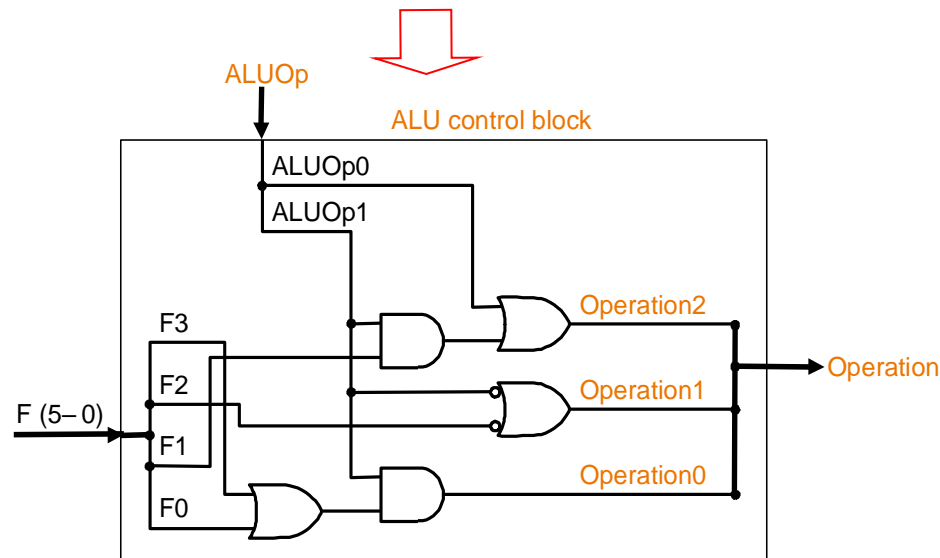


Implementation: ALU Control Block

ALUOp		Funct field						Operation
ALUOp1	ALUOp0	F5	F4	F3	F2	F1	F0	
0	0	X	X	X	X	X	X	010
0*	1	X	X	X	X	X	X	110
1	X	X	X	0	0	0	0	010
1	X	X	X	0	0	1	0	110
1	X	X	X	0	1	0	0	000
1	X	X	X	0	1	0	1	001
1	X	X	X	1	0	1	0	111

Truth table for ALU control bits

*Typo in text
Fig. 5.15: if it is X
then there is potential
conflict between
line 2 and lines 3-7!

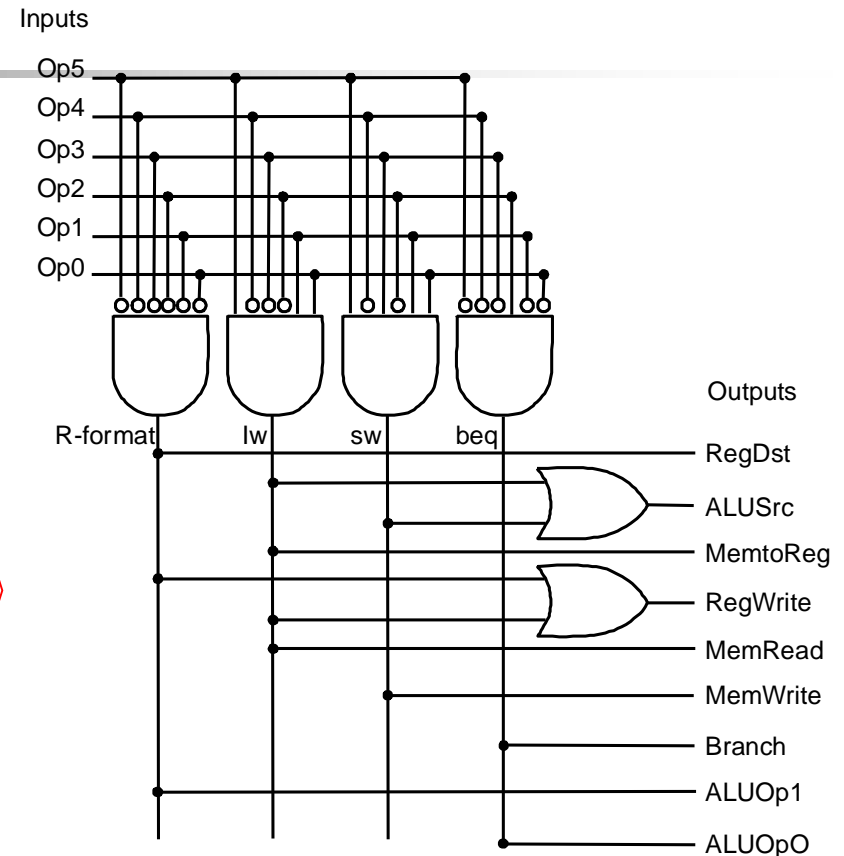


ALU control logic

Implementation: Main Control Block

	Signal name	R-format	lw	sw	beq
Inputs	Op5	0	1	1	0
	Op4	0	0	0	0
	Op3	0	0	1	0
	Op2	0	0	0	1
	Op1	0	1	1	0
	Op0	0	1	1	0
Outputs	RegDst	1	0	x	x
	ALUSrc	0	1	1	0
	MemtoReg	0	1	x	x
	RegWrite	1	1	0	0
	MemRead	0	1	0	0
	MemWrite	0	0	1	0
	Branch	0	0	0	1
	ALUOp1	1	0	0	0
	ALUOp2	0	0	0	1

Truth table for main control signals



Main control PLA (programmable logic array): principle underlying PLAs is that any logical expression can be written as a sum-of-products



Single-Cycle Design Problems

- Assuming fixed-period clock every instruction datapath uses one clock cycle implies:
 - $CPI = 1$
 - cycle time determined by length of the longest instruction path (load)
 - but several instructions could run in a shorter clock cycle: *waste of time*
 - consider if we have more complicated instructions like floating point!
 - resources used more than once in the same cycle need to be duplicated
 - *waste of hardware and chip area*

Example: Fixed-period clock vs. variable-period clock in a single-cycle implementation

- Consider a machine with an additional floating point unit. Assume functional unit delays as follows
 - *memory: 2 ns., ALU and adders: 2 ns., FPU add: 8 ns., FPU multiply: 16 ns., register file access (read or write): 1 ns.*
 - *multiplexors, control unit, PC accesses, sign extension, wires: no delay*
- Assume instruction mix as follows
 - all loads take same time and comprise 31%
 - all stores take same time and comprise 21%
 - R-format instructions comprise 27%
 - branches comprise 5%
 - jumps comprise 2%
 - FP adds and subtracts take the same time and totally comprise 7%
 - FP multiplies and divides take the same time and totally comprise 7%
- *Compare the performance of (a) a single-cycle implementation using a fixed-period clock with (b) one using a variable-period clock where each instruction executes in one clock cycle that is only as long as it needs to be (not really practical but pretend it's possible!)*



Solution

Instruction class	Instr. mem.	Register read	ALU oper.	Data mem.	Register write	FPU add/sub	FPU mul/div	Total time ns.
Load word	2	1	2	2	1			8
Store word	2	1	2	2				7
R-format	2	1	2	0	1			6
Branch	2	1	2					5
Jump	2							2
FP mul/div	2	1			1		16	20
FP add/sub	2	1			1	8		12

- Clock period for fixed-period clock = longest instruction time = 20 ns.
- Average clock period for variable-period clock = $8 \times 31\% + 7 \times 21\% + 6 \times 27\% + 5 \times 5\% + 2 \times 2\% + 20 \times 7\% + 12 \times 7\%$
= 7.0 ns.
- Therefore, $\text{performance}_{\text{var-period}} / \text{performance}_{\text{fixed-period}} = 20/7 = 2.9$




Fixing the problem with single-cycle designs

- One solution: a variable-period clock with different cycle times for each instruction class
 - *unfeasible*, as implementing a variable-speed clock is technically difficult
- Another solution:
 - use a smaller cycle time...
 - ...have different instructions take different numbers of cycles by breaking instructions into steps and fitting each step into one cycle
 - *feasible: multicyle approach!*



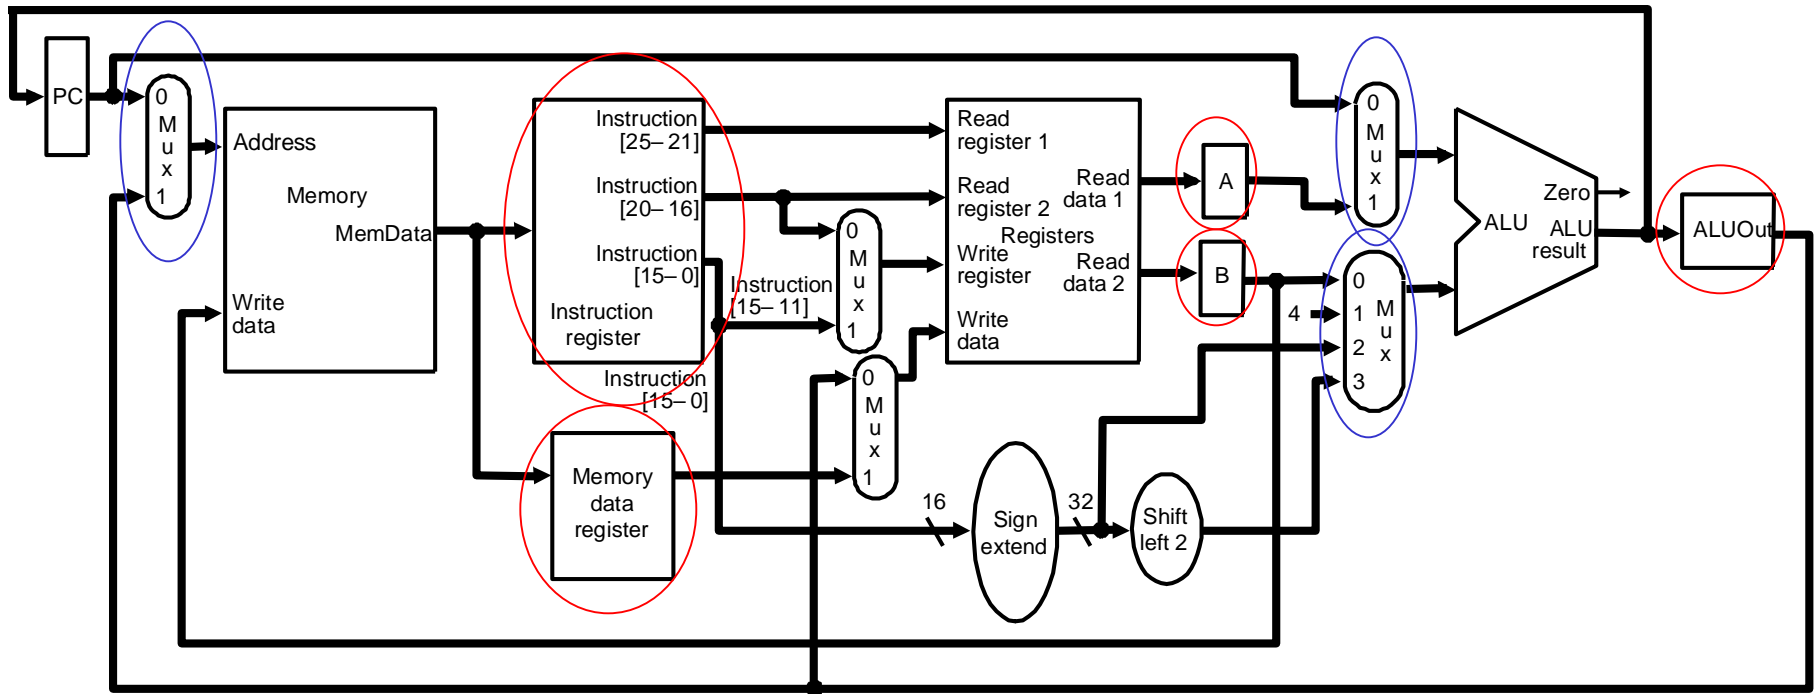
Multicycle Approach

- Break up the instructions into *steps*
 - each step takes one clock cycle
 - balance the amount of work to be done in each step/cycle so that they are about equal
 - restrict each cycle to use at most once each major functional unit so that such units do not have to be replicated
 - functional units can be shared between different cycles within one instruction
- Between steps/cycles
 - At the end of one cycle store data to be used in *later cycles of the same* instruction
 - need to introduce additional *internal* (programmer-invisible) registers for this purpose
 - Data to be used in *later instructions* are stored in programmer-visible state elements: the register file, PC, memory

- single memory for data and instructions
 - single ALU, no extra adders
 - extra registers to hold data between clock cycles
- 
- The diagram shows a partial view of a processor block. A box labeled 'PC' (Program Counter) is visible, with an arrow pointing to it from the left. Another arrow points from the 'PC' box to the right, indicating the flow of the next instruction.



Multicycle Datapath



Basic multicycle MIPS datapath handles R-type instructions and load/stores:
new internal register in **red ovals**, new multiplexors in **blue ovals**



Breaking instructions into steps

- Our goal is to break up the instructions into *steps* so that
 - each step takes one clock cycle
 - the amount of work to be done in each step/cycle is about equal
 - each cycle uses at most once each major functional unit so that such units do not have to be replicated
 - functional units can be shared between different cycles within one instruction
- Data at end of one cycle to be used in next *must be stored* !!



Breaking instructions into steps

- We break instructions into the following *potential* execution steps – not all instructions require all the steps – each step takes one clock cycle
 1. Instruction fetch and PC increment (**IF**)
 2. Instruction decode and register fetch (**ID**)
 3. Execution, memory address computation, or branch completion (**EX**)
 4. Memory access or R-type instruction completion (**MEM**)
 5. Memory read completion (**WB**)
- Each MIPS instruction takes from 3 – 5 cycles (steps)



Step 1: Instruction Fetch & PC Increment (**IF**)

- Use PC to get instruction and put it in the instruction register. Increment the PC by 4 and put the result back in the PC.
- Can be described succinctly using *RTL (Register-Transfer Language)*:

```
IR = Memory[PC];
```

```
PC = PC + 4;
```

Step 2: Instruction Decode and Register Fetch (**ID**)

- Read registers rs and rt in case we need them.
Compute the branch address in case the instruction is a branch.
- RTL:
 $A = \text{Reg}[\text{IR}[25-21]];$
 $B = \text{Reg}[\text{IR}[20-16]];$
 $\text{ALUOut} = \text{PC} + (\text{sign-extend}(\text{IR}[15-0]) \ll 2);$

Step 3: Execution, Address Computation or Branch Completion

(EX)

- ALU performs one of four functions *depending* on instruction type
 - memory reference:
`ALUOut = A + sign-extend(IR[15-0]);`
 - R-type:
`ALUOut = A op B;`
 - branch (instruction *completes*):
`if (A==B) PC = ALUOut;`
 - jump (instruction *completes*):
`PC = PC[31-28] || (IR(25-0) << 2)`



Step 4: Memory access or R-type Instruction Completion (MEM)

- Again *depending* on instruction type:
- Loads and stores access memory
 - load
 - MDR = Memory[ALUOut];
 - store (instruction *completes*)
 - Memory[ALUOut] = B;
- R-type (instructions *completes*)
 - Reg[IR[15-11]] = ALUOut;



Step 5: Memory Read Completion (**WB**)

- Again *depending* on instruction type:
- Load writes back (instruction *completes*)

`Reg[IR[20-16]] = MDR;`

Important: There is no reason from a datapath (or control) point of view that Step 5 cannot be eliminated by performing

`Reg[IR[20-16]] = Memory[ALUOut];`

for loads in Step 4. This would eliminate the MDR as well.

The reason this is not done is that, to keep steps balanced in length, the design restriction is to allow each step to contain *at most* one ALU operation, or one register access, or one memory access.

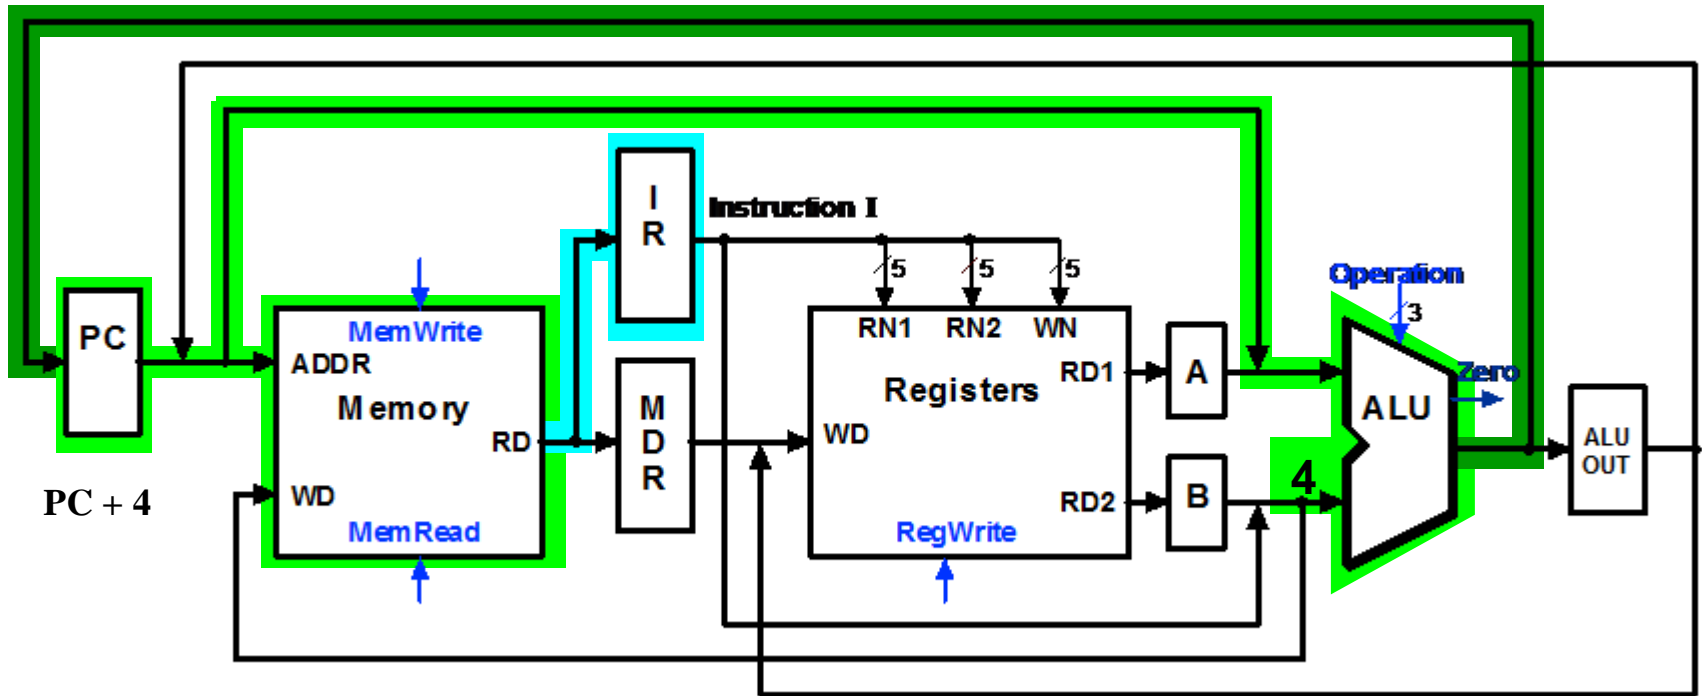


Summary of Instruction Execution

Step	Step name	Action for R-type instructions	Action for memory-reference instructions	Action for branches	Action for jumps
1: IF	Instruction fetch	$IR = \text{Memory}[PC]$ $PC = PC + 4$			
2: ID	Instruction decode/register fetch	$A = \text{Reg}[IR[25-21]]$ $B = \text{Reg}[IR[20-16]]$ $ALUOut = PC + (\text{sign-extend}(IR[15-0]) \ll 2)$			
3: EX	Execution, address computation, branch/jump completion	$ALUOut = A \text{ op } B$	$ALUOut = A + \text{sign-extend}(IR[15-0])$	if $(A == B)$ then $PC = ALUOut$	$PC = PC[31-28] \parallel (IR[25-0] \ll 2)$
4: MEM	Memory access or R-type completion	$\text{Reg}[IR[15-11]] = ALUOut$	Load: $MDR = \text{Memory}[ALUOut]$ or Store: $\text{Memory}[ALUOut] = B$		
5: WB	Memory read completion		Load: $\text{Reg}[IR[20-16]] = MDR$		

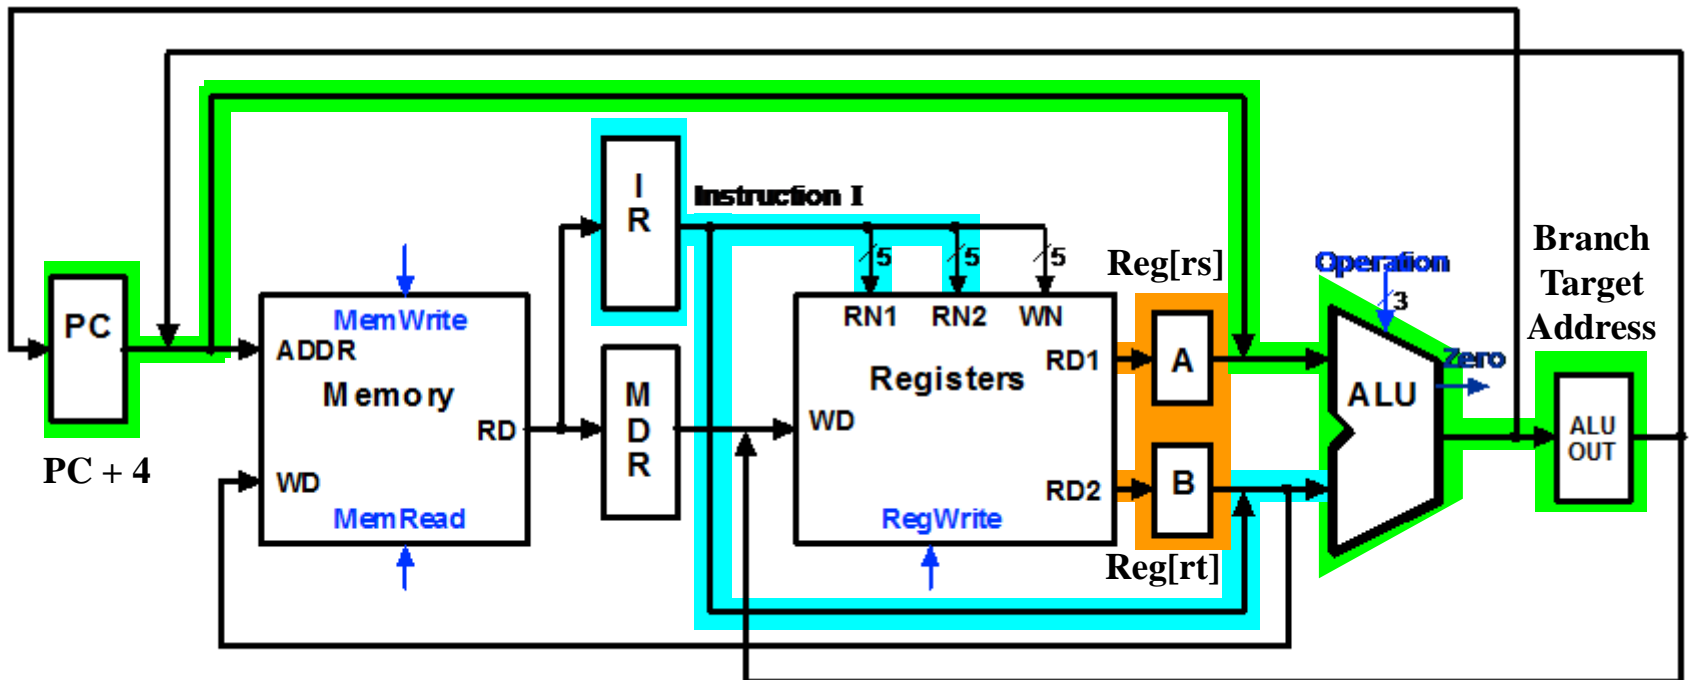
Multicycle Execution Step (1): Instruction Fetch

$IR = \text{Memory}[PC];$
 $PC = PC + 4;$



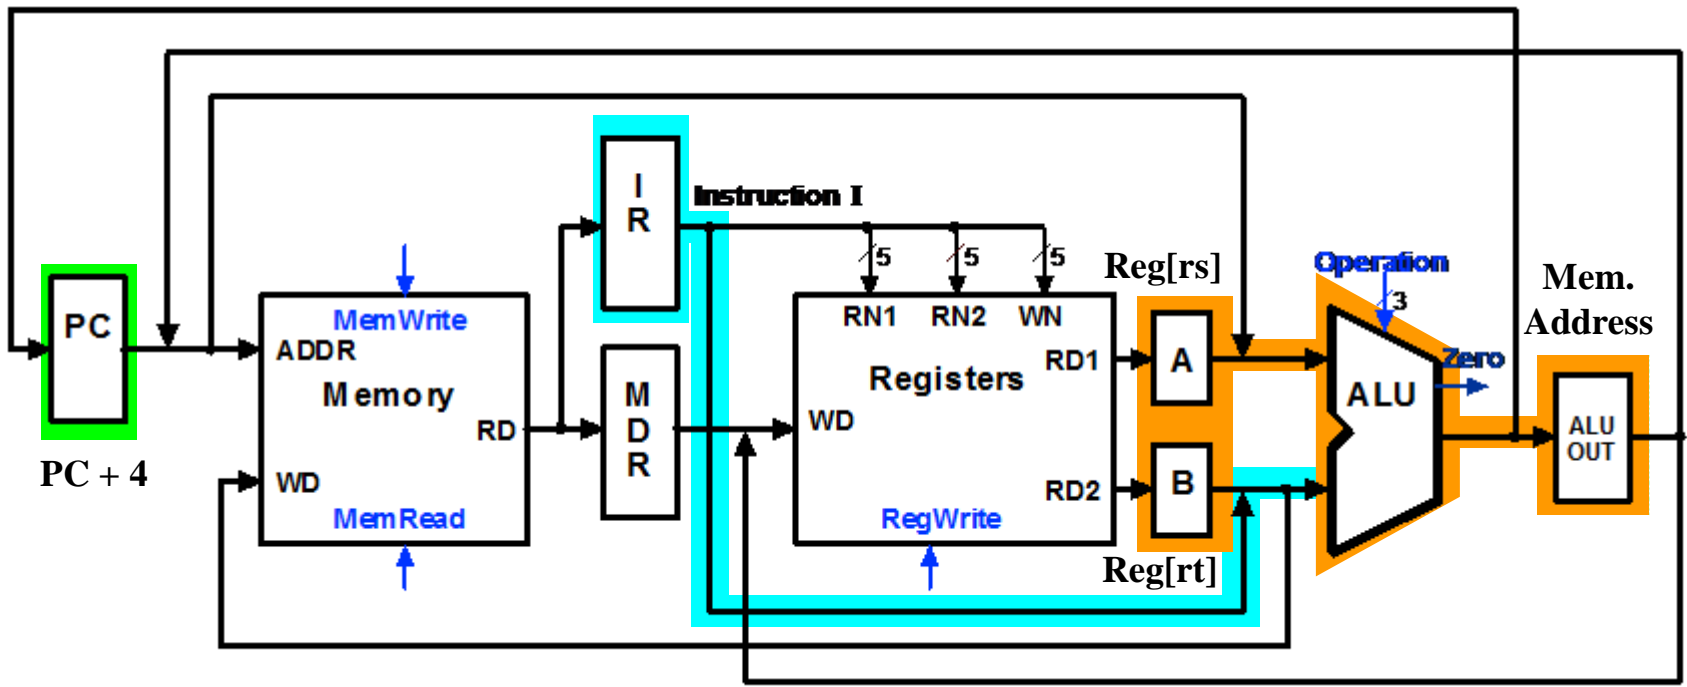
Multicycle Execution Step (2): Instruction Decode & Register Fetch

```
A = Reg[IR[25-21]];           (A = Reg[rs])  
B = Reg[IR[20-15]];          (B = Reg[rt])  
ALUOut = (PC + sign-extend(IR[15-0]) << 2)
```



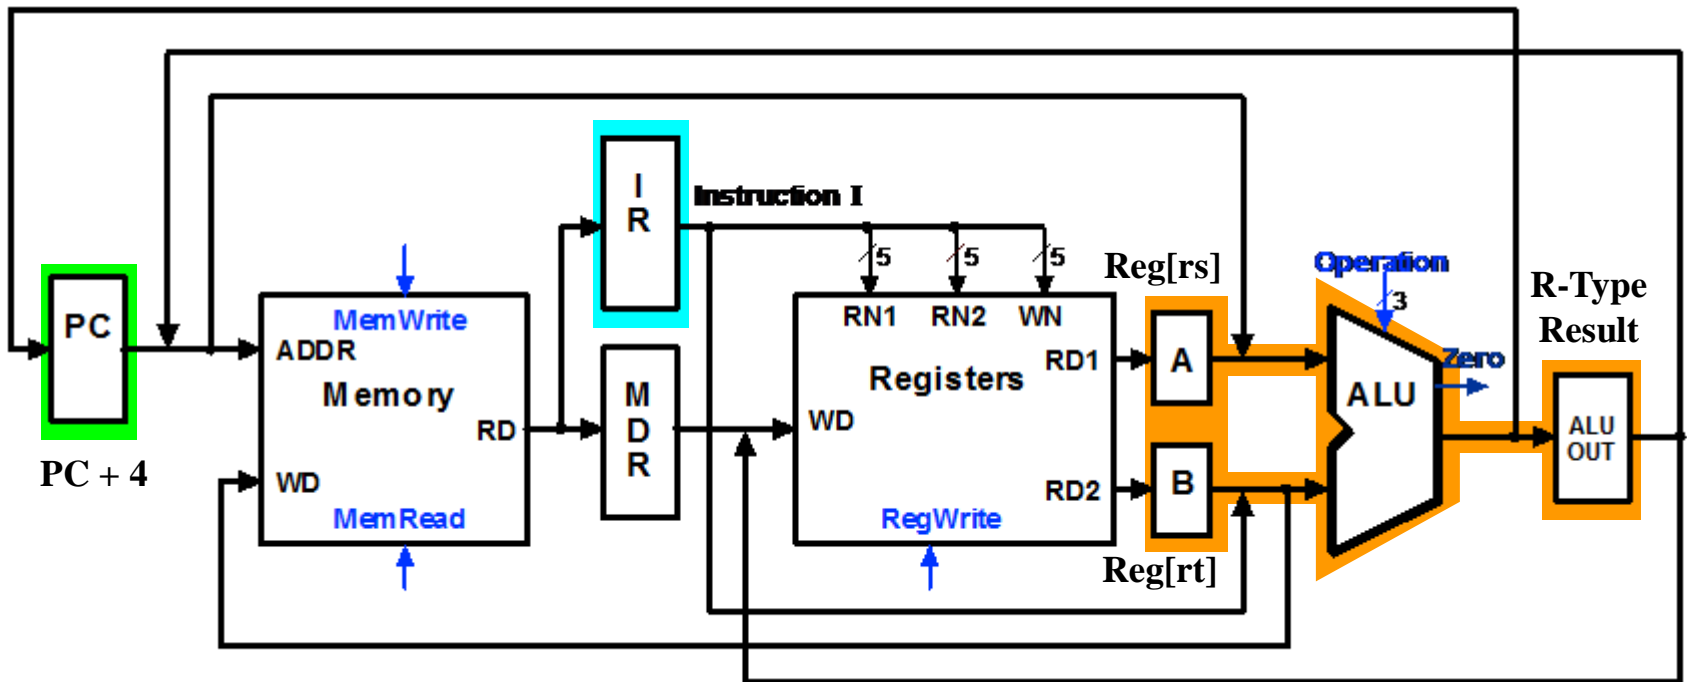
Multicycle Execution Step (3): Memory Reference Instructions

$ALUOut = A + \text{sign-extend}(IR[15-0]);$



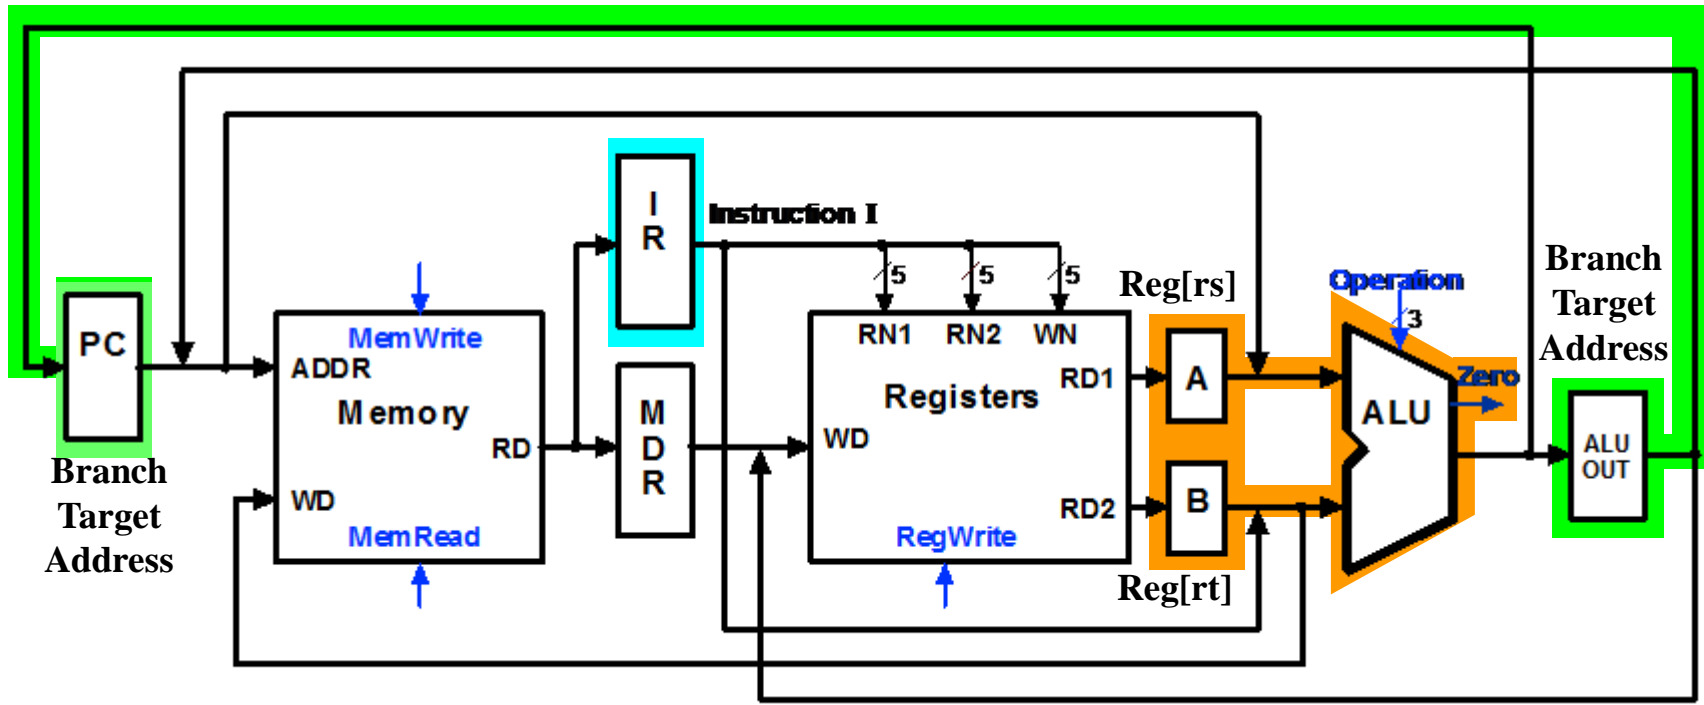
Multicycle Execution Step (3): ALU Instruction (R-Type)

$$\text{ALUOut} = A \text{ op } B$$



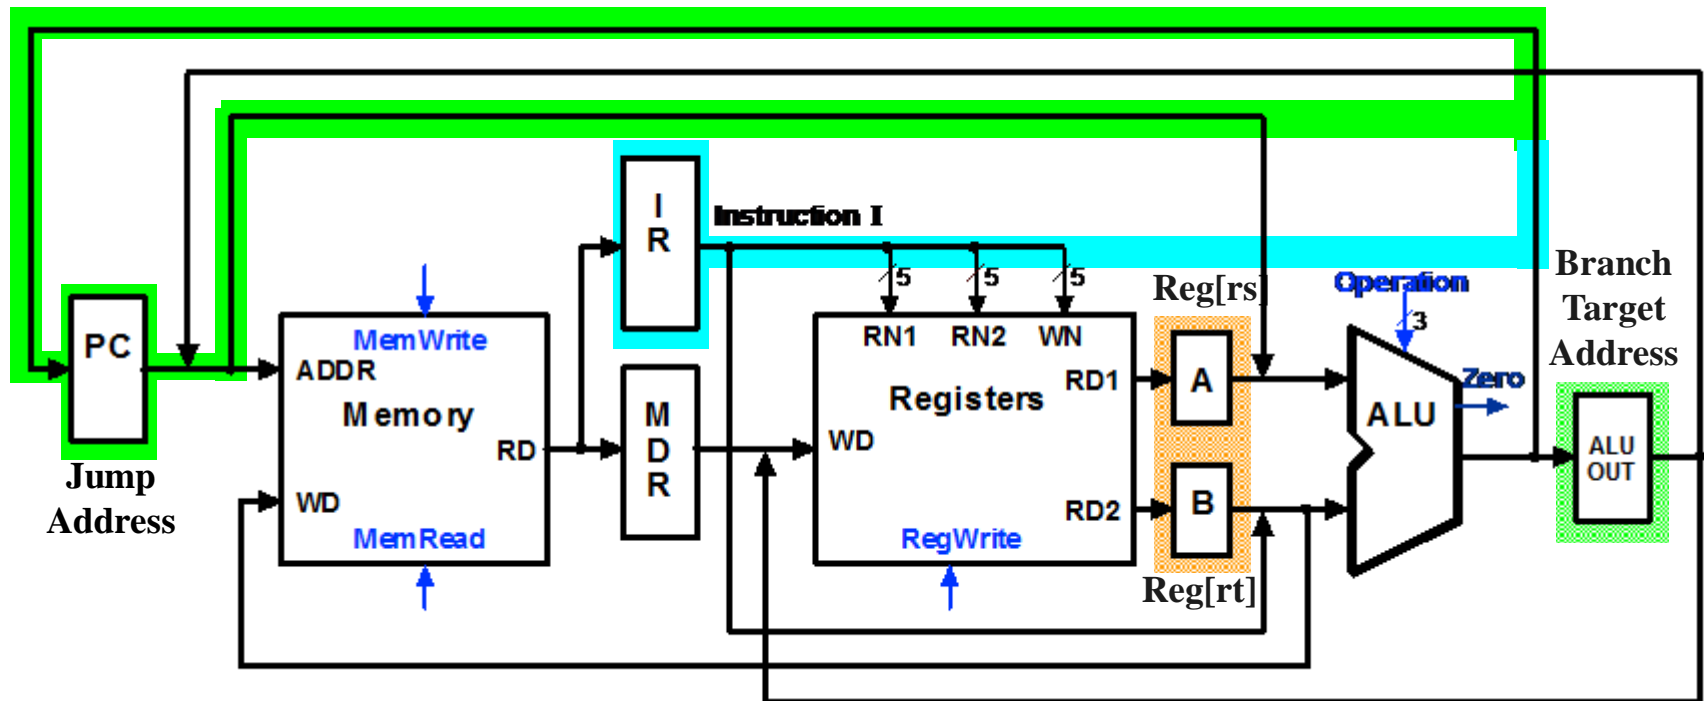
Multicycle Execution Step (3): Branch Instructions

```
if (A == B) PC = ALUOut;
```



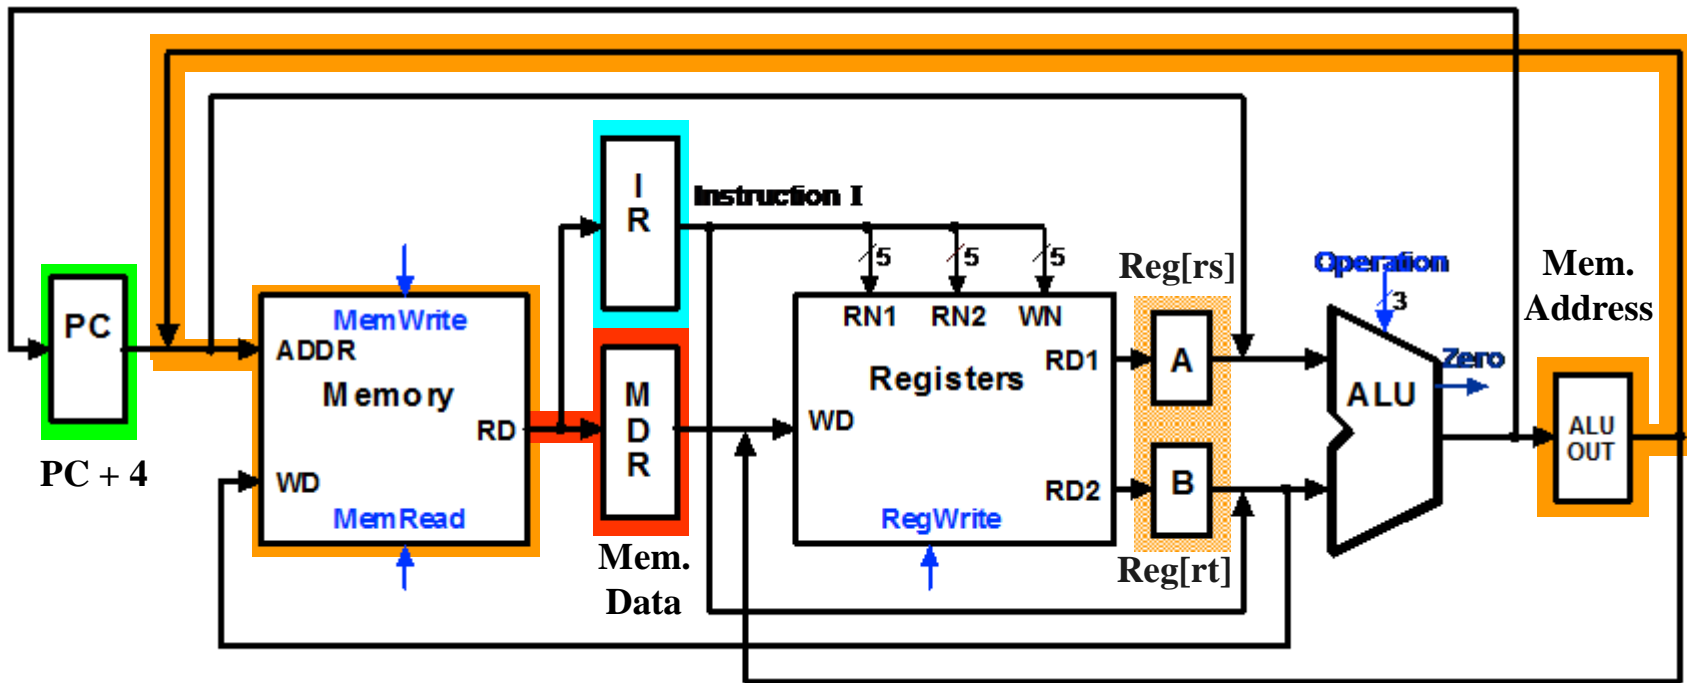
Multicycle Execution Step (3): Jump Instruction

$PC = PC[31-28] \text{ concat } (IR[25-0] \ll 2)$



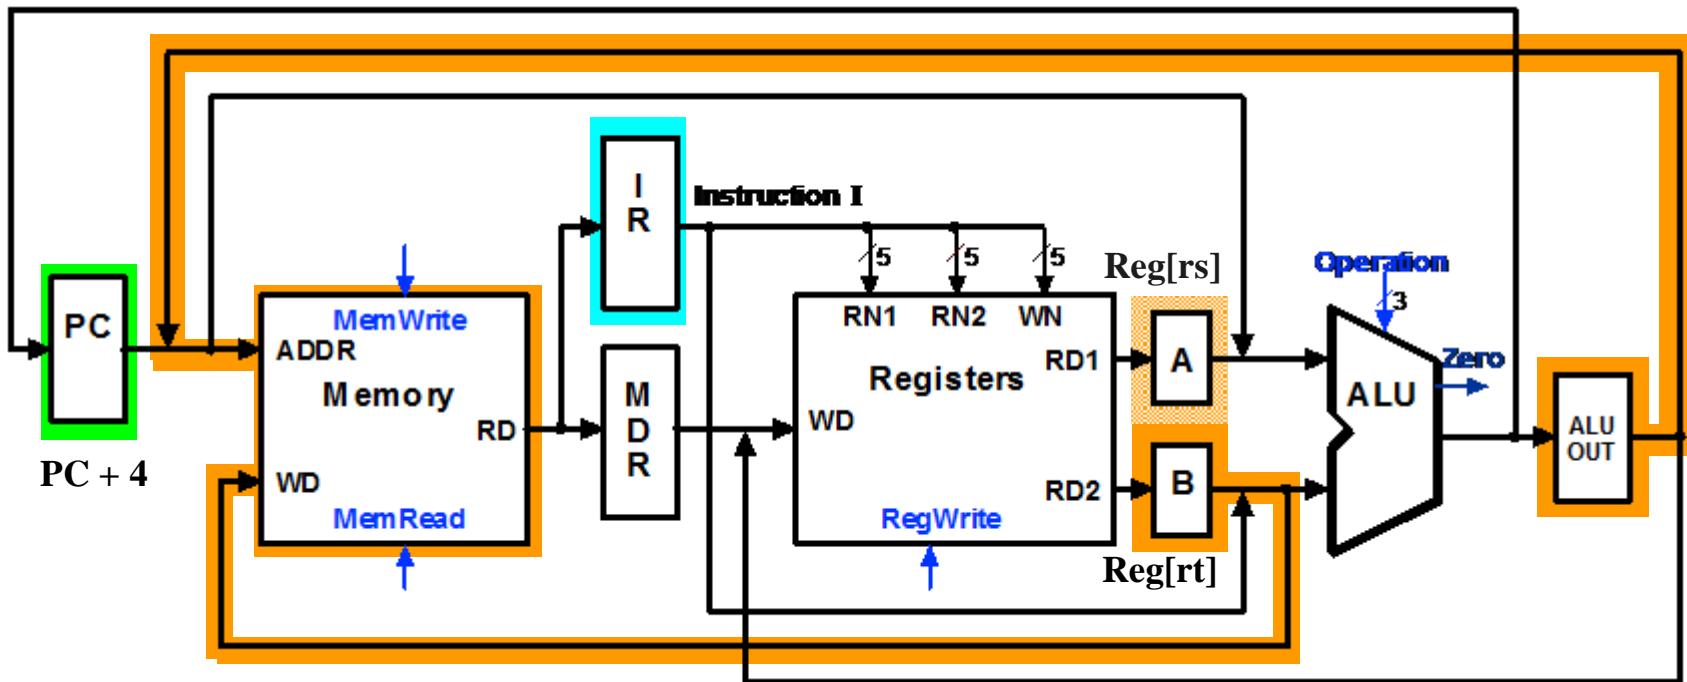
Multicycle Execution Step (4): Memory Access - Read (1_w)

$MDR = Memory[ALUOut];$



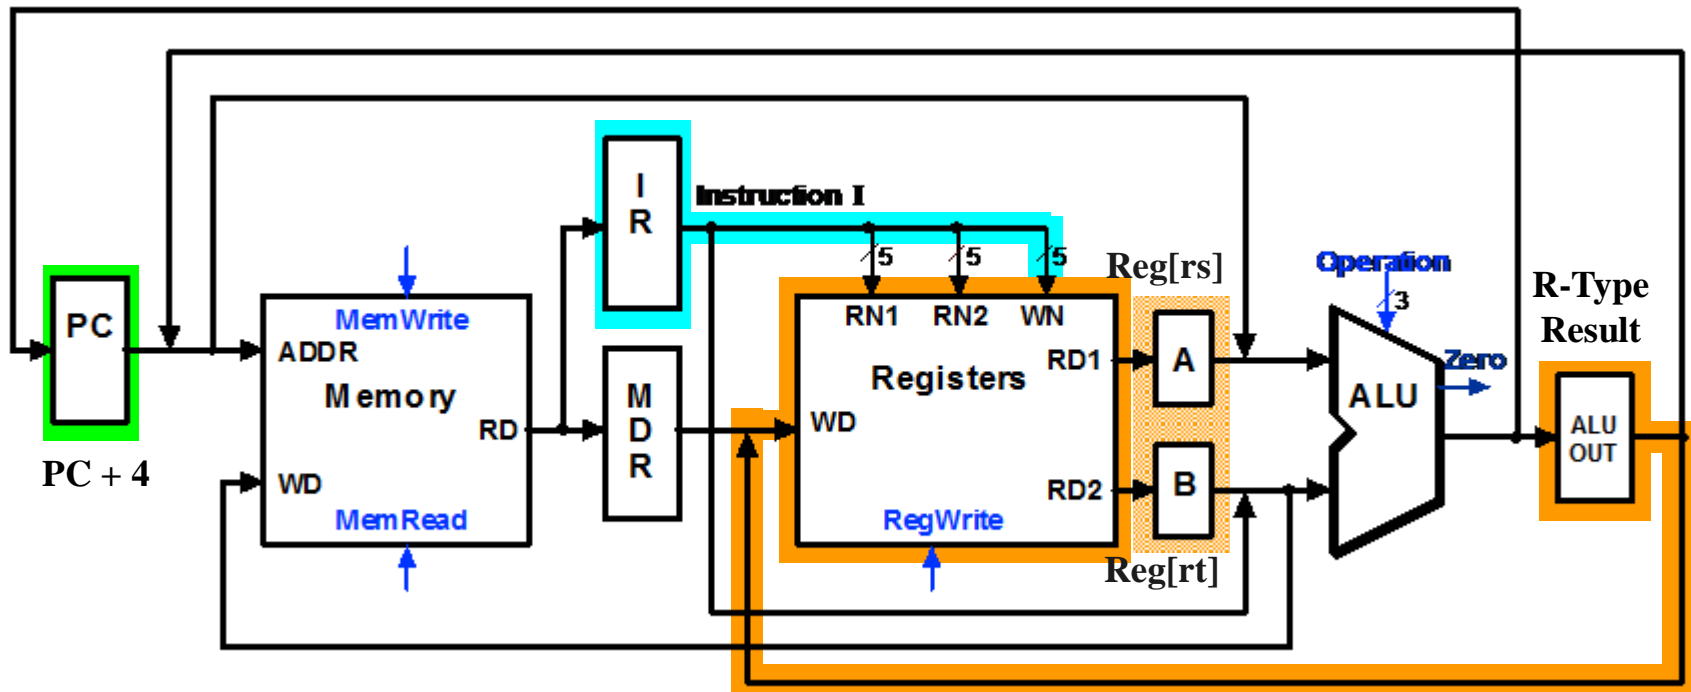
Multicycle Execution Step (4): Memory Access - Write (S_W)

Memory[ALUOut] = B;



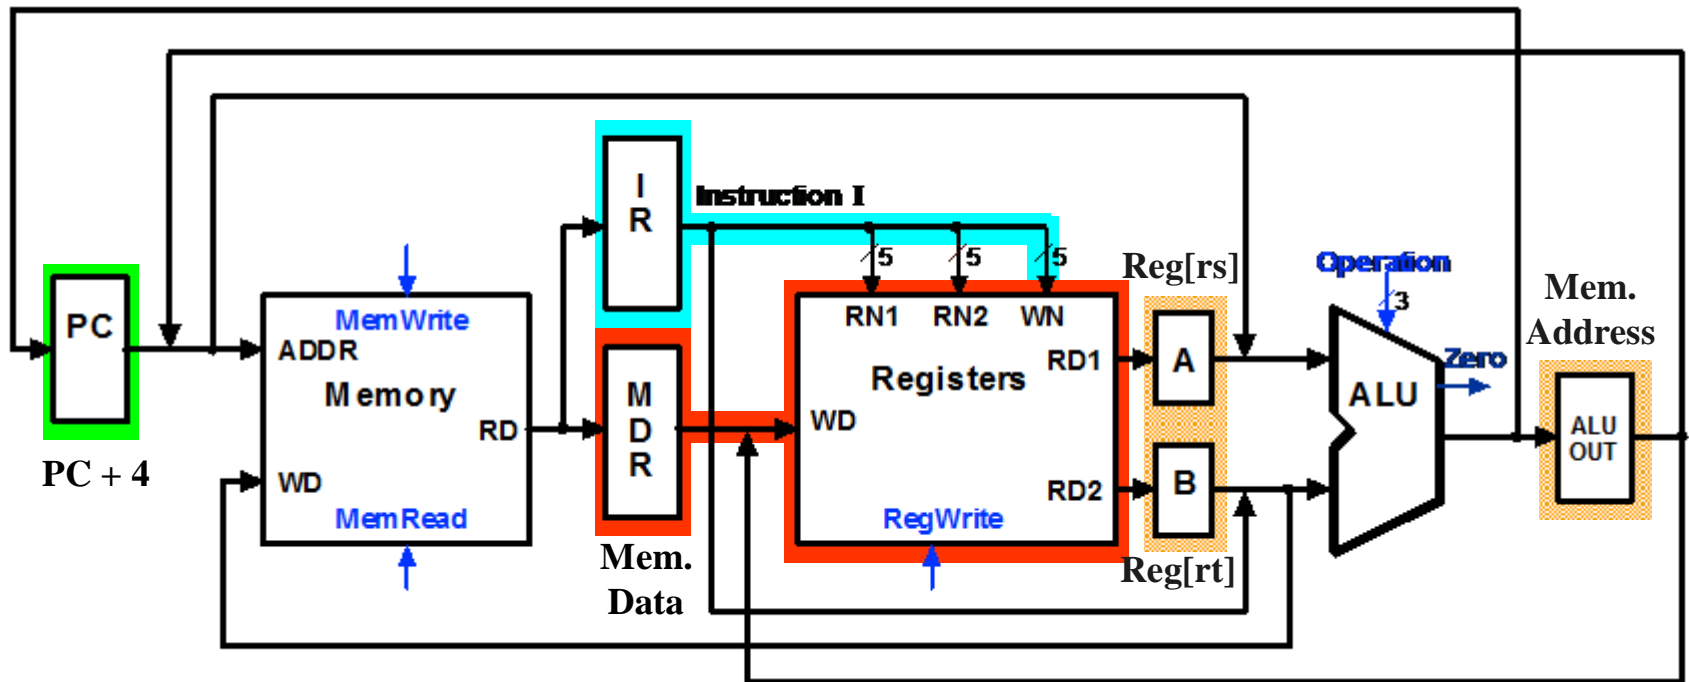
Multicycle Execution Step (4): ALU Instruction (R-Type)

$\text{Reg}[\text{IR}[15:11]] = \text{ALUOUT}$

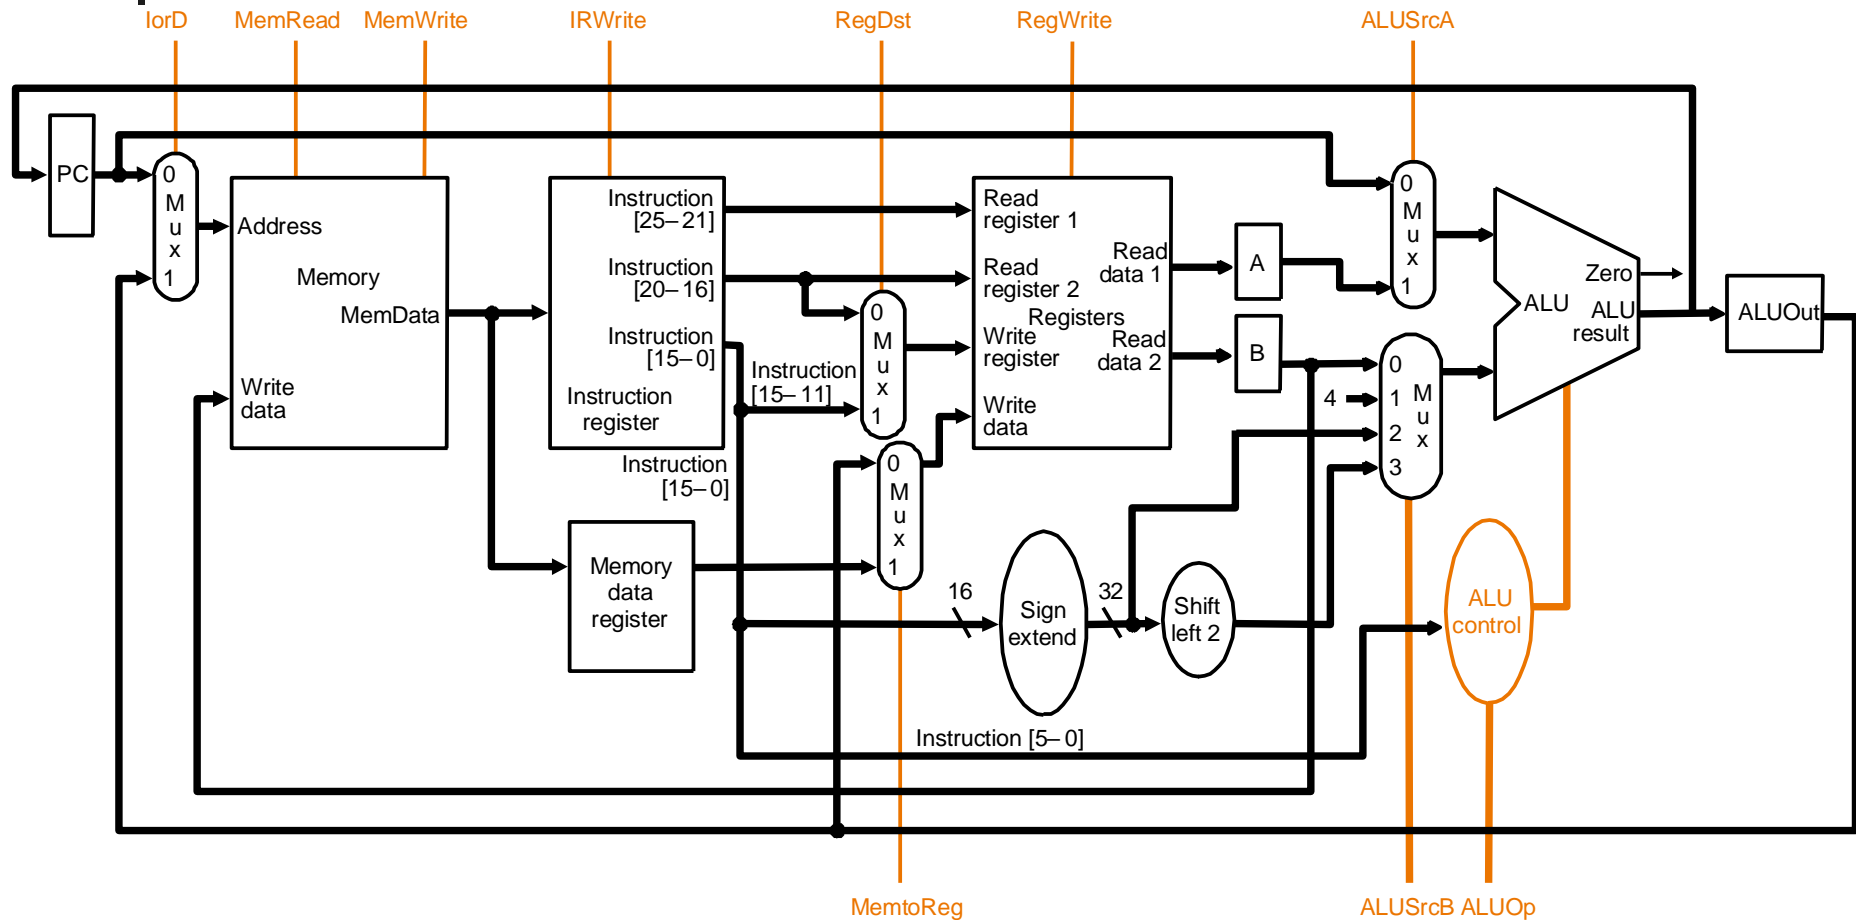


Multicycle Execution Step (5): Memory Read Completion (1_w)

`Reg[IR[20-16]] = MDR;`

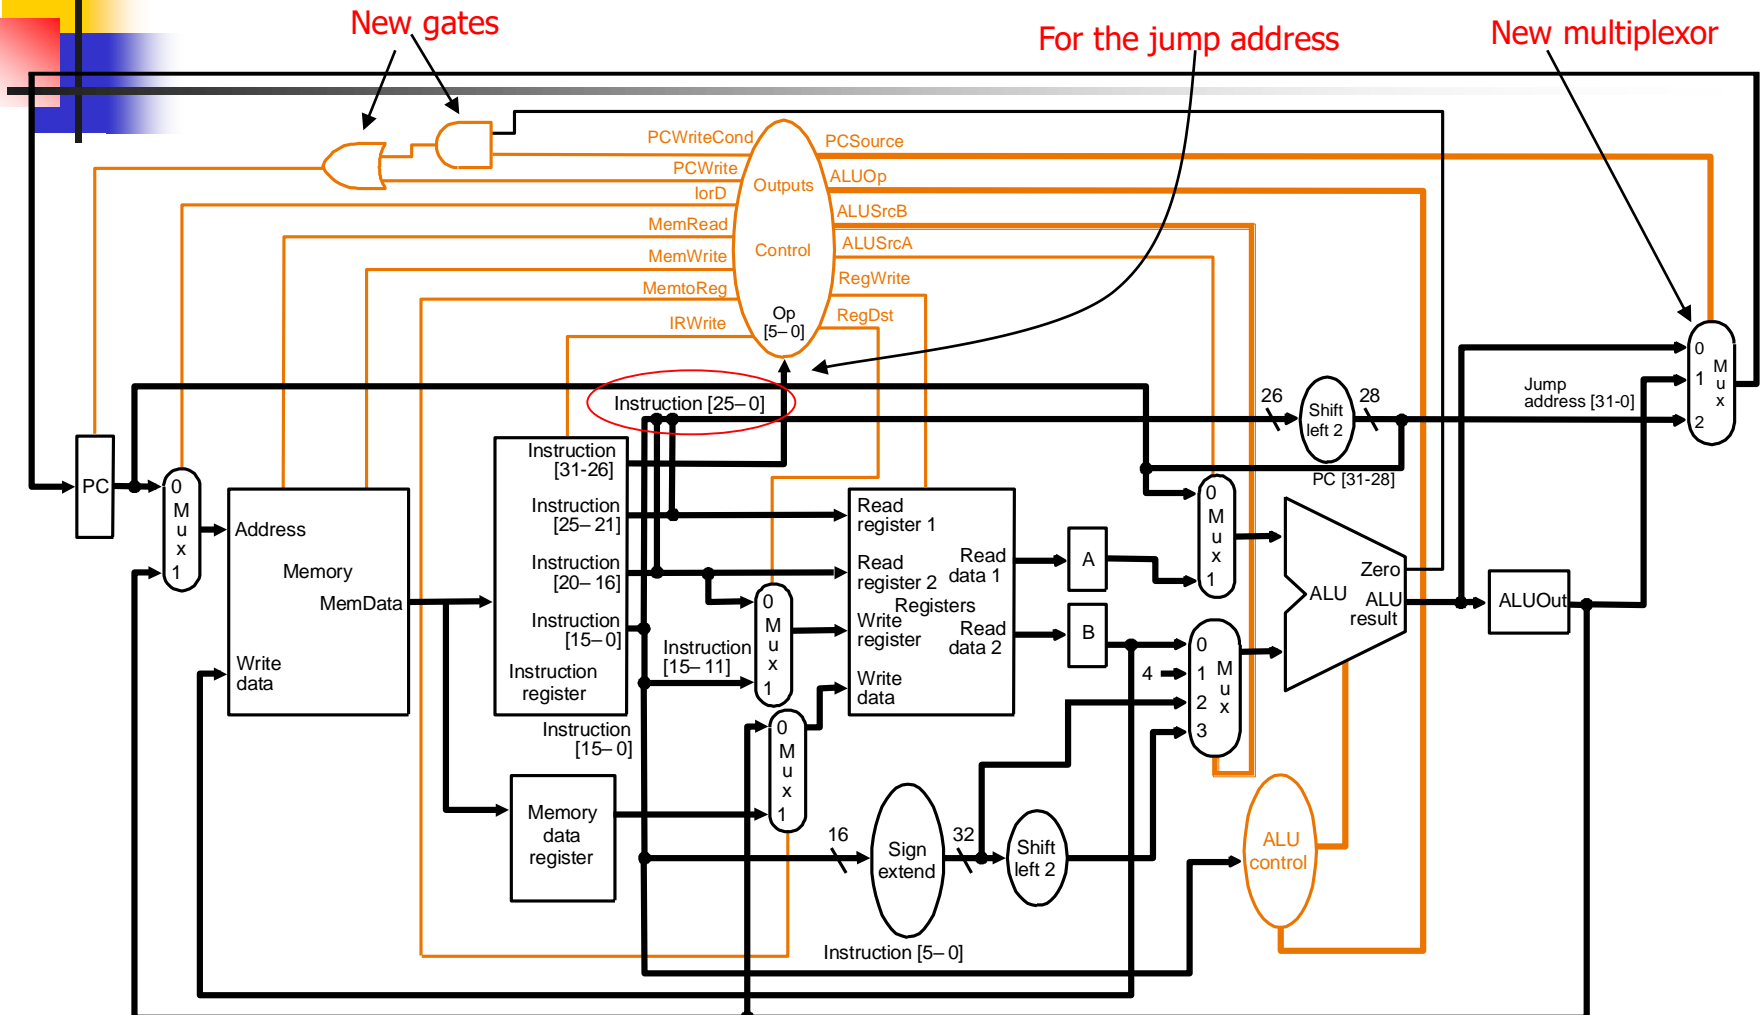


Multicycle Datapath with Control I

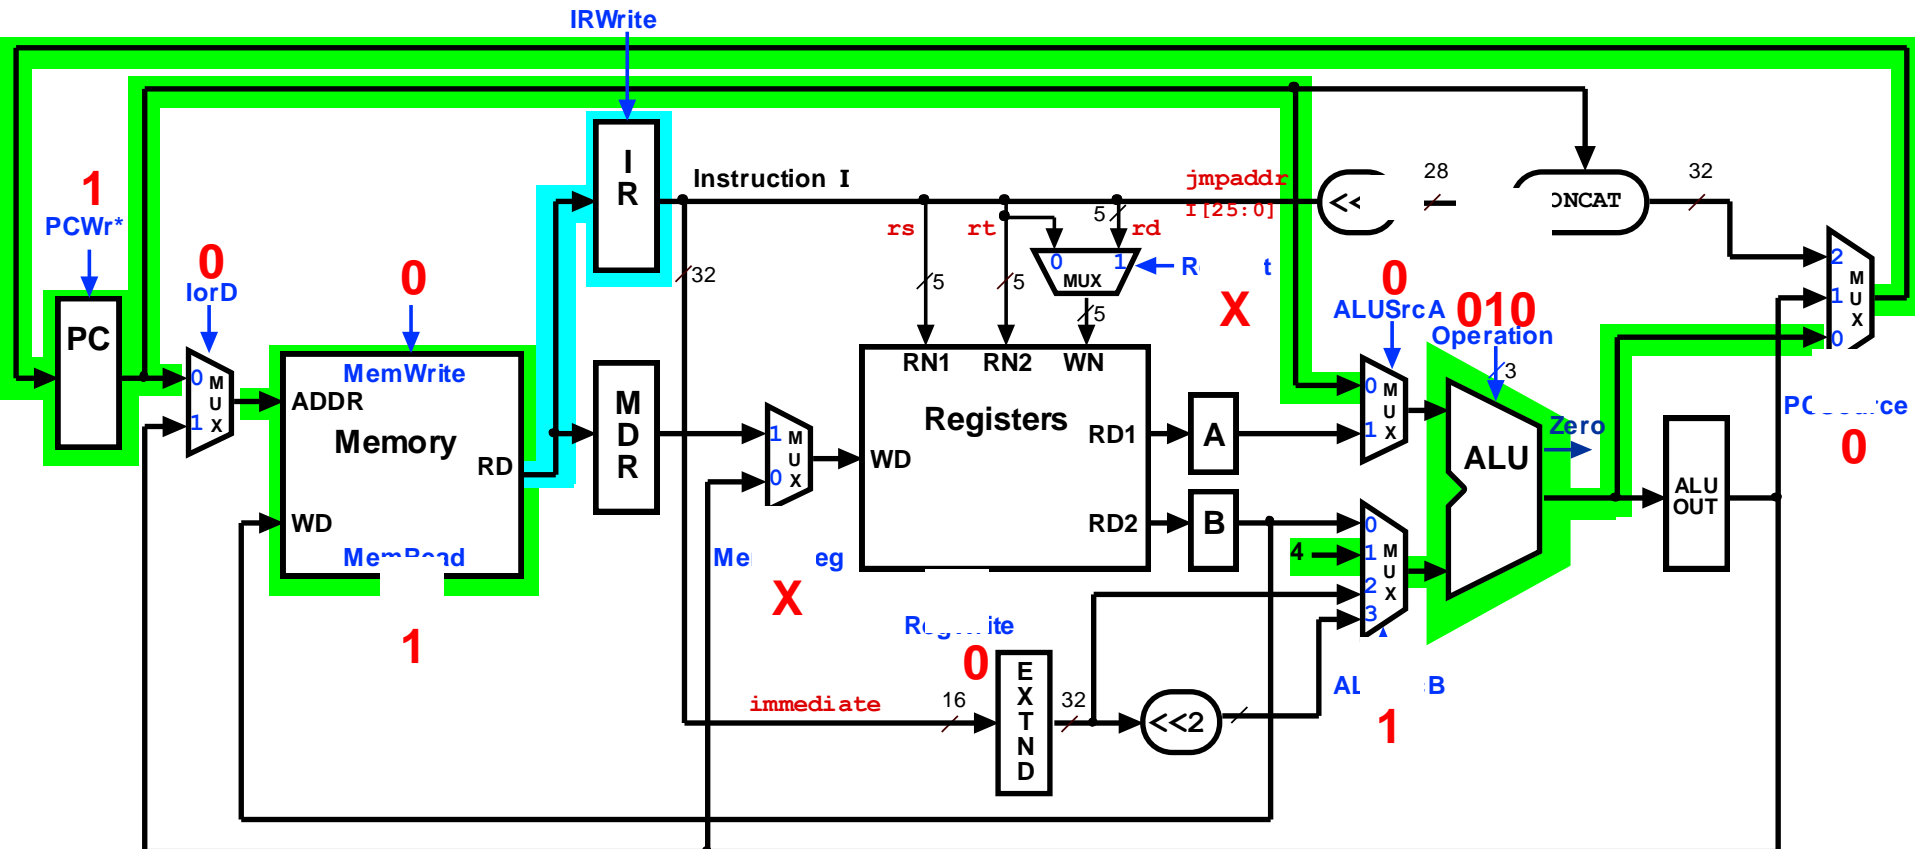


... with control lines and the ALU control block added – *not all* control lines are shown

Multicycle Datapath with Control II



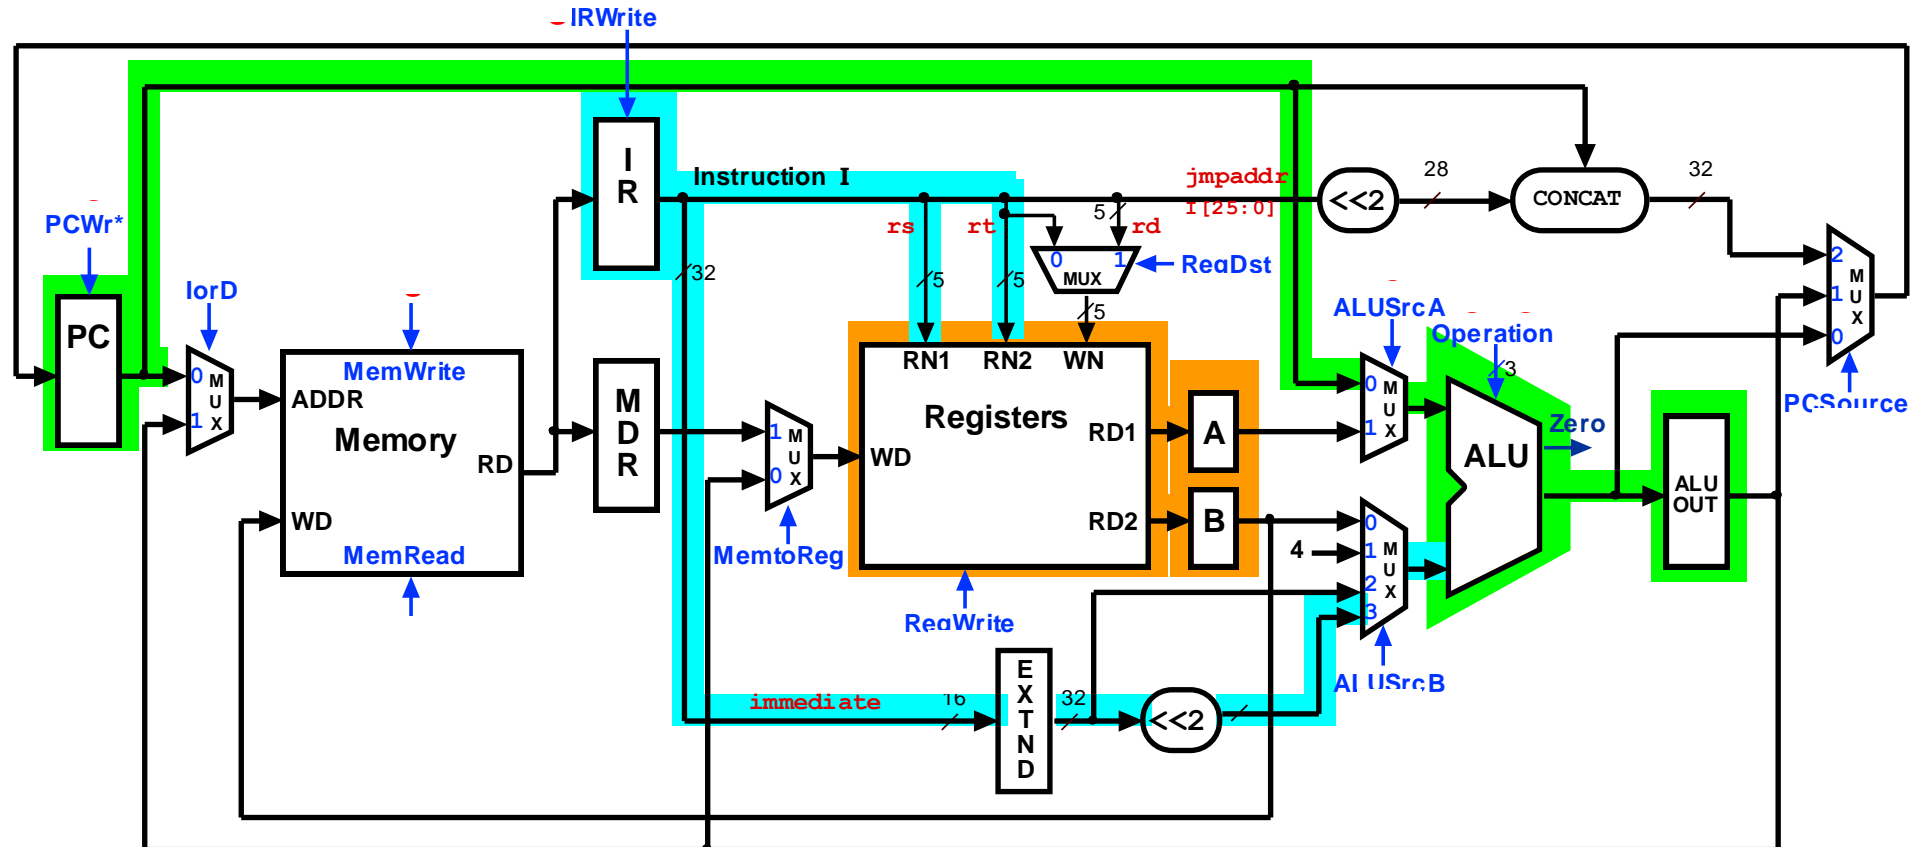
Complete multicycle MIPS datapath (with branch and jump capability) and showing the main control block and all control lines

$$\text{PC} = \text{PC} + 4$$


Multicycle Control Step (2): Instruction Decode & Register Fetch

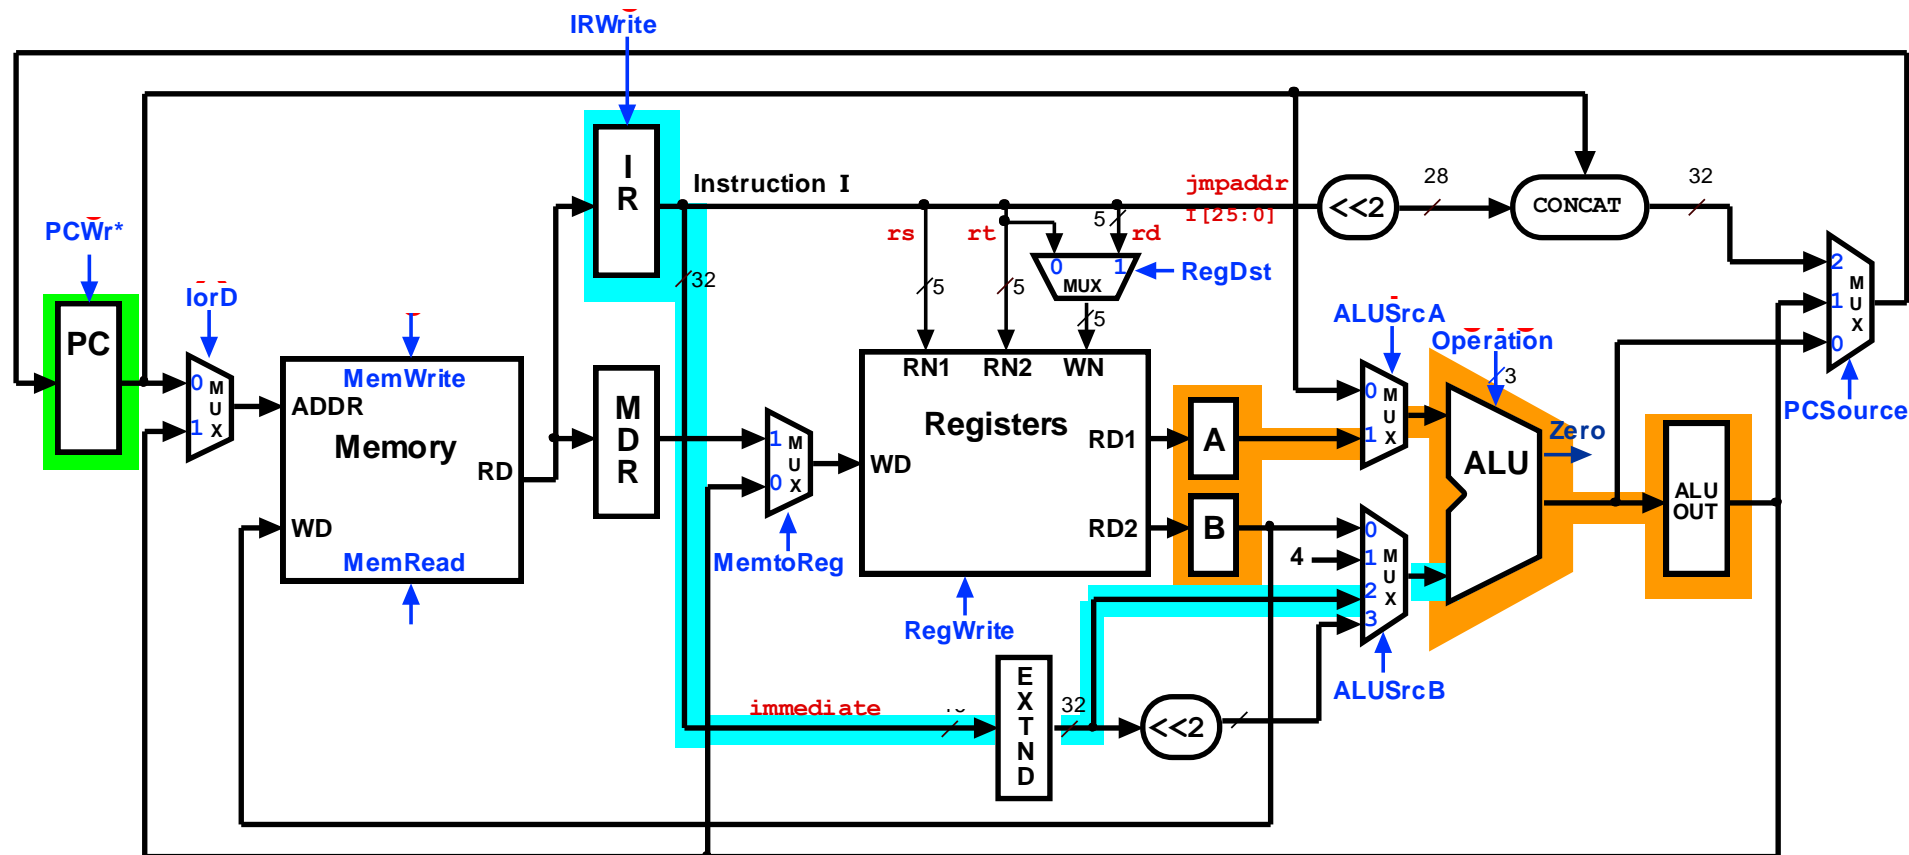
```

A = Reg[IR[25-21]];          (A = Reg[rs])
B = Reg[IR[20-15]];          (B = Reg[rt])
ALUOut = (PC + sign-extend(IR[15-0]) << 2);
    
```



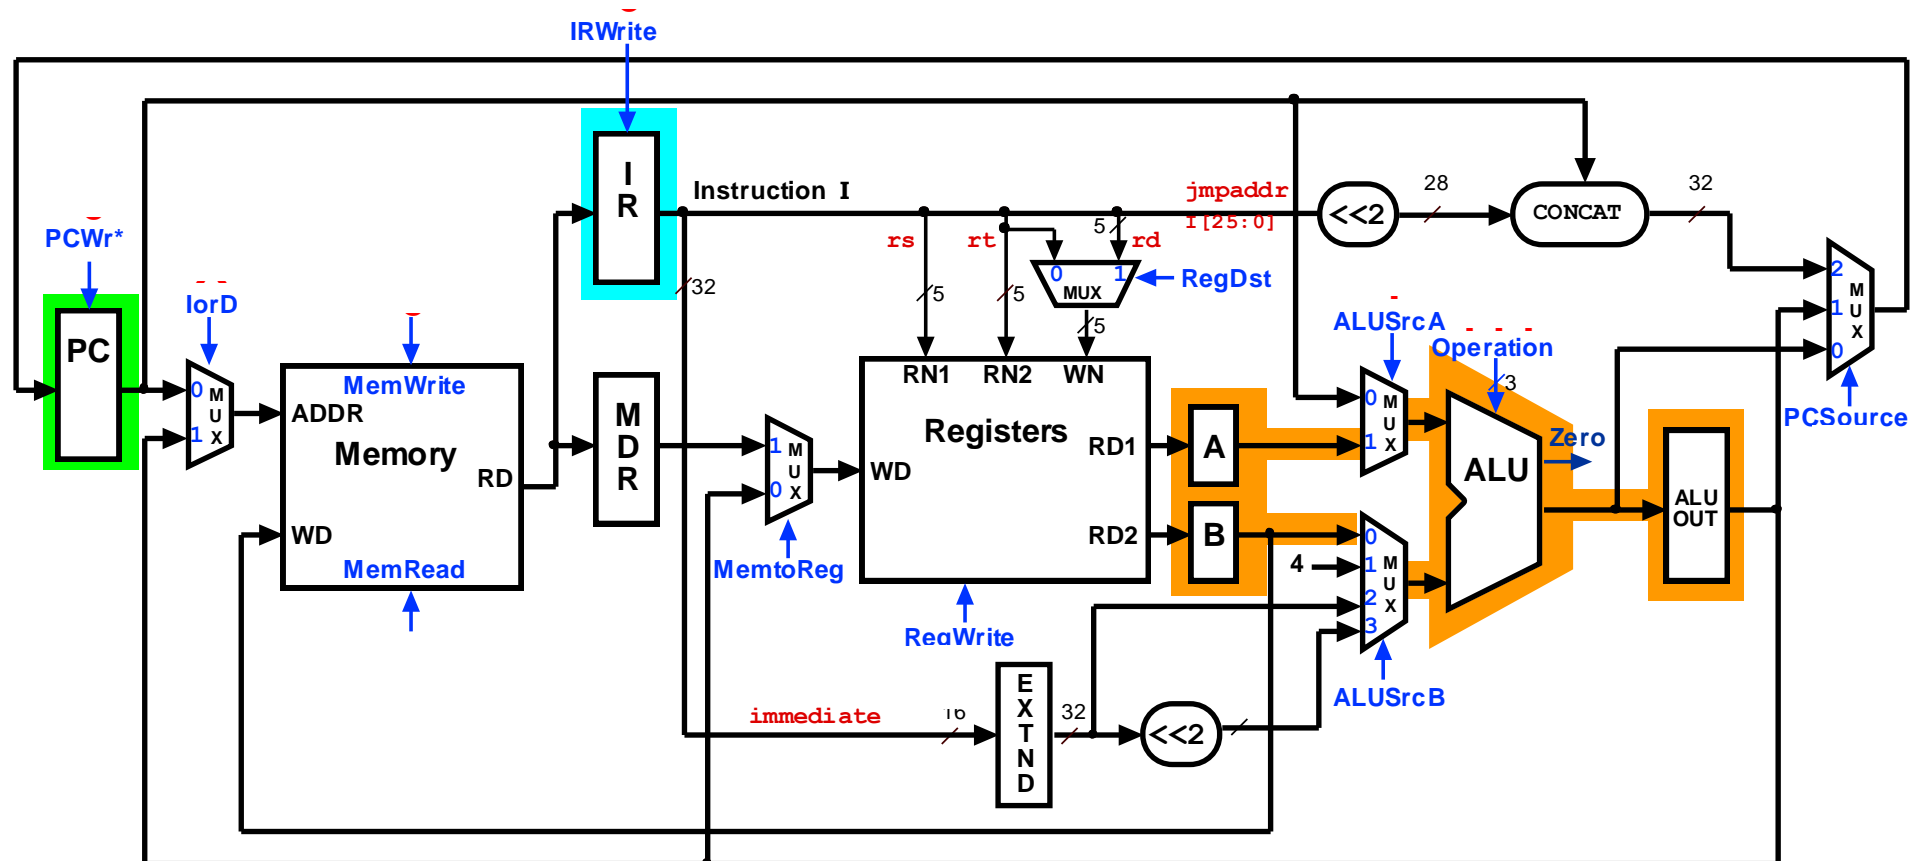
Multicycle Control Step (3): Memory Reference Instructions

$ALUOut = A \cdot \text{sign-extend}(IR[15:0]);$



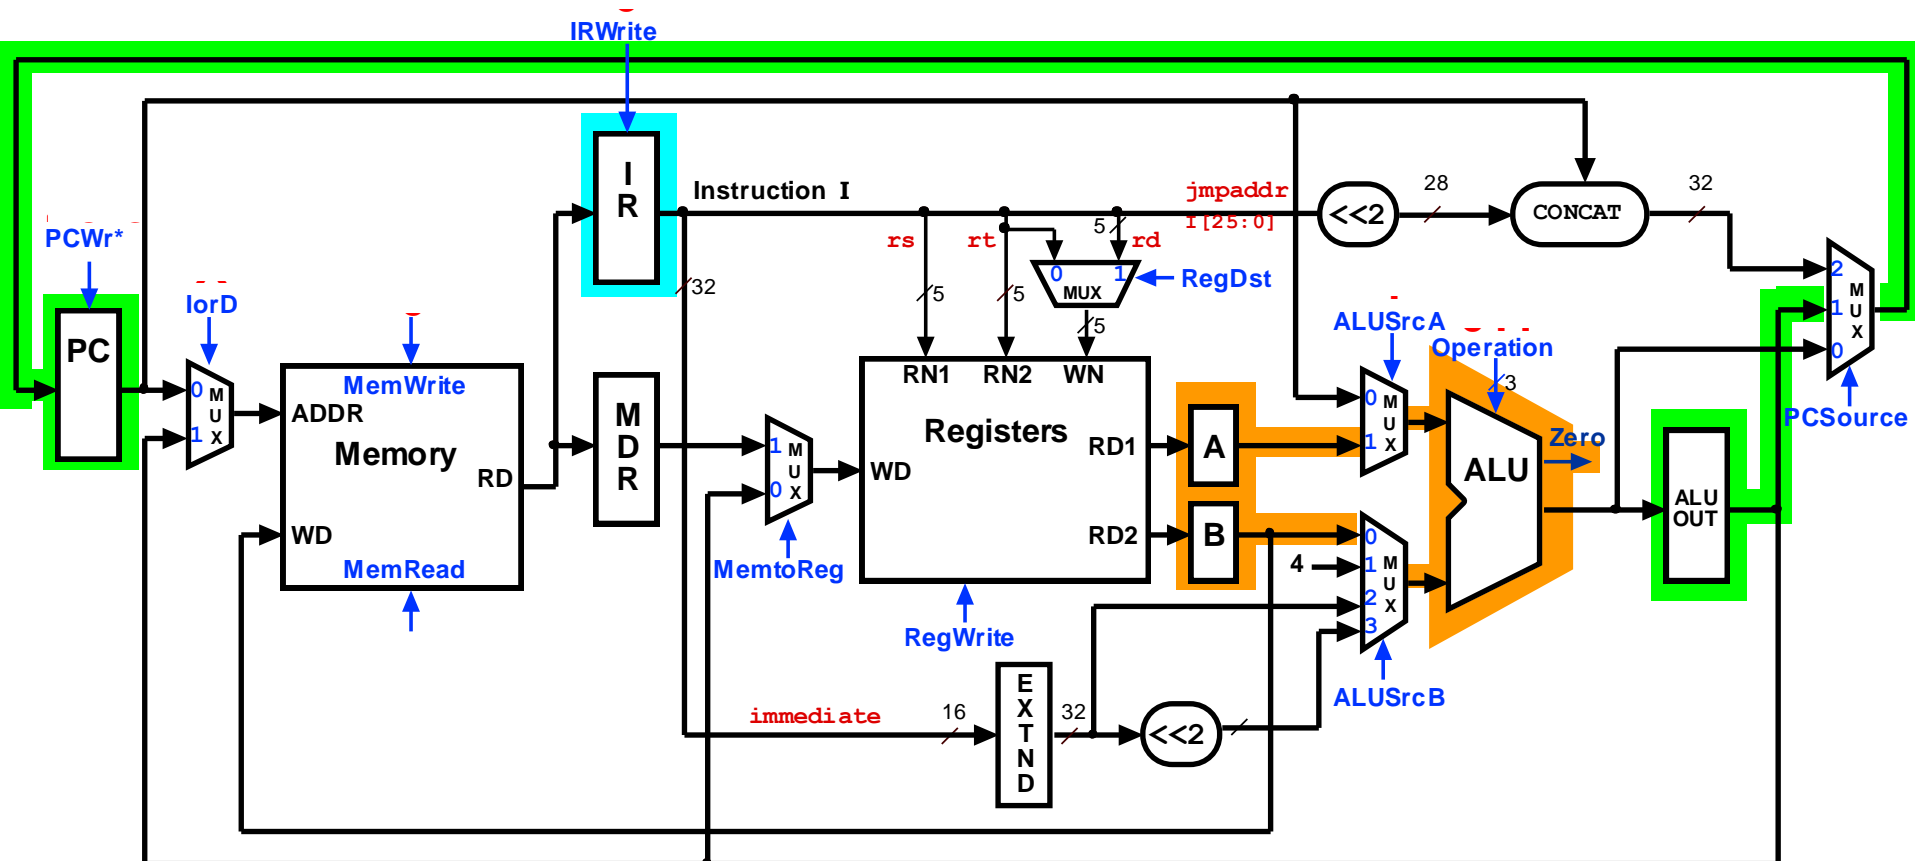
Multicycle Control Step (3): ALU Instruction (R-Type)

ALUOut = A p B;



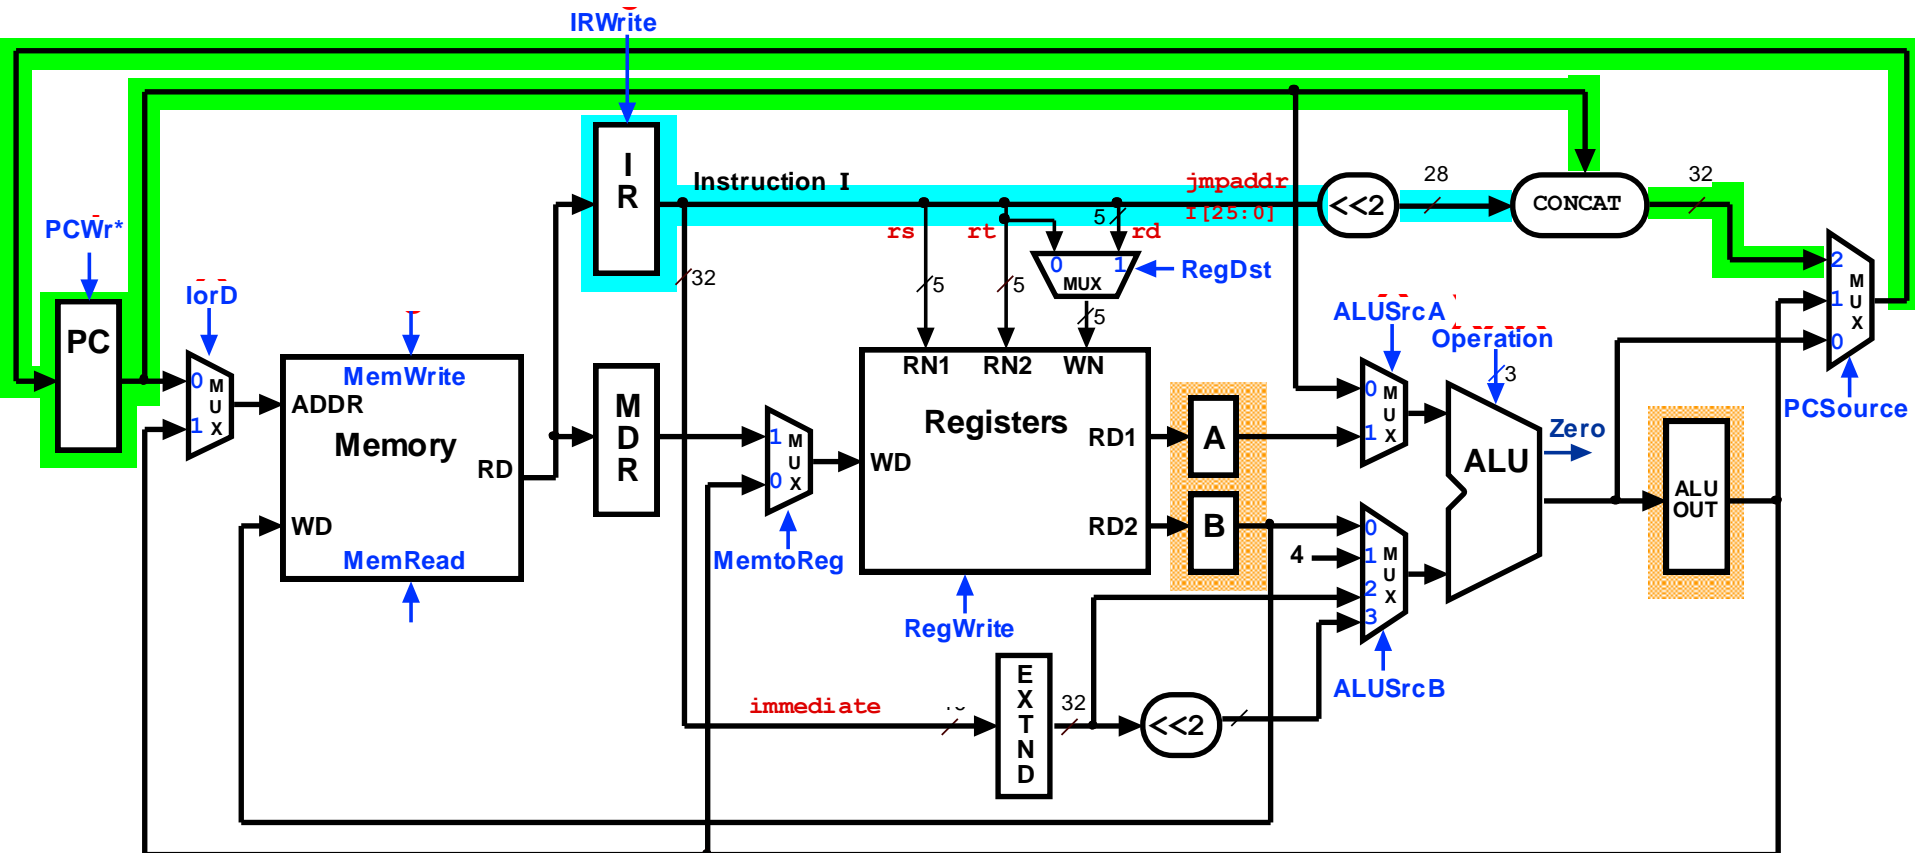
Multicycle Control Step (3): Branch Instructions

```
if (A == B) PC = ALUOut;
```



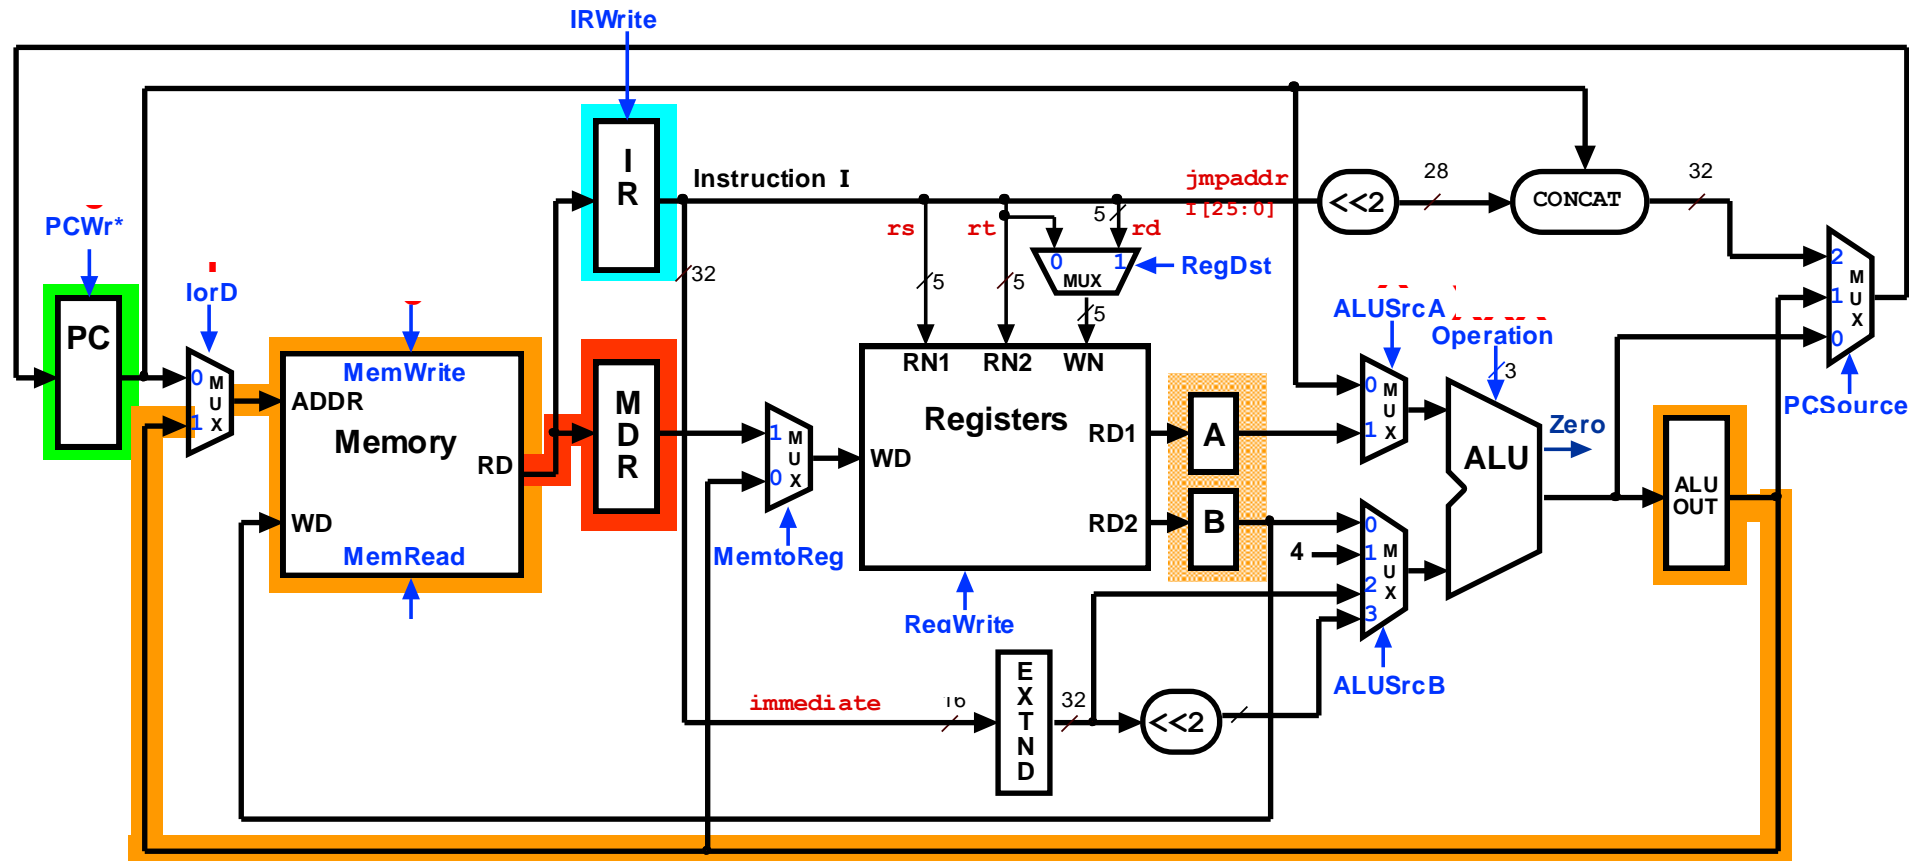
Multicycle Execution Step (3): Jump Instruction

$PC = PC[21-28] \text{ concat } (IR[25-0] \ll 2);$



Multicycle Control Step (4): Memory Access - Read (1_w)

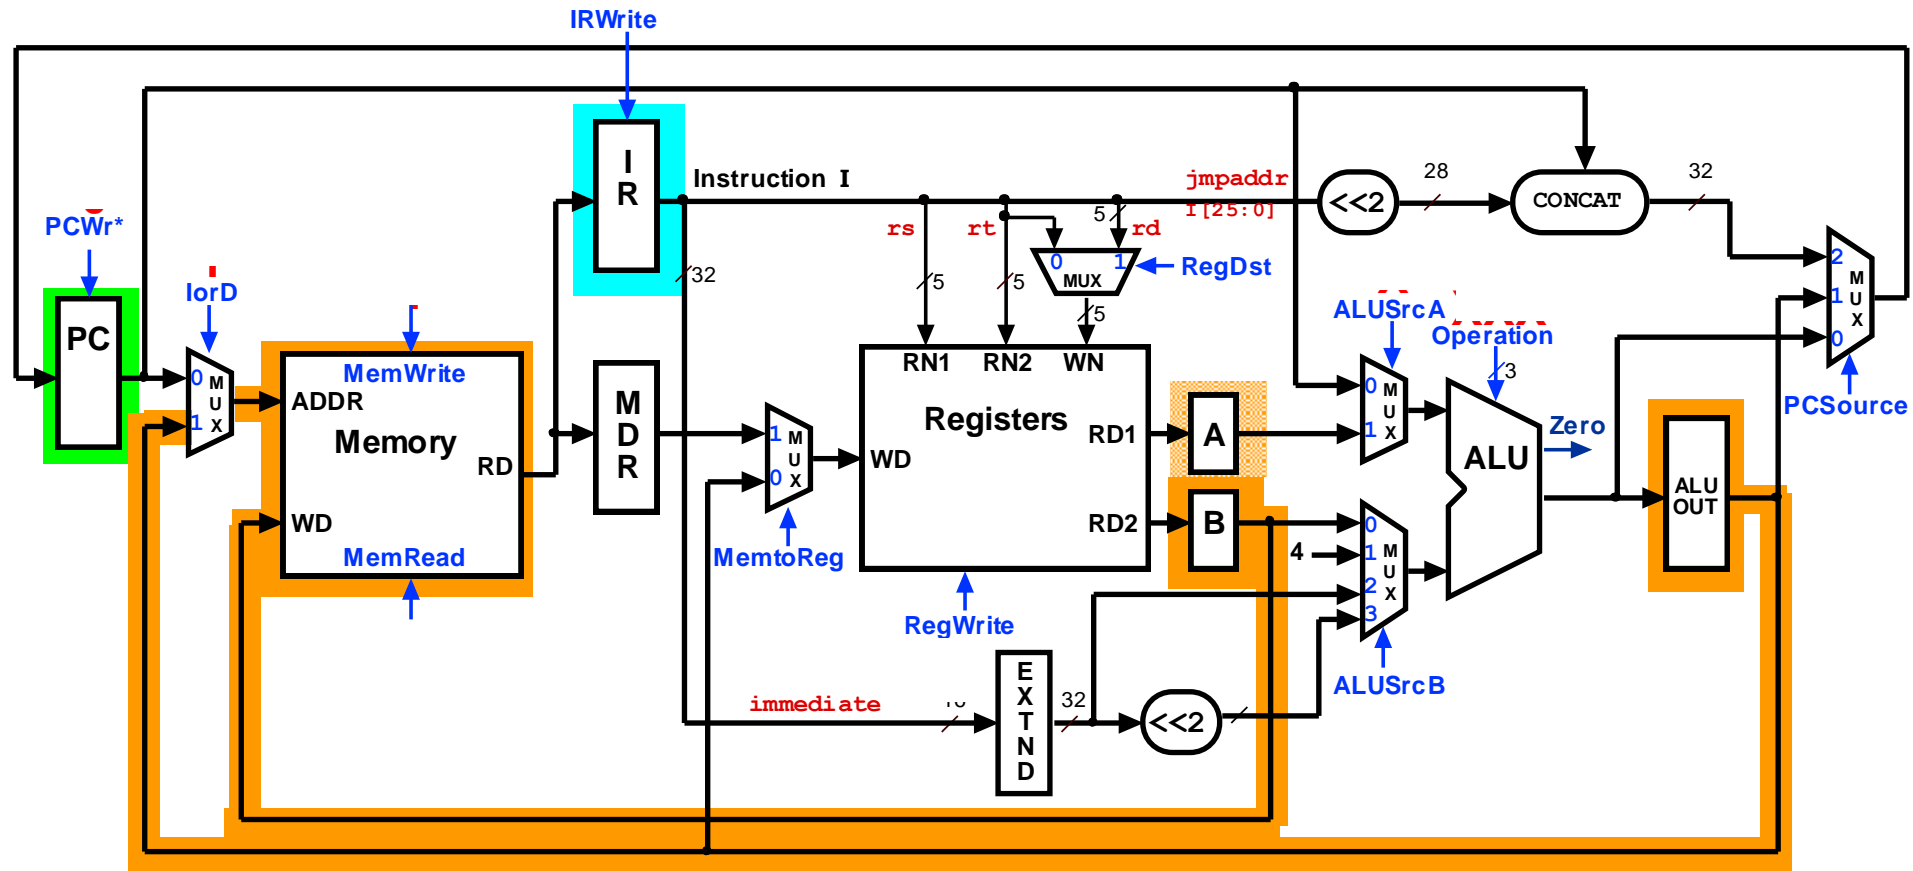
`MDR = Memory[ALUOut];`



Multicycle Execution Steps (4)

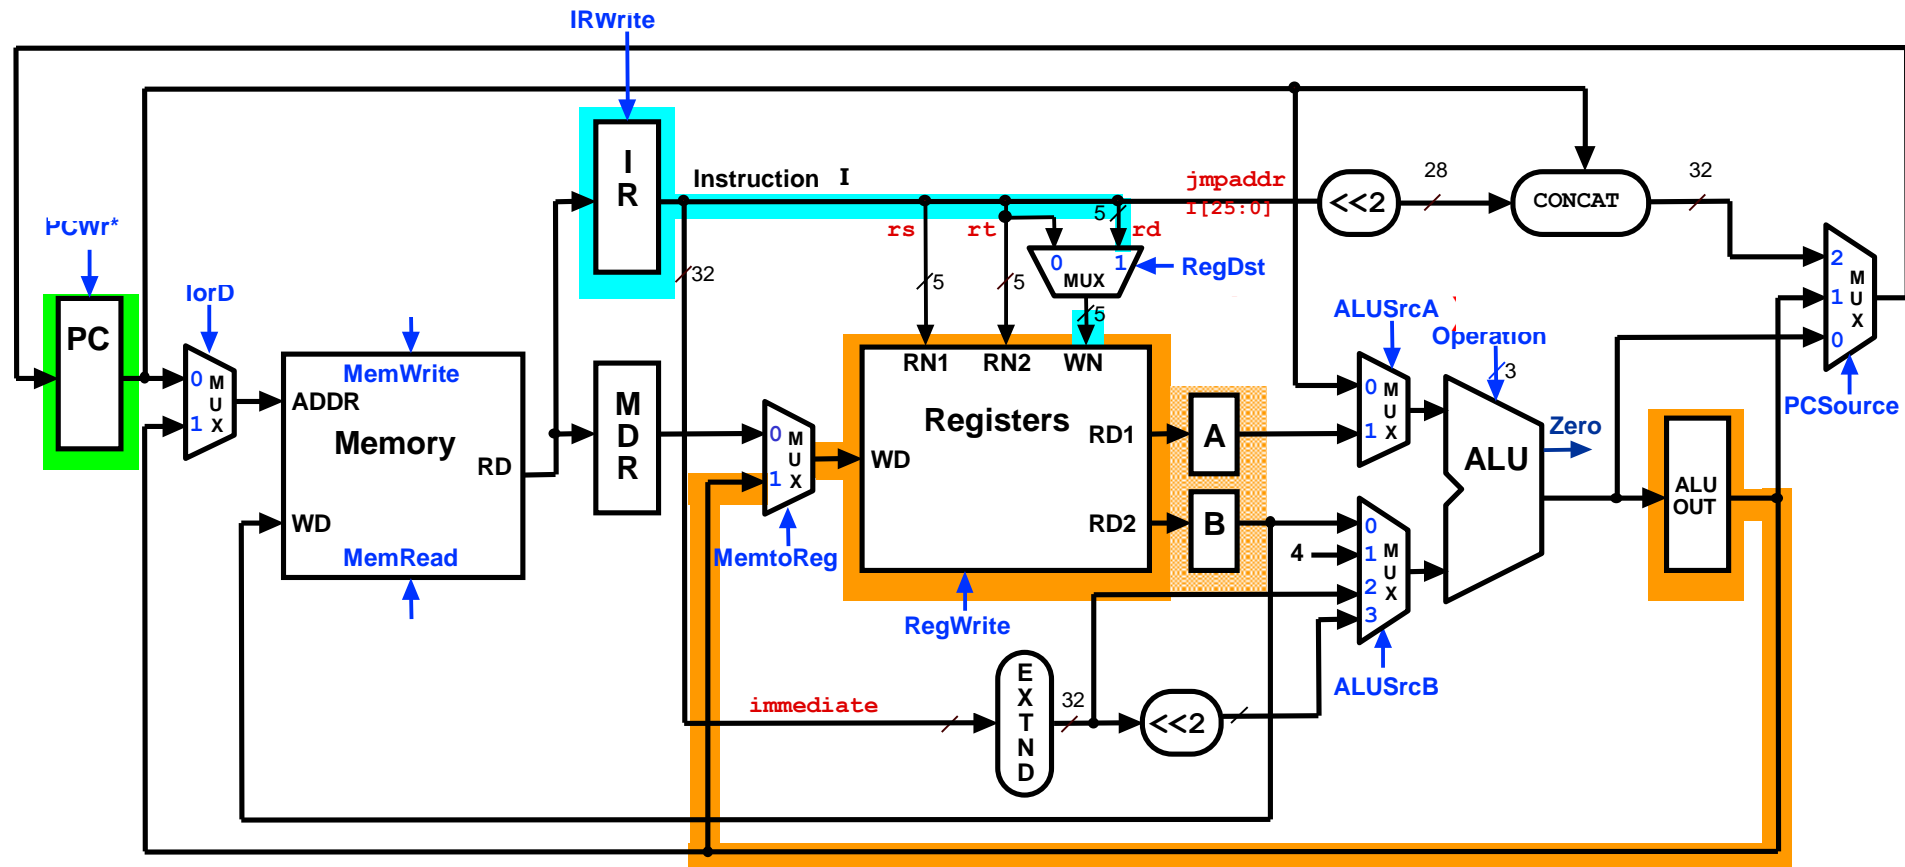
Memory Access - Write (sw)

Memory[ALUOut] = B;



Multicycle Control Step (4): ALU Instruction (R-Type)

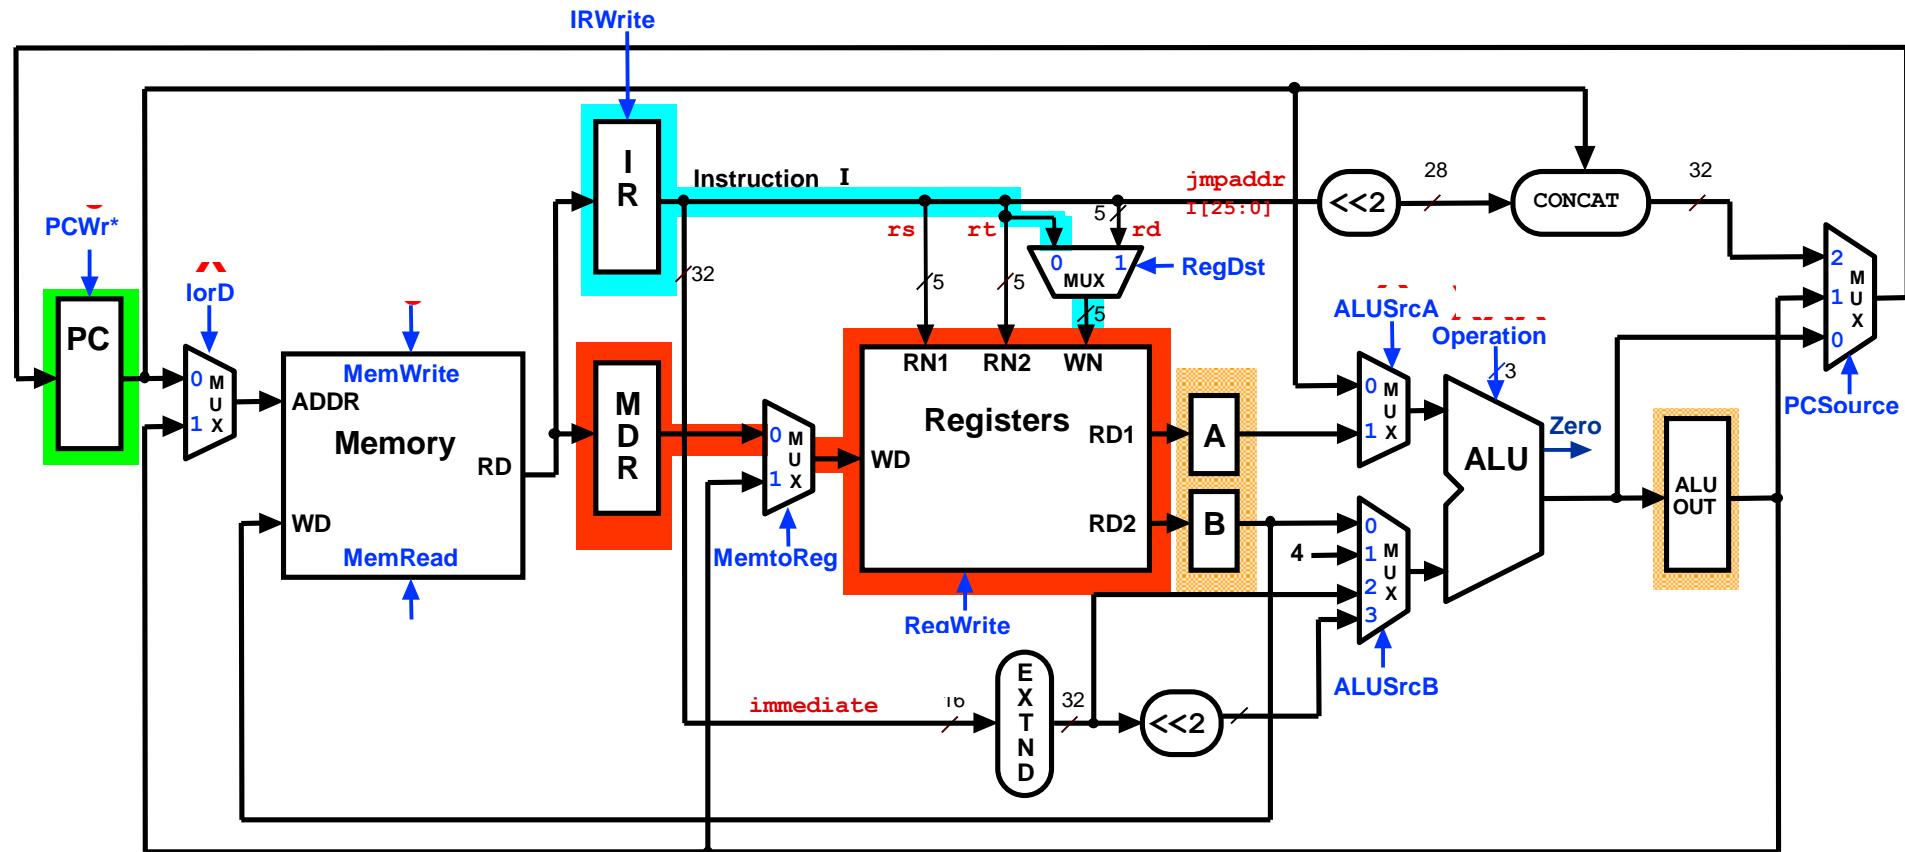
$\text{Reg}[\text{IR}[15:11]] = \text{ALUOut};$ $(\text{Reg}[\text{Rd}] = \text{ALUOut})$



Multicycle Execution Steps (5)

Memory Read Completion (lw)

$\text{Reg}[\text{IR}[20-16]] = \text{MDR};$



Simple Questions

- *How many cycles will it take to execute this code?*

```
        lw $t2, 0($t3)
        lw $t3, 4($t3)
        beq $t2, $t3, Label #assume not equal
        add $t5, $t2, $t3
        sw $t5, 8($t3)
Label:   ...
```

- *What is going on during the 8th cycle of execution?*



Clock time-line

- *In what cycle does the actual addition of \$t2 and \$t3 takes place?*



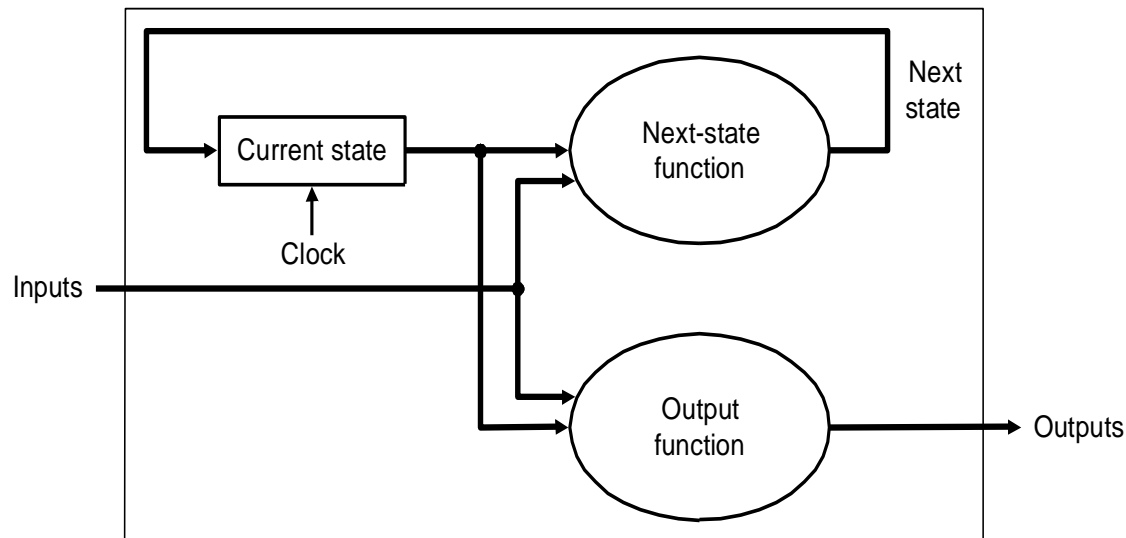
Implementing Control

- Value of control signals is dependent upon:
 - what instruction is being executed
 - which step is being performed
- Use the information we have accumulated to specify a finite state machine
 - specify the finite state machine graphically, or
 - use microprogramming
- Implementation is then derived from the specification

Review: Finite State Machines

Finite state machines (FSMs):

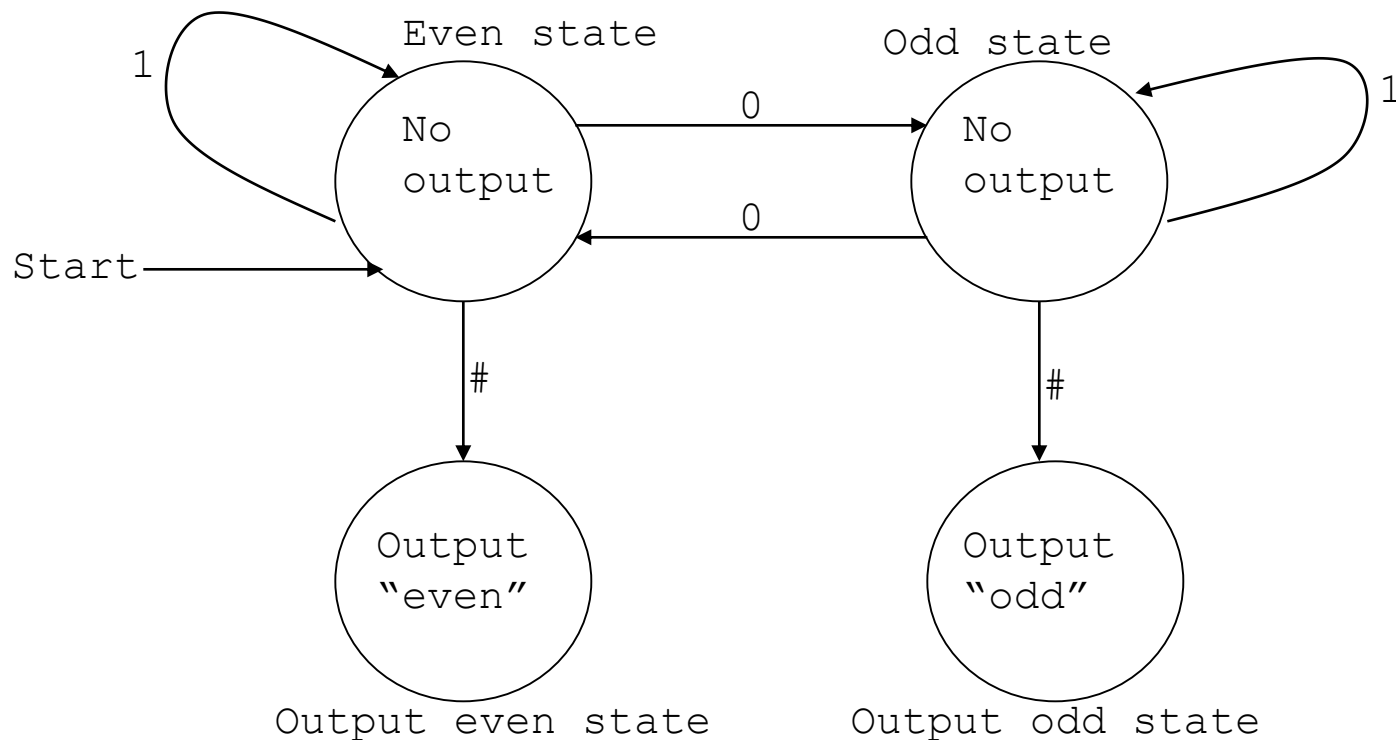
- a set of states and
- next state function, determined by current state and the input
- output function, determined by current state and possibly input





- We'll use a *Moore machine* – output based *only* on current state

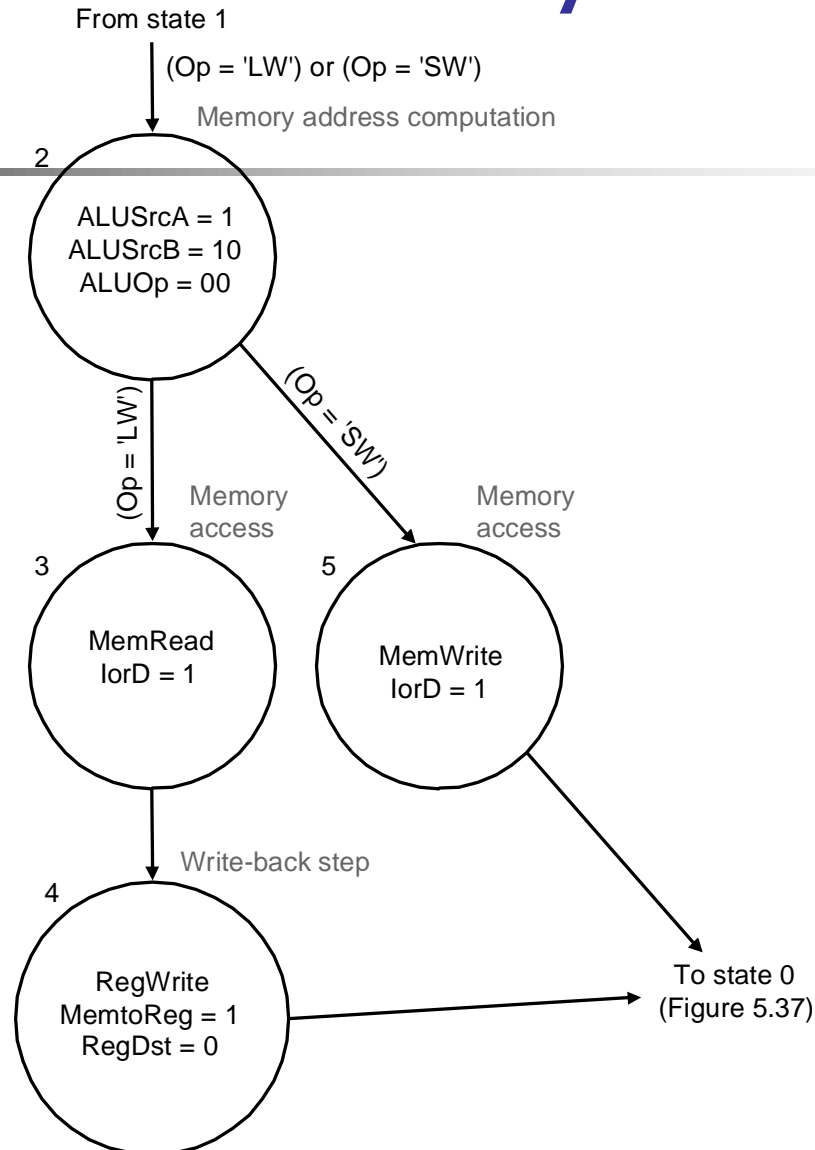
Example: Moore Machine

- The Moore machine below, given *input* a binary string terminated by "#", will *output* "even" if the string has an even number of 0's and "odd" if the string has an odd number of 0's



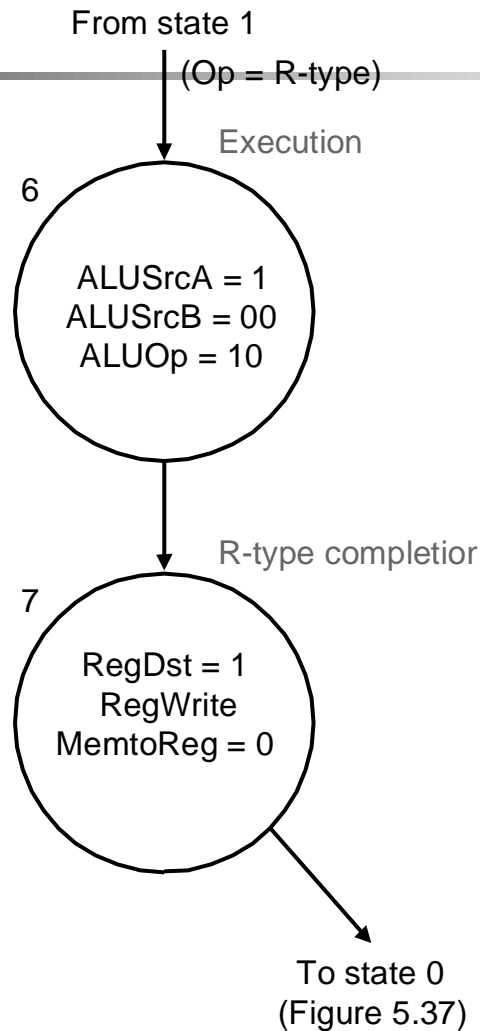


FSM Control: Memory Reference




FSM control for memory-reference has 4 states

FSM Control: R-type Instruction




FSM control to implement R-type instructions has 2 states

FSM Control: Branch Instruction



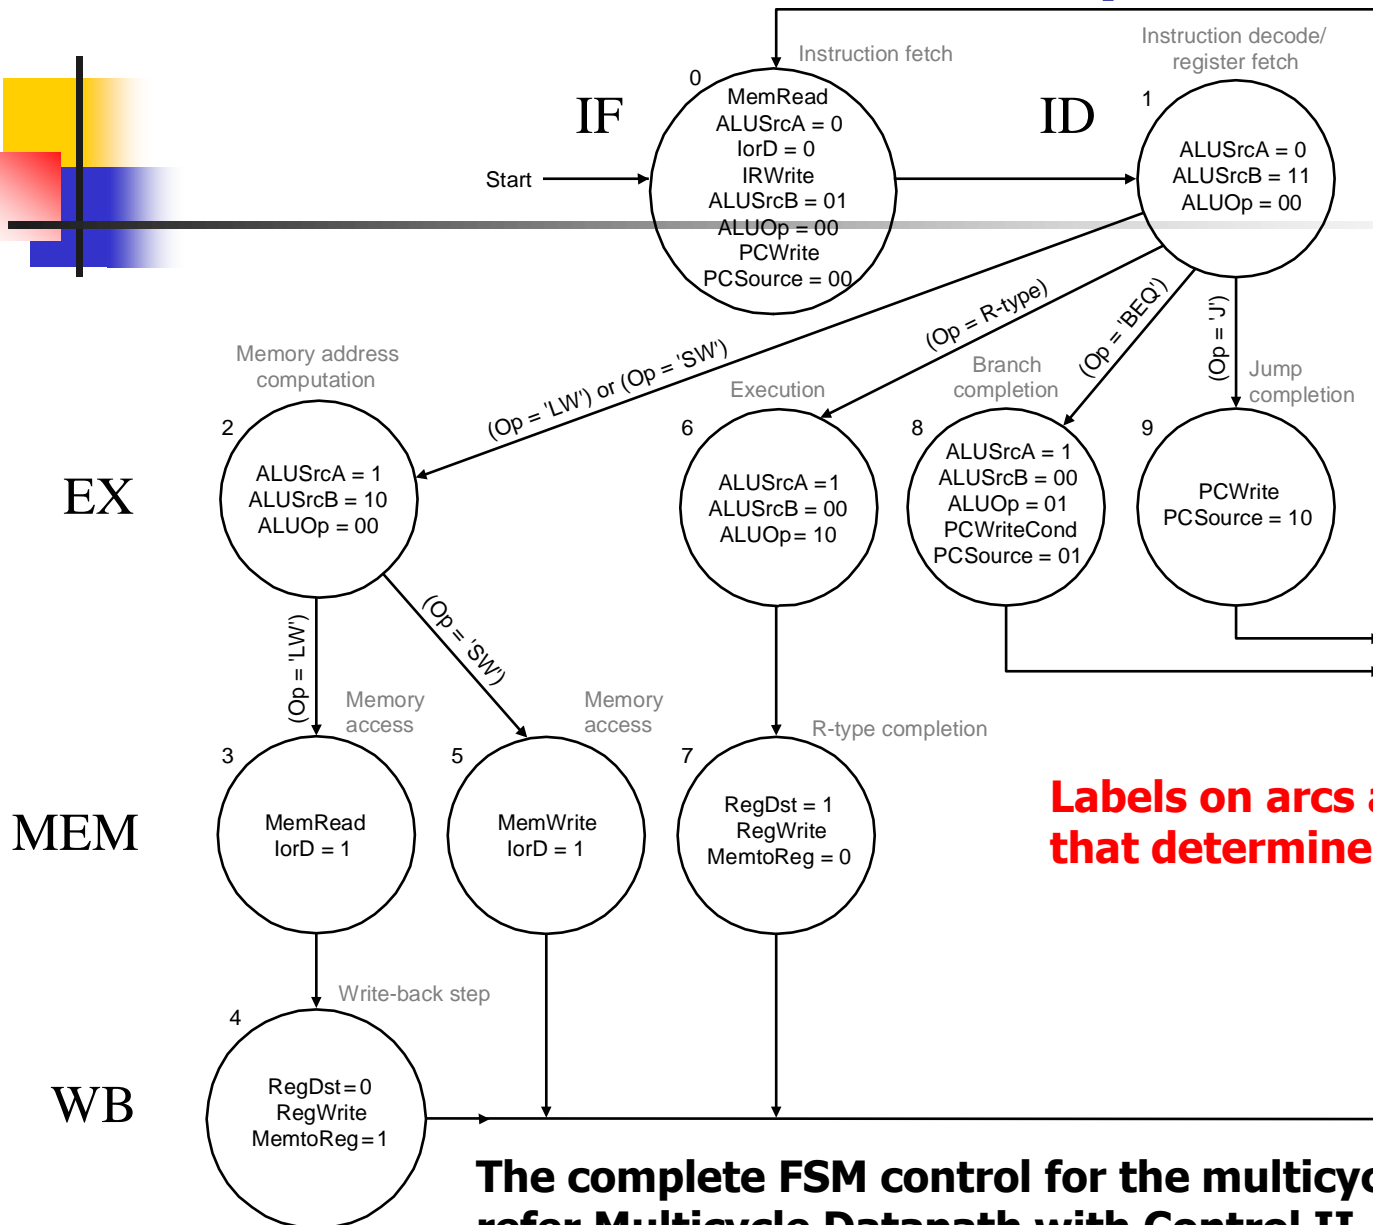
FSM control to implement branches has 1 state

FSM Control: Jump Instruction



FSM control to implement jumps has 1 state

FSM Control: Complete View



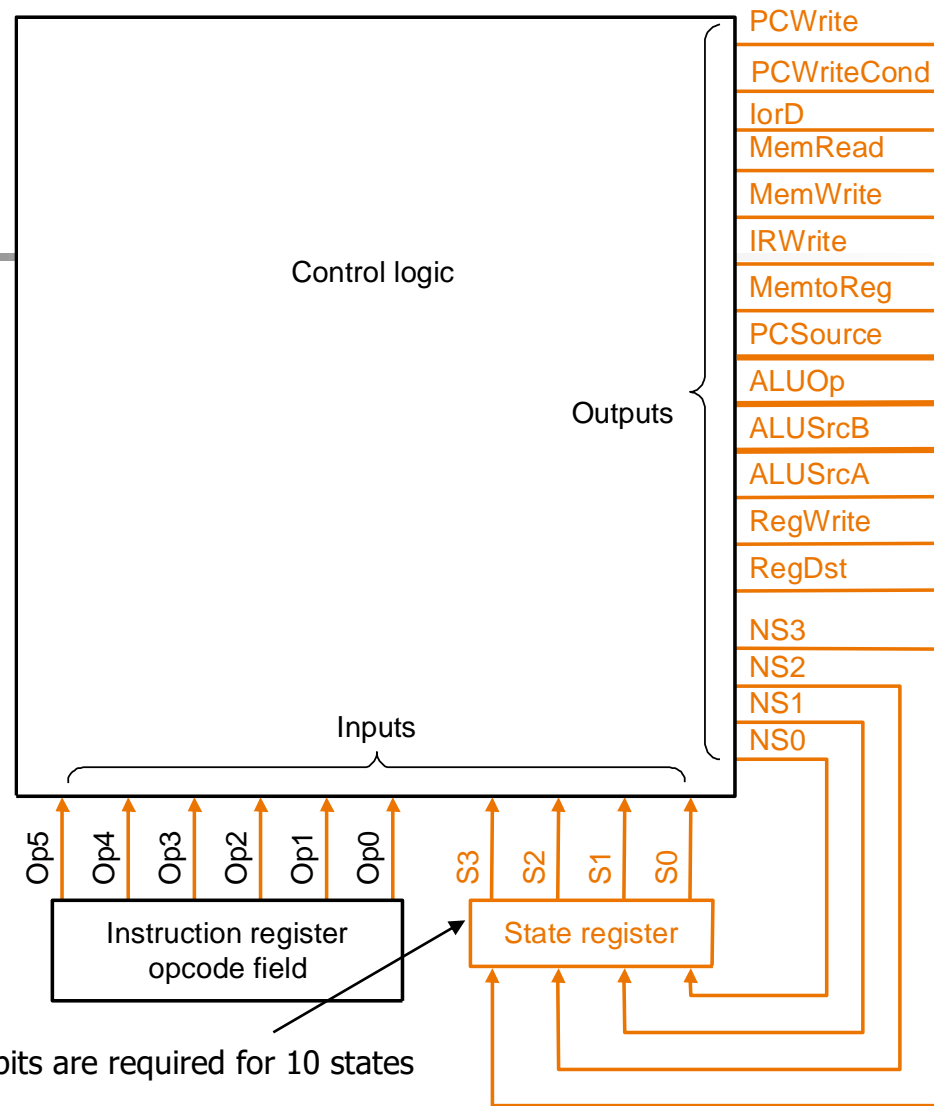
**The complete FSM control for the multicycle MIPS datapath:
refer Multicycle Datapath with Control II**



Example: CPI in a multicycle CPU

- Assume
 - the control design of the previous slide
 - An instruction mix of 22% *loads*, 11% *stores*, 49% *R-type operations*, 16% *branches*, and 2% *jumps*
- What is the CPI assuming each step requires 1 clock cycle?
- Solution:
 - Number of clock cycles from previous slide for each instruction class:
 - *loads* 5, *stores* 4, *R-type instructions* 4, *branches* 3, *jumps* 3
 - $$\begin{aligned}\text{CPI} &= \text{CPU clock cycles} / \text{instruction count} \\ &= \sum (\text{instruction count}_{\text{class } i} \times \text{CPI}_{\text{class } i}) / \text{instruction count} \\ &= \sum (\text{instruction count}_{\text{class } i} / \text{instruction count}) \times \text{CPI}_{\text{class } i} \\ &= 0.22 \times 5 + 0.11 \times 4 + 0.49 \times 4 + 0.16 \times 3 + 0.02 \times 3 \\ &= 4.04\end{aligned}$$

FSM Control: Implementation



Four state bits are required for 10 states


High-level view of FSM implementation: inputs to the combinational logic block are the current state number and instruction opcode bits; outputs are the next state number and control signals to be asserted for the current state

FSM

Control:

PLA

Implem- entation

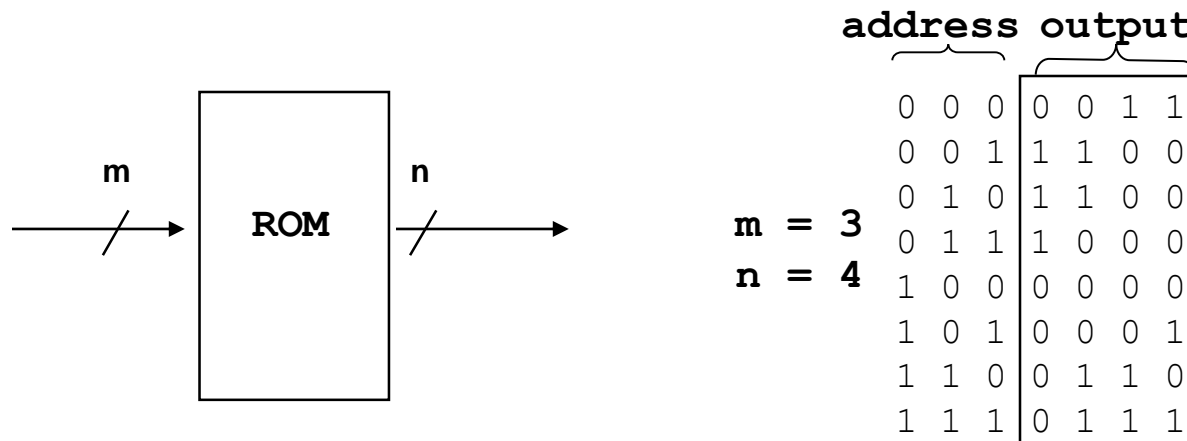


Upper half is the AND plane that computes all the products. The products are carried to the lower OR plane by the vertical lines. The sum terms for each output is given by the corresponding horizontal line

E.g., $IorD = S0.S1.S2.S3 + S0.S1.S2.\overline{S3}$

FSM Control: ROM Implementation

- ROM (Read Only Memory)
 - values of memory locations are fixed ahead of time
- A ROM can be used to implement a truth table
 - if the address is m -bits, we can address 2^m entries in the ROM
 - outputs are the bits of the entry the address points to



The size of an m -input n -output ROM is $2^m \times n$ bits – such a ROM can be thought of as an array of size 2^m with each entry in the array being n bits



FSM Control: ROM vs. PLA

- First improve the ROM: break the table into two parts
 - 4 state bits give the 16 output signals – $2^4 \times 16$ bits of ROM
 - all 10 input bits give the 4 next state bits – $2^{10} \times 4$ bits of ROM
 - Total – 4.3K bits of ROM
- PLA is much smaller
 - can share product terms
 - only need entries that produce an active output
 - can take into account don't cares
- PLA size = $(\text{\#inputs} \times \text{\#product-terms}) + (\text{\#outputs} \times \text{\#product-terms})$
 - FSM control PLA = $(10 \times 17) + (20 \times 17) = 460$ PLA cells
- PLA cells usually about the size of a ROM cell (slightly bigger)



Microprogramming

- Microprogramming is a method of *specifying* FSM control that resembles a programming language – textual rather graphic
 - this is appropriate when the FSM becomes very large, e.g., if the instruction set is large and/or the number of cycles per instruction is large
 - in such situations graphical representation becomes difficult as there may be thousands of states and even more arcs joining them
 - a microprogram is *specification* : *implementation* is by ROM or PLA
- A *microprogram is a sequence of microinstructions*
 - each microinstruction has eight fields (label + 7 functional)
 - Label: used to control microcode sequencing
 - ALU control: specify operation to be done by ALU
 - SRC1: specify source for first ALU operand
 - SRC2: specify source for second ALU operand
 - Register control: specify read/write for register file
 - Memory: specify read/write for memory
 - PCWrite control: specify the writing of the PC
 - Sequencing: specify choice of next microinstruction



Microprogramming

- The *Sequencing* field value determines the execution order of the microprogram
 - value *Seq* : control passes to the sequentially next microinstruction
 - value *Fetch* : branch to the first microinstruction to begin the next MIPS instruction, i.e., the first microinstruction in the microprogram
 - value *Dispatch i* : branch to a microinstruction based on control input and a dispatch table entry (called *dispatching*):
 - Dispatching is implemented by means of creating a table, called *dispatch table*, whose entries are microinstruction labels and which is indexed by the control input. There may be multiple dispatch tables – the value *Dispatch i* in the sequencing field indicates that the *i*th dispatch table is to be used

Control Microprogram

The microprogram corresponding to the FSM control shown graphically earlier:

Label	ALU control	SRC1	SRC2	Register control	Memory	PCWrite control	Sequencing
Fetch	Add	PC	4		Read PC	ALU	Seq
	Add	PC	Extshft	Read			Dispatch 1
Mem1	Add	A	Extend				Dispatch 2
LW2					Read ALU		Seq
				Write MDR			Fetch
SW2					Write ALU		Fetch
Rformat1	Func code	A	B				Seq
				Write ALU			Fetch
BEQ1	Subt	A	B			ALUOut-cond	Fetch
JUMP1						Jump address	Fetch

Microprogram containing 10 microinstructions

Dispatch ROM 1		
Op	Opcode name	Value
000000	R-format	Rformat1
000010	jmp	JUMP1
000100	beq	BEQ1
100011	lw	Mem1
101011	sw	Mem1

Dispatch Table 1

Dispatch ROM 2		
Op	Opcode name	Value
100011	lw	LW2
101011	sw	SW2

Dispatch Table 2



Microcode: Trade-offs

- Specification advantages
 - easy to design and write
 - typically manufacturer designs architecture and microcode in parallel
- Implementation advantages
 - easy to change since values are in memory (e.g., off-chip ROM)
 - can emulate other architectures
 - can make use of internal registers
- Implementation disadvantages
 - control is implemented nowadays on same chip as processor so the advantage of an off-chip ROM does not exist
 - ROM is no longer faster than on-board cache
 - there is little need to change the microcode as general-purpose computers are used far more nowadays than computers designed for specific applications



Summary

- *Techniques described in this chapter to design datapaths and control are at the core of all modern computer architecture*
- Multicycle datapaths offer two great advantages over single-cycle
 - functional units can be reused within a single instruction if they are accessed in different cycles – reducing the need to replicate expensive logic
 - instructions with shorter execution paths can complete quicker by consuming fewer cycles
- Modern computers, in fact, take the multicycle paradigm to a higher level to achieve greater instruction throughput:
 - *pipelining* (next topic) where multiple instructions execute simultaneously by having cycles of different instructions overlap in the datapath
 - *the MIPS architecture was designed to be pipelined*