



Computer Organization and Architecture

Instructor: Dr. Rushdi abu
zneit

Slide Sources: Based on CA:
aQA by Hennessy/Patterson.



Advanced Topic:

Dynamic Hardware Branch

Prediction

Ca:aQA Sec. 3.4




Static Prediction

- Does not take into account the *run-time history* of the particular branch instruction – whether it was taken or not taken recently, how often it was taken or not taken, etc.
- Simplest static prediction:
 - predict *always taken*
 - predict *always not taken*
- More complex static prediction:
 - performed at compile time by analyzing the program...

Dynamic prediction: 1-bit Predictor

- *Branch- prediction buffer* or *branch history table (BHT)* is a cache indexed by a fixed lower portion of the address of the branch instruction
- *1-bit prediction*: for each index the BHT contains one *prediction bit* (also called *history bit*) that says if the branch was last taken or not – *prediction is that branch will do the same again*



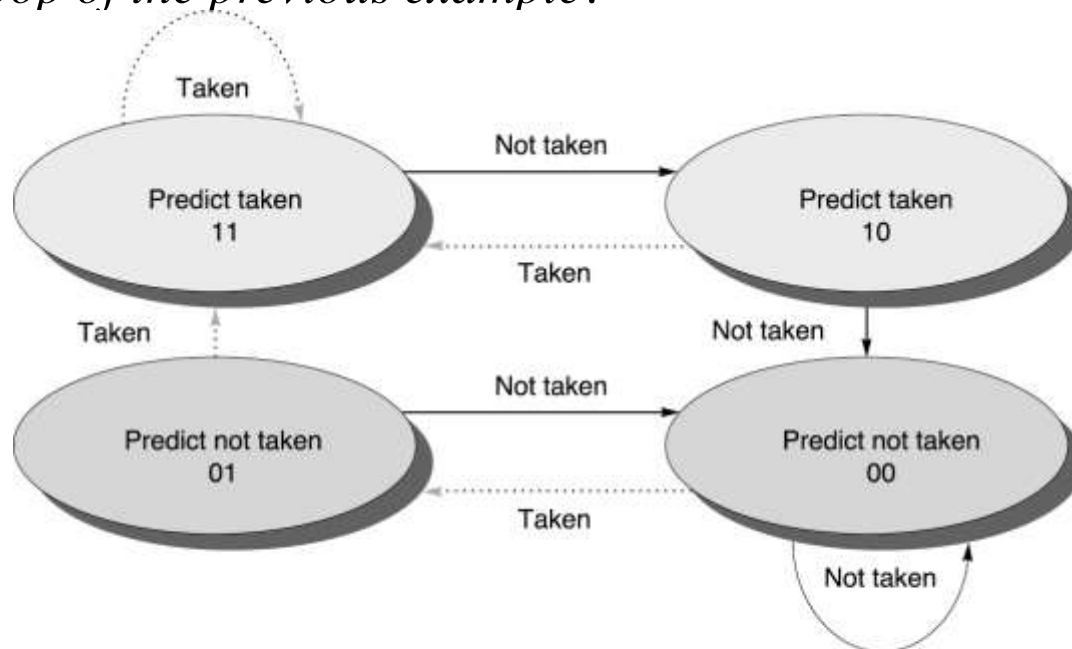


Dynamic prediction: 1-bit Predictor


- Meaning of prediction bit
 - 1 = branch was last taken
 - 0 = branch was last not taken
- Using the BHT
 - index into the BHT and use the prediction bit to predict branch behavior
 - note the prediction bit may have been set by a different branch instruction with the same lower address bits but that does not matter – the history bit is simply a *hint*
 - if prediction is wrong, invert prediction bit
- Example: *Consider a loop branch that is taken 9 times in a row and then not taken once. What is the prediction accuracy of 1-bit predictor for this branch assuming only this branch ever changes its corresponding prediction bit?*
 - Answer: 80%. Because there are two mispredictions – one on the first iteration and one on the last iteration. *Why?*

Dynamic prediction: 2-bit Predictor

- *2-bit prediction*: for each index the BHT contains two prediction bits that change as in the figure below
- Key idea: the prediction must be wrong *twice* for it to be changed
- Example: *What is the prediction accuracy of a 2-bit predictor on the loop of the previous example?*




2-bit Predictor Statistics



Prediction accuracy of 4K-entry 2-bit prediction buffer on SPEC89 benchmarks: accuracy is lower for integer programs (gcc, espresso, eqntott, li) than for FP

2-bit Predictor Statistics



Prediction accuracy of 4K-entry 2-bit prediction buffer vs. "infinite" 2-bit buffer: increasing buffer size from 4K does not significantly improve performance



n-bit Predictors

- Use an n-bit counter which, therefore, represents a value X where $0 \leq X \leq 2^n - 1$
 - increment X if branch is taken (to a max of 2^n)
 - decrement X if branch is not taken (to a min of 0)
 - If $X \geq 2^{n-1}$, then predict taken; otherwise, untaken
- Studies show that there is no significant improvement in performance using n-bit predictors with $n > 2$, so *2-bit predictors are implemented in most systems*



Correlating Predictors

if (aa == 2)		DSUBUI	R3, R1, #2	
aa = 0;		BNEZ	R3, L1	; branch b1 (aa != 2)
if (bb == 2)		DADD	R1, R0, R0	; aa = 0
bb=0;	L1:	DSUBUI	R3, R2, #2	
if (aa != bb) { ...		BNEZ	R3, L2	; branch b2 (bb != 2)
		DADD	R2, R0, R0	; bb = 0
	L2:	DSUB	R3, R2, R1	; R3 = aa - bb
		BEQZ	R3, L3	; branch b3 (aa == bb)

**Code fragment from
eqntott SPEC89 benchmark**

**Corresponding MIPS code:
aa is in R1, bb is in R2**

- Key idea: branch b3 behavior is correlated with the behavior of branches b1 and b2
 - because *if branches b1 and b2 are both not taken*, then the statements following the branches will set aa=0 and bb=0
⇒ *b3 will be taken*



Correlating Predictors: Simple Example

```
if (d == 0)                BNEZ      R1, L1      ; branch b1   (d != 0)
    d = 1;                DADDIU     R1, R0, #1   ; d==0, so d=1
if (d == 1) {              L1: DADDIU     R3, R1, #-1
    ...                    BNEZ      R3, L2      ; branch b2   (d != 1)
    L2
```

**Simple code
fragment**

**Corresponding MIPS code:
d is in R1**

Initial Value of d	Values of d				
	d==0?	b1	before b2	d==1?	b2
0	yes	not taken	1	yes	not taken
1	no	taken	1	yes	not taken
2	no	taken	2	no	taken

Possible execution sequences assuming d is one of 0, 1, or 2

Correlating Predictors: Impact of Ignoring Correlation

Initial Value of d	d==0?	b1	Values of d before b2	d==1?	b2
0	yes	not taken	1	yes	not taken
1	no	taken	1	yes	not taken
2	no	taken	2	no	taken

Possible execution sequences assuming d is one of 0, 1, or 2

d=	b1 prediction	b1 action	new b1 prediction	b2 prediction	b2 action	new b2 prediction
2	NT	T	T	NT	T	T
0	T	NT	NT	T	NT	NT
2	NT	T	T	NT	T	T
0	T	NT	NT	T	NT	NT

Behavior of 1-bit predictor initialized to not taken with d alternating between 2 and 0: 100% misprediction!

Correlating Predictors: Taking Correlation into Account

Prediction bits	Prediction if last branch not taken	Prediction if last branch taken
NT/NT	NT	NT
NT/T	NT	T
T/NT	T	NT
T/T	T	T

Meaning of 1-bit predictor with 1 bit of correlation: equivalent to assuming two separate prediction bits – one assuming last branch executed was not taken and one assuming the last branch executed was taken

d=	b1 prediction	b1 action	new b1 prediction	b2 prediction	b2 action	new b2 prediction
2	NT/NT	T	T/NT	NT/NT	T	NT/T
0	T/NT	NT	T/NT	NT/T	NT	NT/T
2	T/NT	T	T/NT	NT/T	T	NT/T
0	T/NT	NT	T/NT	NT/T	NT	NT/T

Behavior of 1-bit predictor with 1-bit of correlation, assuming initially NT/NT and d alternating between 0 and 2: mispredictions only on first iteration. Analyze why...!
Predictions used in red.




Correlating Predictors: (m,n) Predictors

- The correlating predictor as before – 1 bit of prediction plus 1 correlating bit – is called a *(1,1) predictor*
- Generalization of the (1,1) predictor is the *(m,n) predictor*
- *(m,n) predictor* : use the *behavior of the last m branches* to choose from one of 2^m branch predictors, *each of which is an n -bit predictor*
- The history of the most recent m branches is recorded in an m -bit shift register called the m -bit *global history register*
 - shift in the behavior bit for the most recent branch, shift out the the bit for the least recent branch
- Index into the BHT by concatenating the lower bits of the branch instruction address with the m -bit global history to access an n -bit entry

Correlating Predictors: Data Structures

A (2,2) BHT uses the 2-bit global history to choose from one of four predictors for each branch address. In the following 16 entry BHT each of the four predictors has 16 2-bit entries and 4 bits of the branch instruction address is used to choose one of these entries.

The BHT is shown as four separate arrays; in actuality it is one 2-bit wide linear array.






Simple Examples

- *Note* : A 2-bit predictor with no global history is simply a (0,2) predictor
- *How many bits are in a (0,2) predictor with 4K entries?*
 - Total bits = $2^0 \times 2 \times 4K = 8K$
- *How many bits are in a (2,2) predictor with 16 entries (shown in previous figure)?*
 - Total bits = $64 \times 2 = 128$
- *How many branch-selected entries are in a (2,2) predictor that has a total of 8K bits in the BHT? How many bits of the branch address are used to access the BHT*
 - Total bits = $8K = 2^2 \times 2 \times \text{no. prediction entries selected by branch}$
Therefore, no. prediction entries selected by branch = 1K
Therefore, no. of bits of branch address used to index BHT = 10

2-bit Predictor Statistics



A (2,2) predictor with 1K entries outperforms not only a 2-bit no-history predictor with 4K entries but also a 2-bit no-history predictor with unlimited entries