

نمودج رقم (2)

# COURSE TITLE

Course Code	:30102315
-------------	-----------

Credit Hours	:3
--------------	----

Prerequisite	:30102214
--------------	-----------



### Instructor Information

Name	: Dr. Rushdi Saleem Abu Zneit				
Office No.	18, building 17 floor 3				
Tel (Ext)	Contact Telephone( Department Telephone)				
E-mail	<a href="mailto:dr.rushdizneit@bau.edu.jo">dr.rushdizneit@bau.edu.jo</a>				
Office Hours	Online				
Class Times	Building	Day	Start Time	End Time	Room No.
		Sun,Tue, Thu	10:00	11:00	Online



## **Course description {From course plan}**

**Course Title; Computer Architecture**  
**Credit Hour(3)**

**[Pre-req. 30102214**

Textbook: Textbook Title

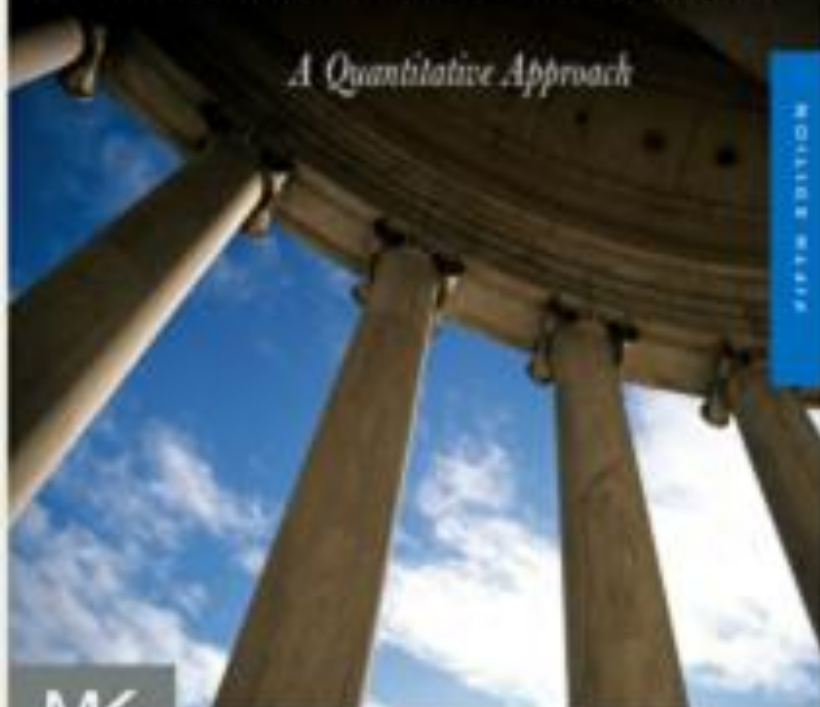
JOHN L. HENNESSY    DAVID A. PATTERSON

# COMPUTER ARCHITECTURE

*A Quantitative Approach*

FOURTH EDITION

MK  
Morgan Kaufmann





## **COURSE OBJECTIVES**

The module aims to provide students with a fundamental knowledge of computer hardware and computer systems, with an emphasis on system design and performance. The module concentrates on the principles underlying systems organization, issues in computer system design, and contrasting implementations of modern systems. The module is central to the aims of the Computing Systems degree course, for which it is core.

## COURSE SYLLABUS


Week	Course Topic
Week 1	Course overview & Introduction
Week 2	Performance, Benchmarks, Measurements
Week 3	Pipelining Review
Week 4	Review: Instruction Set Design & Memory Hierarchy Design
Week 5	Memory Hierarchy Design
Week 6	Pipeline hazards, Instruction Re-scheduling
Week 7	ILP - Static: Loop Unrolling
Week 8	Midterm Exam
Week 9	Introduction to Branch Predictors
Week 10	Dynamic Instruction Level Parallelism: Scoreboarding Technique
Week 11	Dynamic Instruction Level Parallelism: Tomoslus' technique
Week 12	Data-Level Parallelism in Vector, SIMD, and GPU Architectures
Week 13	Data-Level Parallelism in Vector, SIMD, and GPU Architectures
Week 14	Thread-Level Parallelism
Week 15	Warehouse-Scale Computers to Exploit Request-Level and Data-Level Parallelism
Week 16	Final Exam

Course Topic
Introduction
Performance, Benchmarks, Measurements
Review: pipelining
Review: Instruction Set Design
Memory Hierarchy Design
Pipeline hazards, Instruction Re-scheduling
ILP – Static: Loop Unrolling
Midterm Exam
Introduction to Branch Predictors
Dynamic Instruction Level Parallelism: Scoreboarding Technique
Dynamic Instruction Level Parallelism: Tomoslus' technique
Data-Level Parallelism in Vector, SIMD, and GPU Architectures
Thread-Level Parallelism
Warehouse-Scale Computers to Exploit Request-Level and Data-Level Parallelism
Final Exam





## Week 5



## Chapter 2

# MEMORY HIERARCHY DESIGN

Week 5

# Introduction

- Programmers want **very large memory** with **low latency**
- **Fast memory** technology is more **expensive** per bit than slower memory
- Solution: organize memory system into a hierarchy
  - Entire addressable memory space available in largest, slowest memory
  - Incrementally smaller and faster memories, **each containing a subset of the memory below it**, proceed in steps up toward the processor
- Temporal and spatial locality insures that nearly all references can be found in smaller memories
  - Gives the allusion of a large, fast memory being presented to the processor


Temporal Locality :

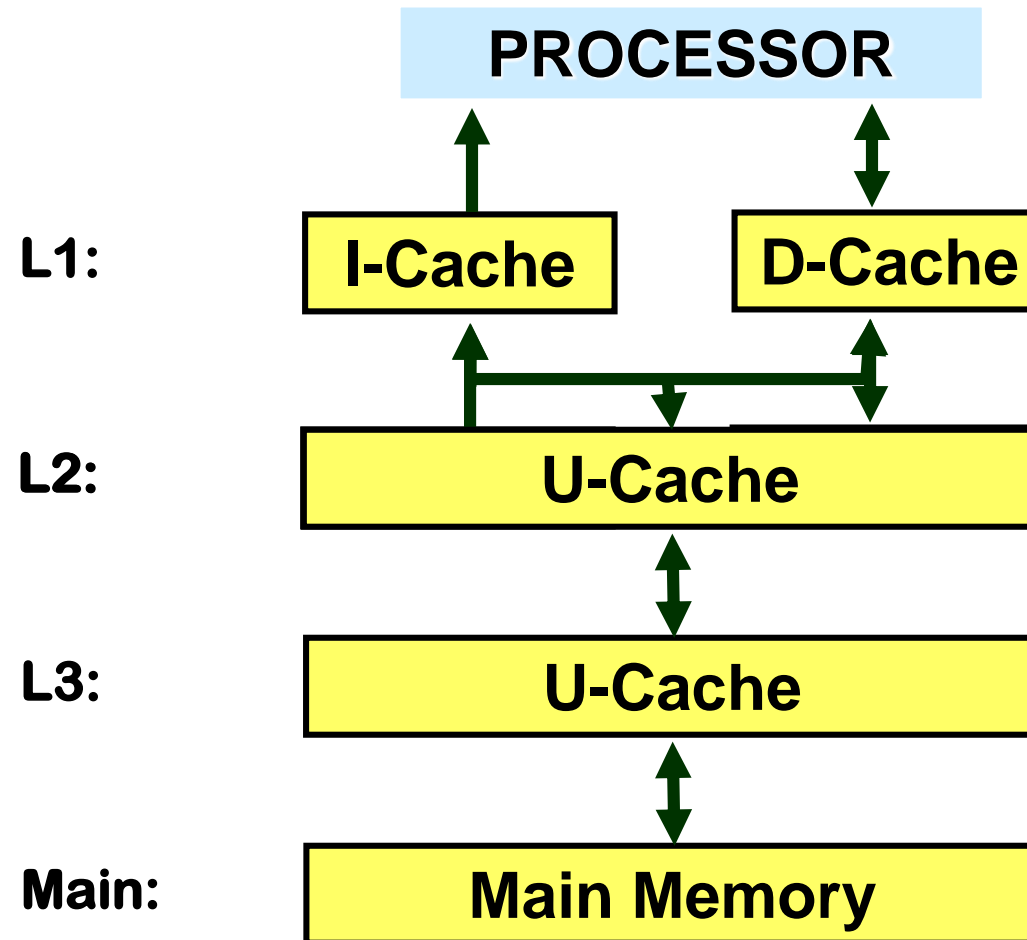
Temporal Locality means that an instruction which is recently executed has high chances of execution again. So the instruction is kept in cache memory such that it can be fetched easily and takes no time in searching for the same instruction.

## Spatial Locality :

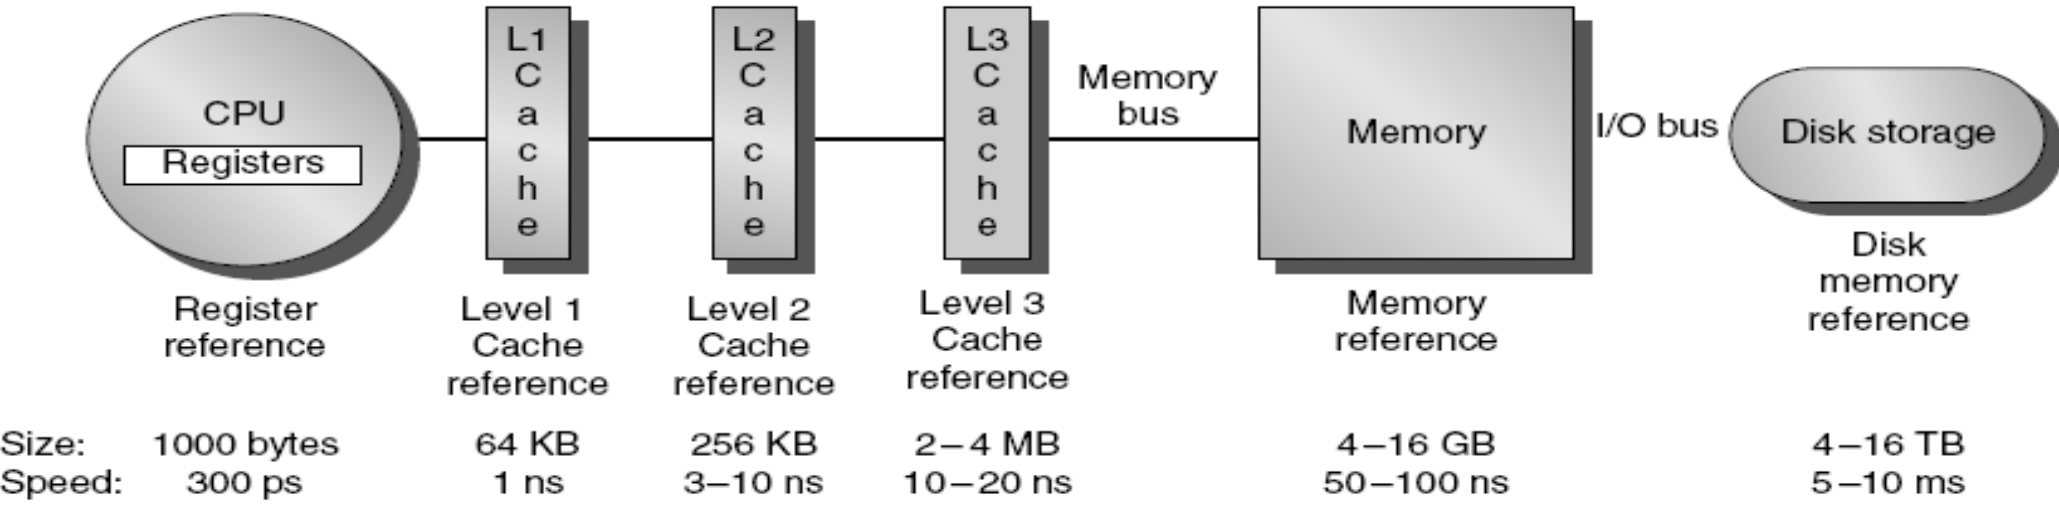
Spatial Locality means that all those instructions which are stored nearby to the recently executed instruction have high chances of execution. It refers to the use of data elements(instructions) which are relatively close in storage locations.

# Memory Hierarchy

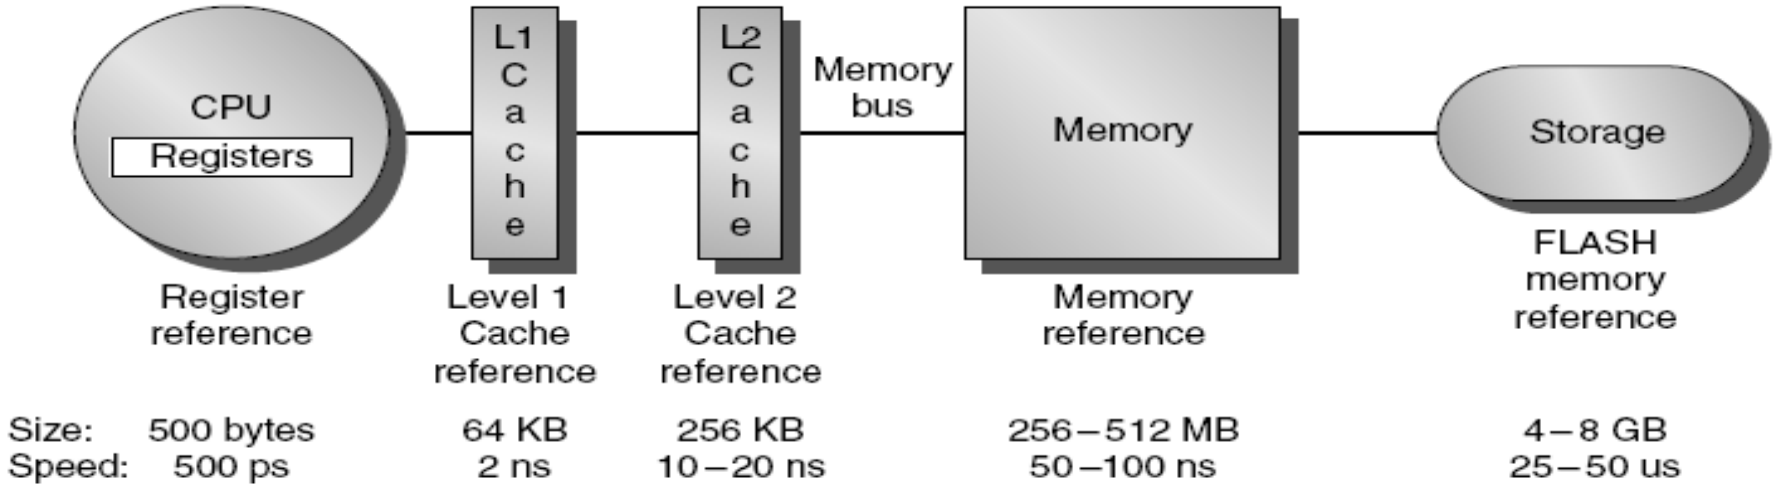




# Memory Hierarchy




(a) Memory hierarchy for server



(b) Memory hierarchy for a personal mobile device



# Memory Performance Gap



# Memory Hierarchy Design

- Memory hierarchy design becomes more crucial with recent multi-core processors:
  - Aggregate peak bandwidth grows with # cores:
    - Intel Core i7 can generate two references per core per clock
    - Four cores and 3.2 GHz clock
      - 25.6 billion 64-bit data references/second +
      - 12.8 billion 128-bit instruction references
      - = 409.6 GB/s!
    - DRAM bandwidth is only 6% of this (25 GB/s)
    - Requires:
      - Multi-port, pipelined caches
      - Two levels of cache per core
      - Shared third-level cache on chip

# Intel Processors (3rd Generation Intel Core)

- **Intel Core i7** \*
- 4 cores 8 threads
- 2.5-3.5 GHz (Normal); 3.7 or 3.9GHz (Turbo)
- **Intel Core i5**
- 4 cores 4 threads (or 2 cores 4 threads)
- 2.3-3.4GHz (Normal); 3.2-3.8Ghz (Turbo)
- **Intel Core i3**
- 2 cores 4 threads
- 3.3 or 3.4 GHz




# Performance and Power

- High-end microprocessors have >10 MB on-chip cache
  - Consumes large amount of area and power budget

# Memory Hierarchy Basics

- When a word is not found in the cache, a *miss* occurs:
  - Fetch word from lower level in hierarchy, requiring a higher latency reference
  - Lower level may be another cache or the main memory
  - Also fetch the other words contained within the *block*
    - Takes advantage of spatial locality
  - Place block into cache in any location within its *set*, determined by address
    - block address MOD number of sets


# Placement Problem



# Placement Policies

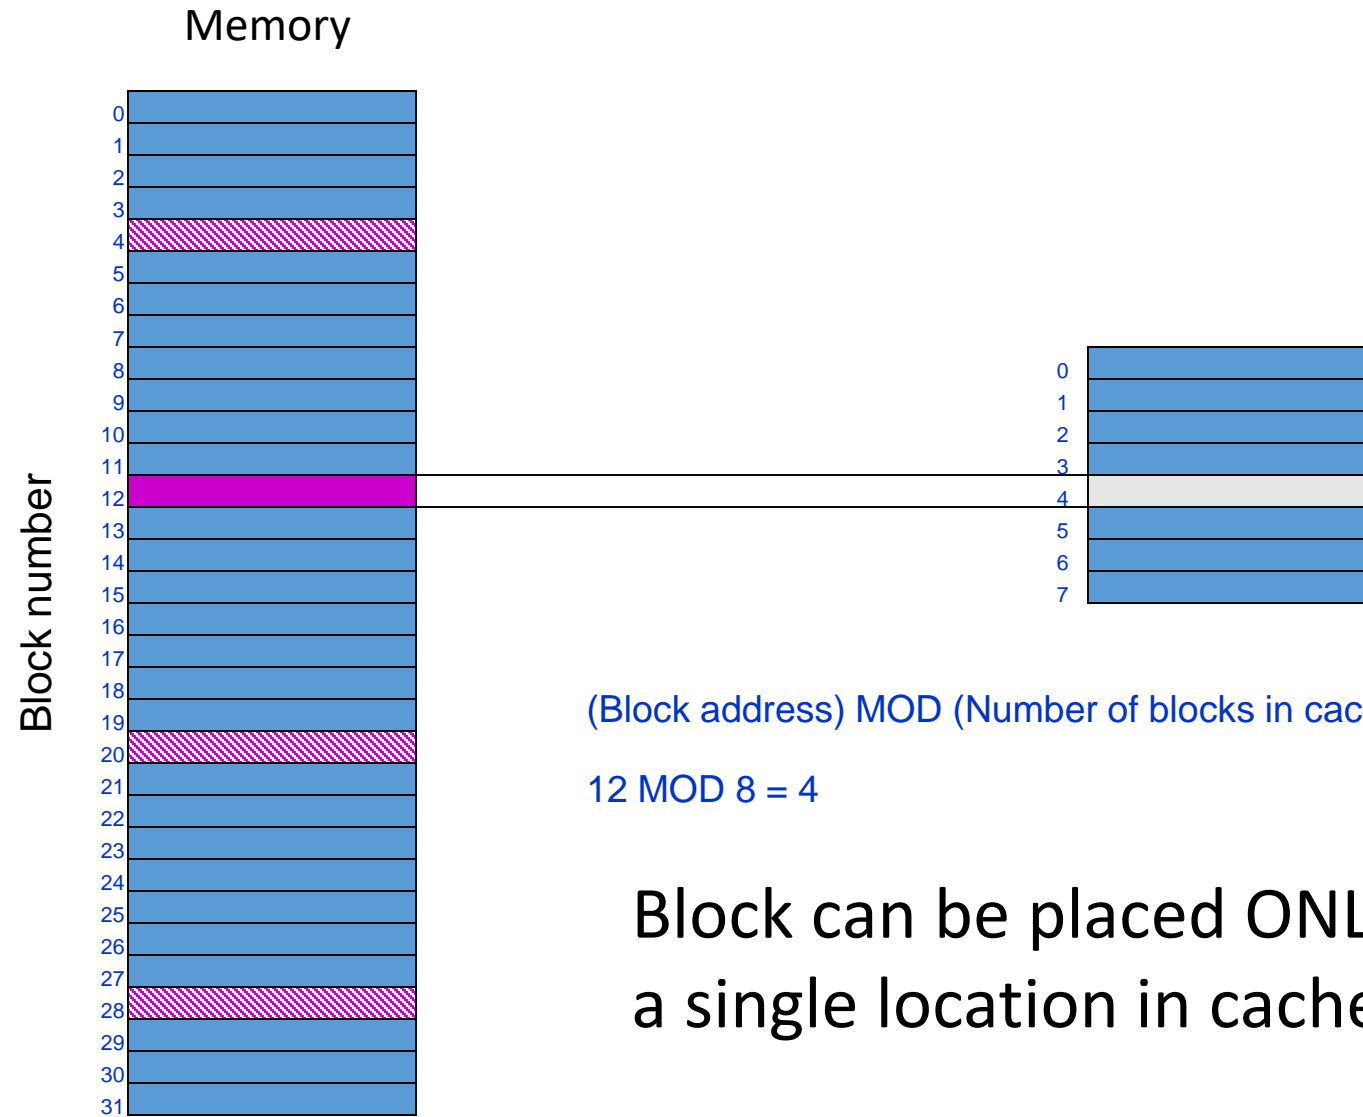
- WHERE to put a block in cache
- **Mapping** between main and cache memories.
- Main memory has a much larger capacity than cache memory.

# Fully Associative Cache

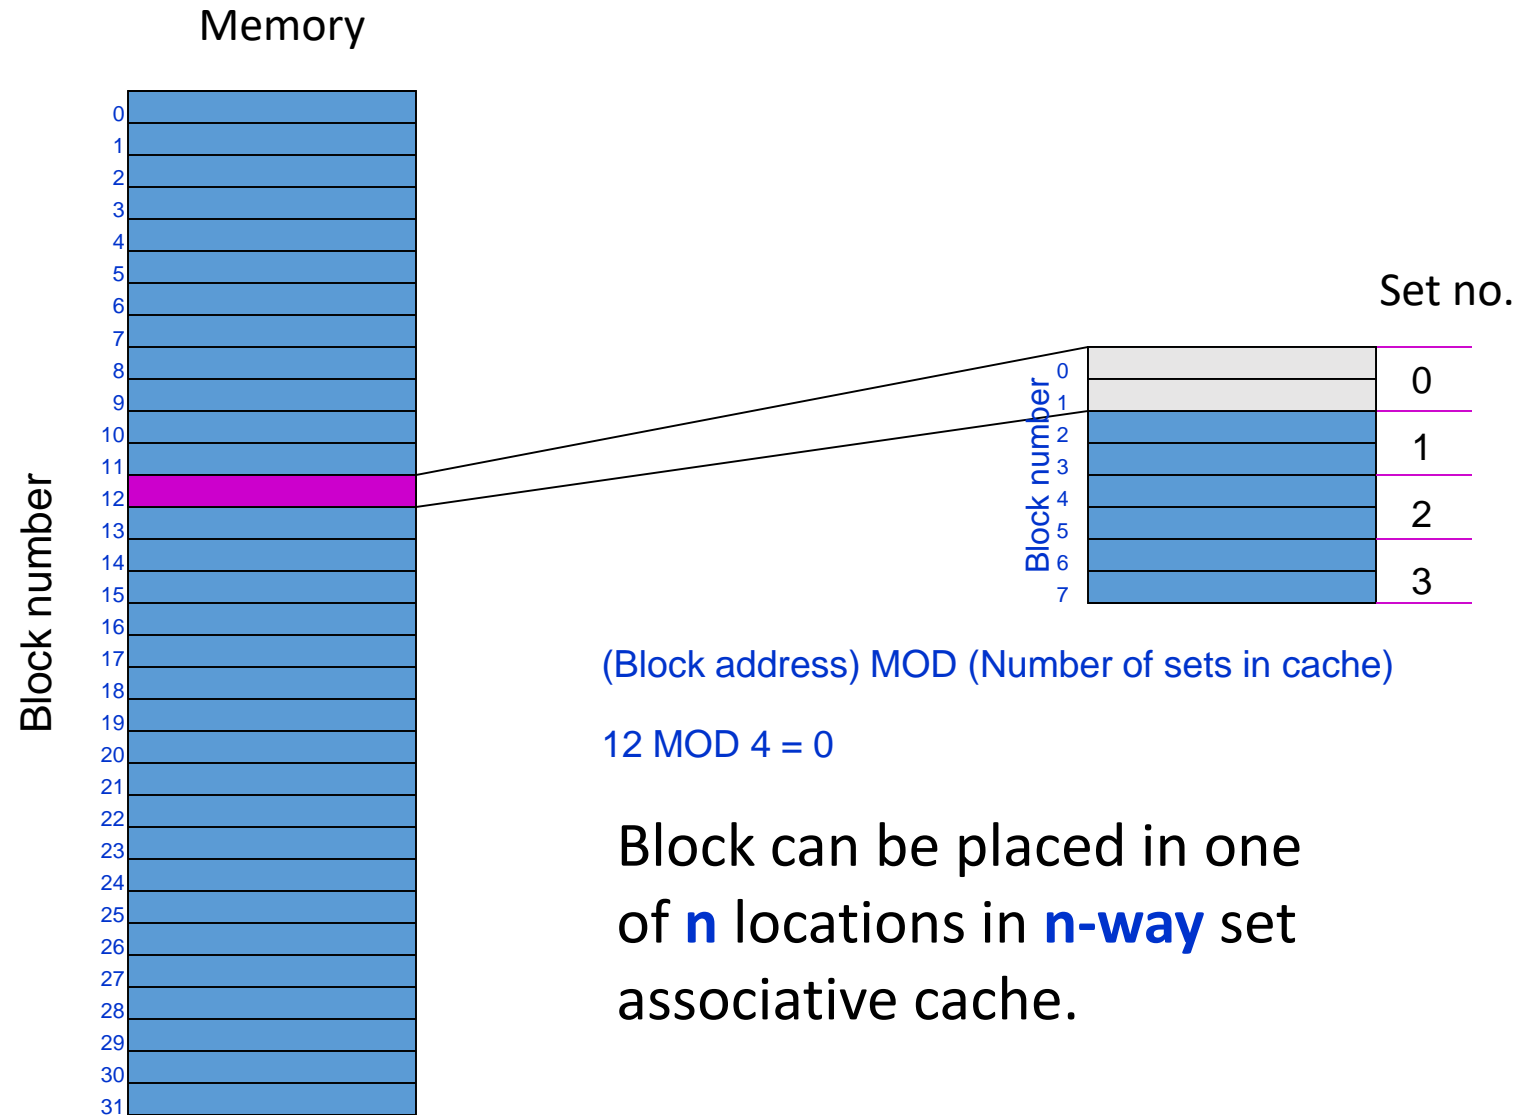




# Direct Mapped Cache



# Set Associative Cache



# Memory Hierarchy Basics


- $n$  sets  $\Rightarrow$   *$n$ -way set associative*
  - *Direct-mapped cache*  $\Rightarrow$  one block per set
  - *Fully associative*  $\Rightarrow$  one set
- Writing to cache: two strategies
  - *Write-through*
    - Immediately update lower levels of hierarchy
  - *Write-back*
    - Only update lower levels of hierarchy when an updated block is replaced
  - Both strategies use *write buffer* to make writes asynchronous

# Dirty bit(s)


- Indicates if the block has been written to.
  - No need in I-caches.
  - No need in write through D-cache.
  - Write back D-cache needs it.



# Write back



# Write through

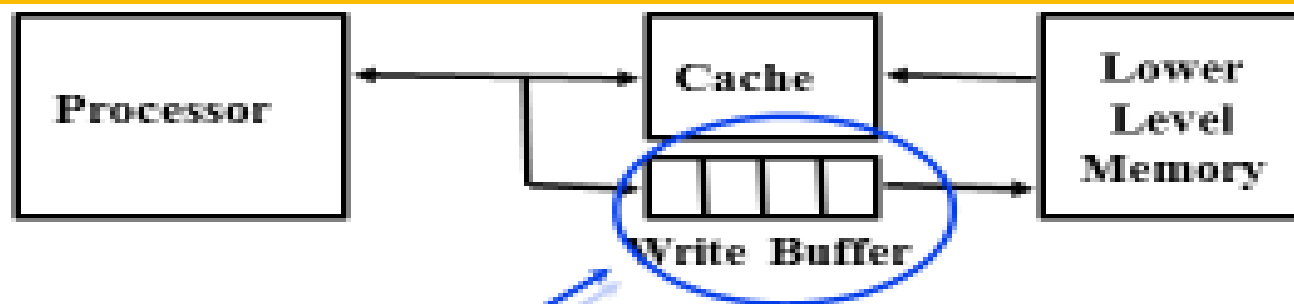


## Q4: What happens on a write?

	Write-Through	Write-Back
Policy	Data written to cache block also written to lower-level memory	Write data only to the cache  Update lower level when a block falls out of the cache
Debug	Easy	Hard
Do read misses produce writes?	No	Yes
Do repeated writes make it to lower level?	Yes	No

**Additional option (on miss)-- let writes to an un-cached address; allocate a new cache line (“write-allocate”).**

# Write Buffers for Write-Through Caches



**Holds data awaiting write-through to lower level memory**

**Q. Why a write buffer ?**

**A. So CPU doesn't stall**

**Q. Why a buffer, why not just one register ?**

**A. Bursts of writes are common.**

**Q. Are Read After Write (RAW) hazards an issue for write buffer?**

**A. Yes! Drain buffer before next read, or send read 1<sup>st</sup> after check write buffers.**



# Memory Hierarchy Basics

- Miss rate
  - Fraction of cache access that result in a miss
- Causes of misses (Three Cs)
  - **Compulsory**
    - First reference to a block
  - **Capacity**
    - Blocks discarded and later retrieved
  - **Conflict**
    - Program makes repeated references to multiple addresses from different blocks that map to the same location in the cache

# Memory Hierarchy Basics


$$\frac{\text{Misses}}{\text{Instruction}} = \frac{\text{Miss rate} \times \text{Memory accesses}}{\text{Instruction count}} = \text{Miss rate} \times \frac{\text{Memory accesses}}{\text{Instruction}}$$

$$\text{Average memory access time} = \text{Hit time} + \text{Miss rate} \times \text{Miss penalty}$$

- Note that speculative and multithreaded processors may execute other instructions during a miss
  - Reduces performance impact of misses

## Q1: Where can a block be placed in the upper level?


- **Block 12 placed in an 8-block cache:**
  - Fully associative, direct mapped, 2-way set associative
  - S.A. Mapping = (Block Number) *Modulo* (Number Sets)



# Direct Mapped Block Placement


---

Cache



address maps to block:

location = (block address MOD #blocks in cache)




Memory


# Fully Associative Block Placement

---

Cache



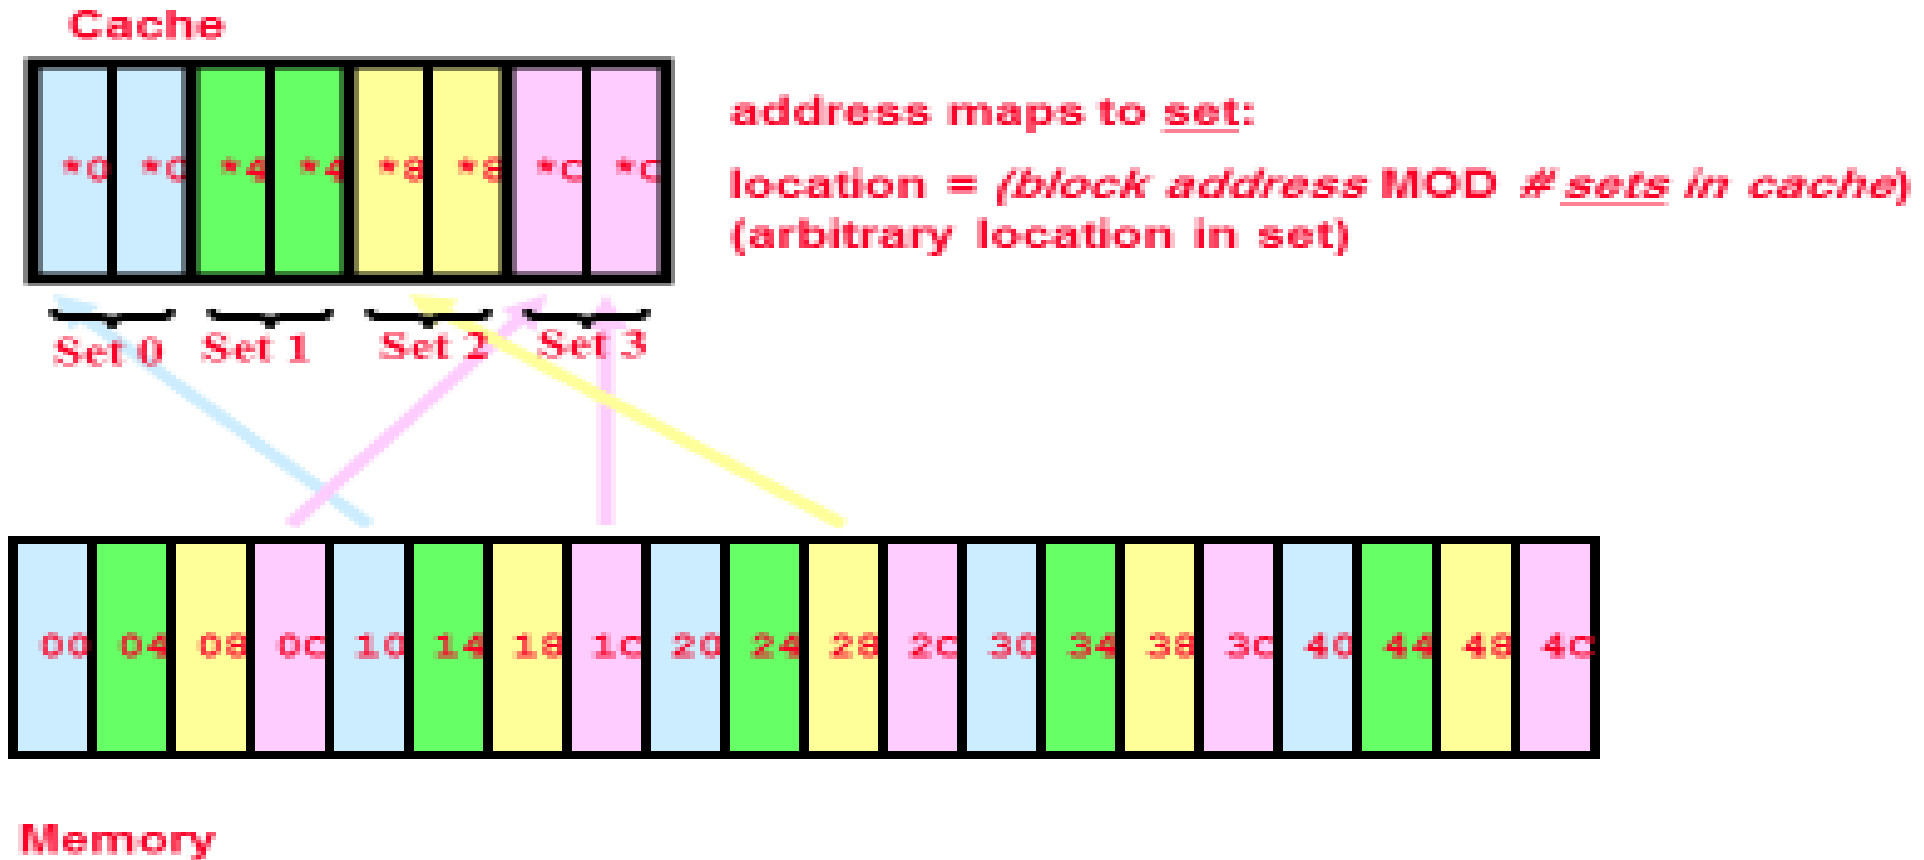
arbitrary block mapping  
location = *any*



Memory

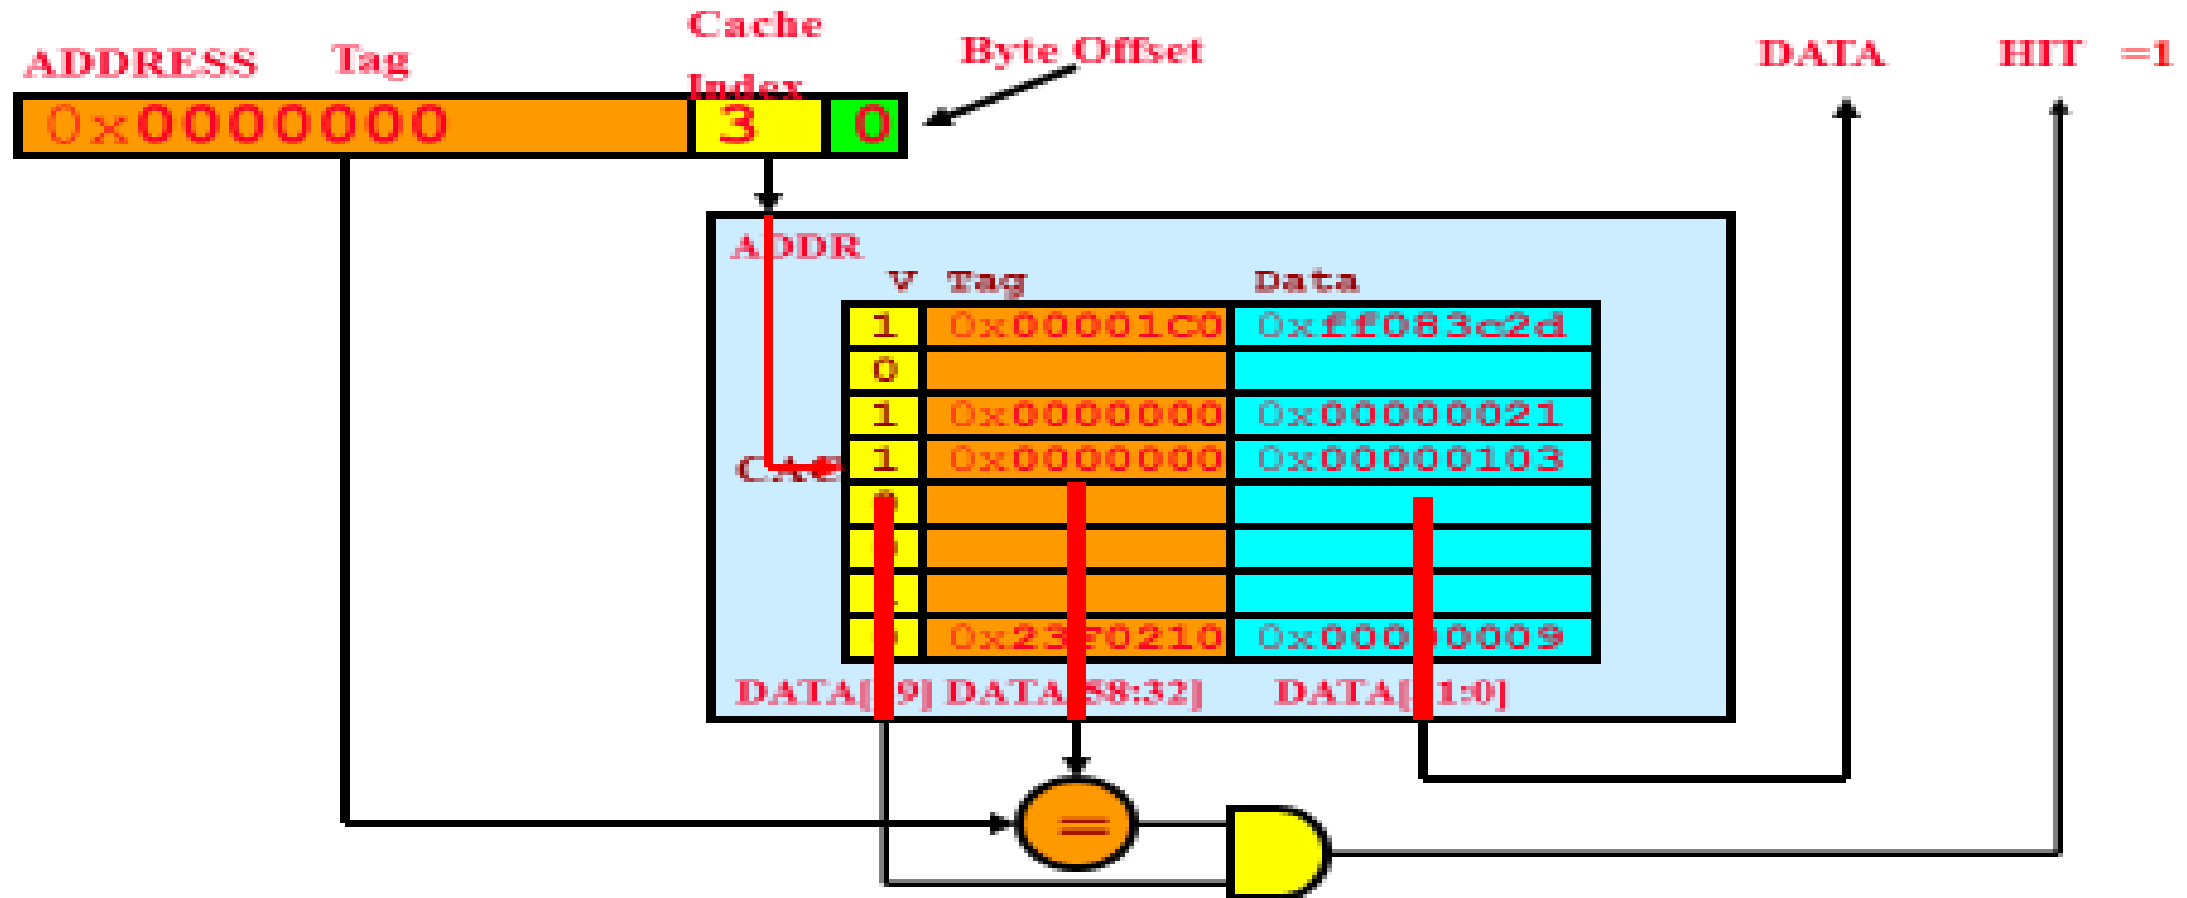
# Set-Associative Block Placement

---





# Direct-Mapped Cache Design

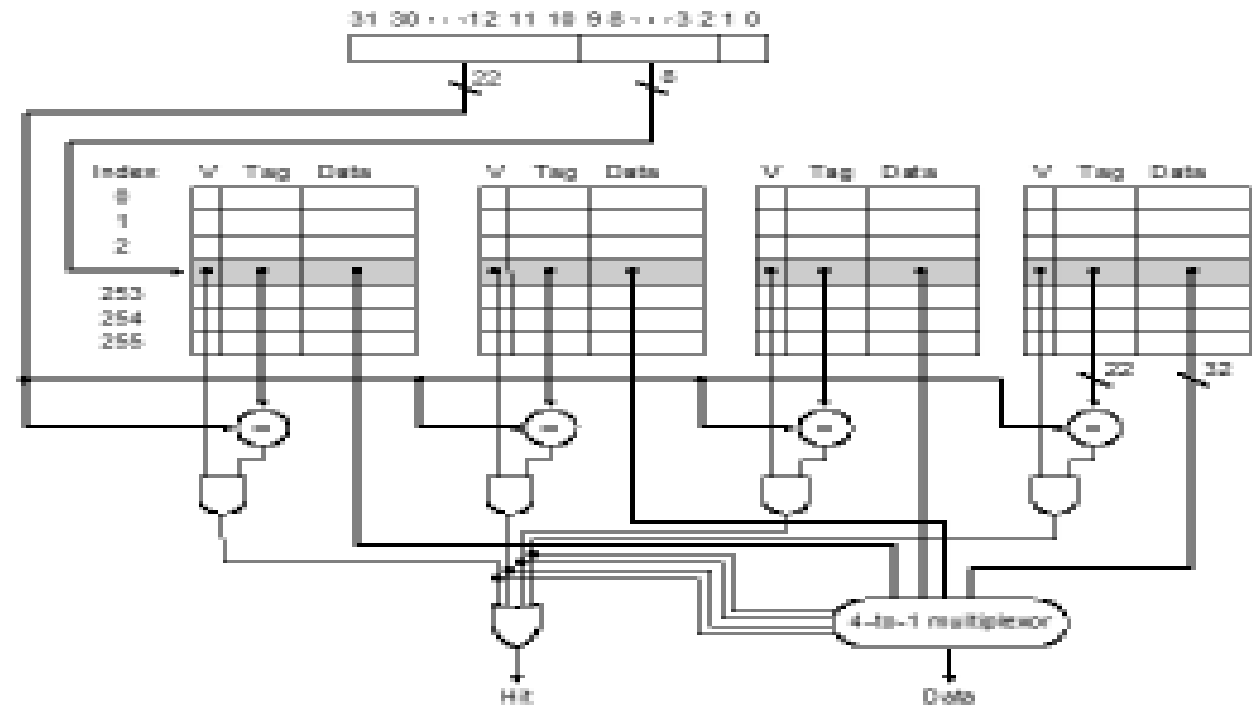




# Four-way cache

## Set Associative Cache Design

- **Key idea:**
  - Divide cache into sets
  - Allow block anywhere in a set
- **Advantages:**
  - Better hit rate
- **Disadvantage:**
  - More tag bits
  - More hardware
  - Higher access time



**A Four-Way Set-Associative Cache**


# Fully Associative Cache Design

---

- Key idea: set size of one block
  - 1 comparator required for each block
  - No address decoding
  - Practical only for small caches due to hardware demands

tag in 11110111

data out 1111000011110000101011



## Q2: How is a block found if it is in the <sup>\*</sup>upper level?

---

- **Tag on each block**
  - No need to check index or block offset
- **Increasing associativity shrinks index, expands tag**



## In-Class Exercise

---

- Given the following requirements for cache design for a 32-bit-address computer: (1) cache contains 16KB of data, and (2) each cache block contains 16 words. (3) Placement policy is 4-way set-associative.
  - What are the lengths (in bits) of the *block offset field* and the *index field* in the address?
  - What are the lengths (in bits) of the *index field* and the *tag field* in the address if the placement is 1-way set-associative?

# Solution

Cache size=16KB= $2^{10} \times 2^4 = 2^{14}$  Bytes

Block size=16 words= $4 \times 16 = 2^2 \times 2^4 = 2^6$  Bytes (word length 32 bit, 4 bytes)

Number of blocks in cache= $\frac{2^{14}}{2^6} = 2^8 = 256$  blocks

4-Way set associative cache then each set has  $\frac{2^8}{2^2} = 2^6$  blocks

Number of bits in offset= $\lg 2^6 = 6$  bits

Number of bits in index field= $\lg 2^6 = 6$  bits

Number of bits in tag field= $32 - 6 - 6 = 20$  bits

## **In one set:**

Number of block from above = 256 block

Number of bits in index= $\lg 2^8 = 8$  bits

Number of bits in tag field= $32 - 8 = 24$  bits

## Q3: Which block should be replaced on a miss?

- Easy for Direct Mapped
- Set Associative or Fully Associative:
  - Random
  - LRU (Least Recently Used)

Assoc:	2-way		4-way		8-way	
Size	LRU	Ran	LRU	Ran	LRU	Ran
16 KB	5.2%	5.7%	4.7%	5.3%	4.4%	5.0%
64 KB	1.9%	2.0%	1.5%	1.7%	1.4%	1.5%
256 KB	1.15%	1.17%	1.13%	1.13%	1.12%	1.12%

**Q3:** After a cache read miss, if there are no empty cache blocks, which block should be removed from the cache?

---

The Least Recently Used (LRU) block? Appealing, but hard to implement for high associativity

A randomly chosen block?  
Easy to implement, how well does it work?

### Miss Rate for 2-way Set Associative Cache

Size	Random	LRU
16 KB	5.7%	5.2%
64 KB	2.0%	1.9%
256 KB	1.17%	1.15%

Also, try other LRU approx.

# Memory Hierarchy Basics

- Six basic cache optimizations:
  - Larger block size
    - Reduces compulsory misses
    - Increases capacity and conflict misses, increases miss penalty
  - Larger total cache capacity to reduce miss rate
    - Increases hit time, increases power consumption
  - Higher associativity
    - Reduces conflict misses
    - Increases hit time, increases power consumption
  - Higher number of cache levels
    - Reduces overall memory access time
  - Giving priority to read misses over writes
    - Reduces miss penalty
  - Avoiding address translation in cache indexing
    - Reduces hit time



Average memory access time = Hit time + Miss rate  $\times$  Miss penalty

Ex: miss rate=0.05, miss penalty=10CC, hit penalty=3CC

$$AMAT = (1 - 0.05) * 3 + 0.05 * 10 = 0.95 * 3 + 0.05 * 10 = 2.85 + 0.5 = 3.35CC$$

If the frequency of CPU=2.5GHz

$$AMAT = 3.35 * 1 / (2.5 * 10^9) = 3.35 * 0.4 * 10^{-9} = 1.34 * 10^{-9} = 1.34ns$$


# Ten Advanced Optimizations

- Metrics:
  - Reducing the hit time
  - Increase cache bandwidth
  - Reducing miss penalty
  - Reducing miss rate
  - Reducing miss penalty or miss rate via parallelism

# 1) Small and simple L1 caches

- Critical timing path:
  - addressing tag memory, then
  - comparing tags, then
  - selecting correct set
- Direct-mapped caches can overlap tag compare and transmission of data
- Lower associativity reduces power because fewer cache lines are accessed

# L1 Size and Associativity



# L1 Size and Associativity

